

# Hands-On with Apache Airflow

## Session 2

# SESSION OVERVIEW

01

Airflow setup/installation on EC2 instance

02

Main motivations behind various operators in Airflow

03

Airflow CLI and UI walkthrough

04

Airflow operators and their demonstrations

# IMPORTANT CONFIGURATIONS

- When you first setup airflow, it creates a file called airflow.cfg
- This file contains Airflow's configurations which can be easily edited to better suit ones requirements
- You can see these configurations from the CLI using:
  - **airflow config** command

# IMPORTANT CONFIGURATIONS

| Config Name             | Description  | Sample Values  |
|-------------------------|--|--|
| dags_folder             | <ul style="list-style-type: none"><li>○ The folder where your airflow pipelines live</li><li>○ This path must be absolute</li></ul>  | /home/ubuntu/airflow/dags                              |
| executor                | <ul style="list-style-type: none"><li>○ The type of executor that airflow should use.</li><li>○ Choices include SequentialExecutor, LocalExecutor, CeleryExecutor</li></ul>  | CeleryExecutor   |
| sql_alchemy_conn        | <ul style="list-style-type: none"><li>○ The SQLAlchemy connection string to the metadata database</li></ul>  | postgresql+psycpg2://airflow:airflow@localhost/airflow |
| parallelism             | <ul style="list-style-type: none"><li>○ The amount of parallelism as a setting to the executor</li><li>○ This defines the max number of task instances that should run simultaneously for the airflow installation</li></ul> | 32   |
| dag_concurrency         | <ul style="list-style-type: none"><li>○ The number of task instances allowed to run concurrently by the scheduler</li></ul>  | 16   |
| max_active_runs_per_dag | <ul style="list-style-type: none"><li>○ The maximum number of active DAG runs per DAG</li></ul>  | 16   |
| base_url                | <ul style="list-style-type: none"><li>○ The base URL of your website as airflow cannot guess what domain or cname you are using.</li></ul>   | http://localhost:8080                                  |
| broker_url              | <ul style="list-style-type: none"><li>○ The Celery broker URL. Celery supports RabbitMQ, Redis and experimentally a sqlalchemy database.</li></ul>   | redis://127.0.0.1:6379/1                               |

# OPERATORS IN AIRFLOW

- Operators **determine what actually gets done** by a task.
- Operators describe a **single task in a workflow** which can be a shell script, a Hive query, a Python function etc.
- The DAG ensures that operators run in the correct order
- They may run on two completely different machines.
- Airflow has a feature of **operator cross communication** known as **Xcom**, which allows two operators to combine or share resources or files
- Operators are only loaded by Airflow if they are assigned to a DAG.

# OPERATOR DEFINITION

```
# Import modules related to DAG and operators
from airflow import DAG
from airflow.operators.python_operator import PythonOperator
from airflow.operators.hive_operator import HiveOperator

# Create a DAG object with necessary configs
ml_dag = DAG('sample', description='Python ML',
             schedule_interval='0 12 * * *',
             start_date=datetime(2017, 3, 20))

# Create task objects with desired task operators
hive_task = HiveOperator(task_id='load_hive_tables', dag=ml_dag, ...)
python_task = PythonOperator(task_id='get_query_result', dag=ml_dag, ...)

# Create task dependencies
hive_task >> python_task
```

# BASH OPERATOR

- The **BashOperator** is used for running any Shell command or script with your DAG.
- It is one of the most versatile and widely used operators in Airflow
- Arguments:
  - **task\_id** : The ID for the task
  - **bash\_command** : The shell command/script to be executed
  - **dag** : The DAG to which this task belongs

# PYTHON OPERATOR


- The **PythonOperator** is used to execute Python callables.
- Arguments:
  - **python\_callable** : the function name to be called
  - **op\_args/op\_kwargs** : to pass additional arguments to the Python callable



# TASK DEPENDENCIES


- If task T1 should finish before T2 starts, you can do this using:

- `T1.set_downstream(T2)`
- `T1 >> T2`
- `T2.set_upstream(T1)`
- `T2 << T1`




- You can also define a chain of dependencies at once:

- `T1 >> T2 >> T3` instead of `T1 >> T2` and `T2 >> T3`



- You can set multiple dependencies at once:

- `T1 >> [T2, T3]`



# PYTHON OPERATOR DAG



# SQOOP OPERATOR

- The **SqoopOperator** is used to execute a Sqoop job
- It is usually used to transfer data between RDBMS and HDFS
- Arguments :
  - **conn\_id** : connection ID
  - **table** : MySQL table name
  - **cmd\_type** : Command type
  - **target\_dir** : Target directory for HDFS

# HIVE OPERATOR

- The **HiveOperator** is used to connect to Hive using `hive_conn_id` and execute Hive queries
- Arguments:
  - **hql** : Hive query to be executed
  - **hive\_cli\_conn\_id** : Hive connection

# SPARK OPERATOR

- The Spark operators are used to schedule spark jobs using Airflow
- **SparkSubmitOperator** launches Spark using the spark-submit CLI on the airflow machine
- It supports all the configurations and arguments needed by spark-submit
- To run Spark SQL, we can make use of **SparkSqlOperator**
- Arguments for SparkSubmitOperator :
  - **application** : The application to be submitted as a job, either jar or py file
  - **conn\_id** : Connection ID
  - **application\_args**: Arguments for the application being submitted

# EMAIL OPERATOR


- The **EmailOperator** is used for alerting task events like task completion, task failure, task retry, task fail on retry
- It is also used to send attachments along with emails with error messages or logs
- Arguments :
  - **to** : list of receiver emails
  - **subject** : subject for the email
  - **html\_content** : content of the email, html markup is used
  - **files** : files to attach in the email

# BASE OPERATOR

- The **BaseOperator** is the abstract base class for all operators, so data members and methods declared in this class are available to all the operators
- It contains methods such as `set_downstream()` and `set_upstream()`
- You can create a custom operator by extending the **`airflow.models.baseoperator.BaseOperator`** and then overriding the `execute()` method
- Some important arguments received by the BaseOperator include :
  - `task_id`, `owner`, `email`, `email_on_retry`, `retries`, `depends_on_past`, `dag`, `params`, `default_args`, `queue`, `sla`, `execution_timeout`, `on_failure_callback`, `on_success_callback`, `on_retry_callback`, `trigger_rule`, `run_as_user`, `task_concurrency`, `do_xcom_push`

# SENSOR OPERATOR

- Sensors are useful to check for a certain condition before we run a task.
- Eg: we have a task in a DAG depending on a task in another DAG or we have a task which should run only after a file has arrived in HDFS
- Sensor tasks continue to execute at a time interval and succeed when a criteria is met or fail when they timeout
- Arguments:
  - **soft\_fail** : if True marks the task as SKIPPED on failure
  - **poke\_interval**: interval between each try (> 1 minute recommended)
  - **timeout** : time before the sensor fails





# SENSOR OPERATOR EXAMPLES

- **WebHdfsSensor**
  - Waits for a file or folder to land in HDFS
- **ExternalTaskSensor**
  - Waits for a different DAG or a task in a different DAG to complete for a specific execution\_date
- **TimeDeltaSensor**
  - Waits for a timedelta after the task's execution\_date + schedule\_interval
- **TimeSensor**
  - Waits until the specified time of the day

# SESSION SUMMARY

**01**

**Airflow Setup**

**02**

**Airflow Configurations**

**03**

**UI Walkthrough**

**04**

**Operators in Airflow with individual demonstrations**

**THANK YOU**