

SparkSubmitOperator Demonstration

This document will guide you through the demonstration of the SparkSubmitOperator in Session 2 of the Airflow module.

Prerequisites:

- Hive table created in the HiveOperator demonstration (credit_card.filtered_transaction).
- sample_spark.py(the code explained in the video)
- transactions_per_item.py (The spark application)
- The version of the enabled JDK is 8.

What are we doing?

In this demonstration, we need to create a DAG with one Spark task.

- transactions_per_item - It will execute a spark job to analyze how many transactions a particular item has been a part of.

Please follow the instructions below:

1. Login to your EMR instance.
2. Activate the Python virtual environment using the following command:

```
source /home/hadoop/airflow/bin/activate
```

3. Now in this demonstration, we will be using the hive tables we created in the HiveOperator demonstration.

You can use the following command to check it:

```
hive -e "select * from credit_card.filtered_transactions;"
```

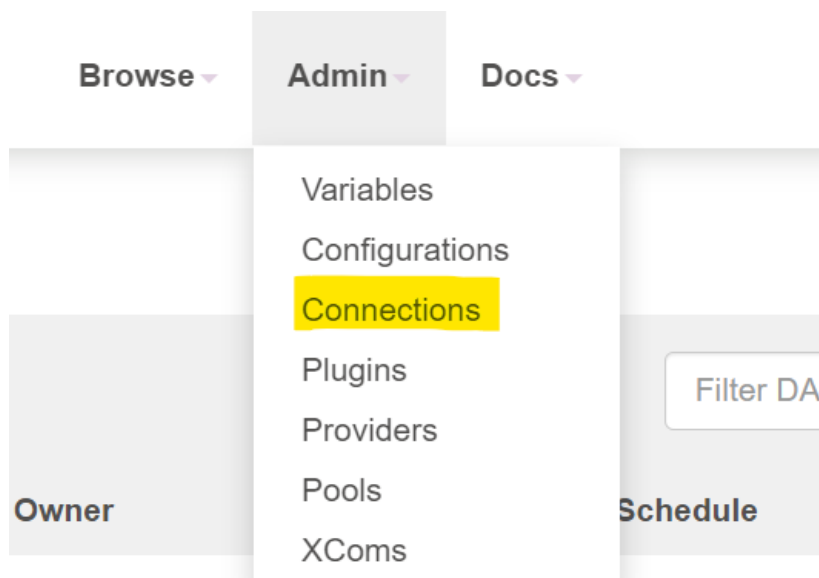
```
hive> select * from filtered_transactions;
OK
SLF4J: Failed to load class "org.slf4j.impl.StaticLoggerBinder".
SLF4J: Defaulting to no-operation (NOP) logger implementation
SLF4J: See http://www.slf4j.org/codes.html#StaticLoggerBinder for
1      U101      I301      1600598377      1600599217      20.0
3      U102      I305      1600588312      1600599326      -100.0
4      U103      I307      1600588342      1600599332      20.0
5      U105      I303      1600588361      1600599325      40.0
7      U107      I302      1600588352      1600599337      60.0
8      U103      I305      1600588336      1600599353      30.0
9      U107      I302      1600588354      1600599338      10.0
10     U105      I302      1600588317      1600599326      50.0
Time taken: 0.298 seconds, Fetched: 8 row(s)
hive> |
```

If you don't have these tables, you will have to finish the HiveOperator demonstration first and then you can continue to the next step

4. In the /home/hadoop directory, we need to place the transactions_per_item.py file. (You can use WinSCP or create a new file and paste the code in that file)
5. Next, you need to set up the Spark connection from the Airflow UI which is hosted in the URL : **your_public_dns:8082**

Note: You can find you_public_ip in your AWS EMR dashboard (IPv4 Public DNS))

Go to the Admin tab and click on Connections



Now click on the edit button to the left of the **spark_default** connection.

<input type="checkbox"/>	 	segment_default	segment
<input type="checkbox"/>	 	sftp_default	sftp
<input type="checkbox"/>	 	spark_default	spark
<input type="checkbox"/>	 	sqlite_default	sqlite
<input type="checkbox"/>	 	sqoop_default	sqoop
<input type="checkbox"/>	 	ssh_default	ssh

Next, fill in the following details and click on Save

Connection Id *	spark_default
Connection Type *	Spark ▼ Connection Type missing? Make sure you've installed the correspond
Description	
Host	yarn
Port	
Extra	{"master": "yarn", "conf": "/etc/spark/conf/spark-defaults.conf"}

Conn Id: spark_default

Conn Type: Spark (Select from the drop-down)

Host: yarn

Extra: `{"master": "yarn", "conf": "/etc/spark/conf/spark-defaults.conf"}`

- Now you need to place the **sample_spark.py** file in the **/home/hadoop/airflow/dags** directory. (You can use WinSCP or create a new file and paste the code in that file)
- To ensure that the file there are no issues/errors with the file is it considered good practice to compile the program using the following command:

python sample_spark.py

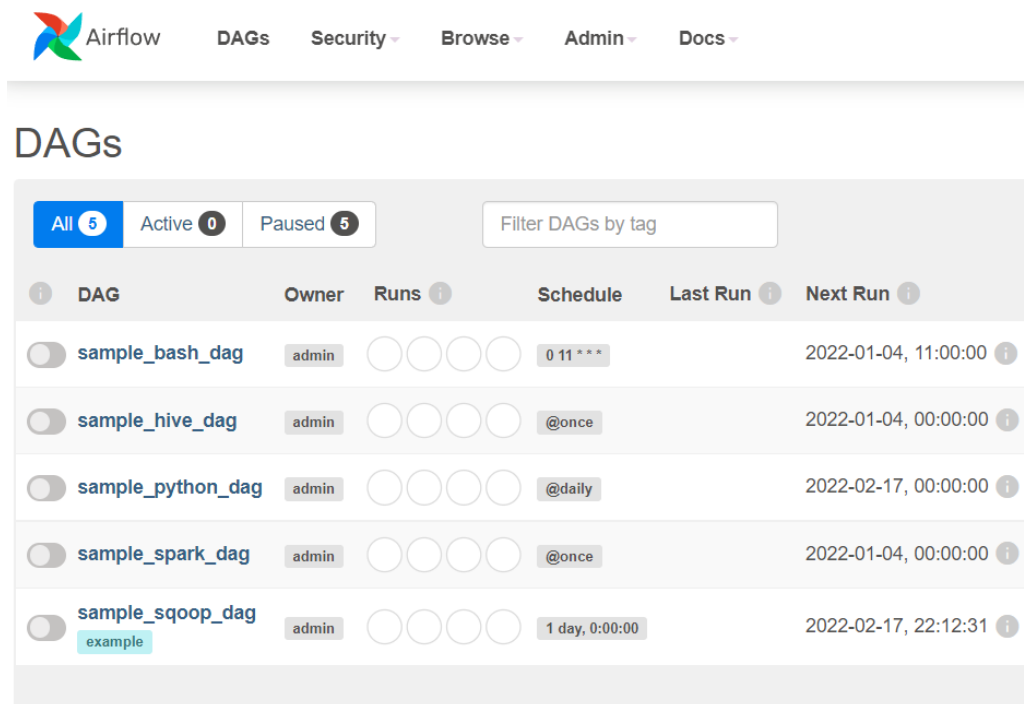
- You can also use the following command to list the dags in your instance:

airflow dags list

- Once you have made sure that your dag file has no issues you can go to the Airflow UI which is hosted in the URL : **your_public_dns:8082**

Note: You can find you_public_ip in your AWS EMR dashboard (IPv4 Public DNS)

- Switch ON the DAG(sample_spark_dag)




The screenshot shows the Airflow web interface. At the top, there's a navigation bar with links: Airflow, DAGs, Security, Browse, Admin, and Docs. Below this, the 'DAGs' section is active. It features a filter bar with 'All' (5), 'Active' (0), and 'Paused' (5) buttons, along with a 'Filter DAGs by tag' input field. The main content is a table listing DAGs:

DAG	Owner	Runs	Schedule	Last Run	Next Run
<input type="checkbox"/> sample_bash_dag	admin	0 11 ***			2022-01-04, 11:00:00
<input type="checkbox"/> sample_hive_dag	admin	@once			2022-01-04, 00:00:00
<input type="checkbox"/> sample_python_dag	admin	@daily			2022-02-17, 00:00:00
<input type="checkbox"/> sample_spark_dag	admin	@once			2022-01-04, 00:00:00
<input type="checkbox"/> sample_sqoop_dag example	admin	1 day, 0:00:00			2022-02-17, 22:12:31


(Note: The DAG might take a while to show up on the UI. Keep refreshing and wait patiently)

11. Click on the sample_spark_dag and go to the graph view

You will see the task is running


DAG: sample_spark_dag
Demo of Spark Operator

Tree
Graph
Calendar
Task Duration
Task Tries


2022-01-04T00:00:01Z


Runs 25

Run scheduled_

SparkSubmitOperator

transactions_per_item

Click on refresh and eventually, it will have successfully completed


DAG: sample_spark_dag
Demo of Spark Operator

Tree
Graph
Calendar
Task Duration
Task


2022-01-04T00:00:01Z

Runs 25

Run sched

SparkSubmitOperator

transactions_per_item

12. Once the DAG has completed execution, the output will be generated in the following HDFS location:

hdfs:///data/credit_card/transactions_per_item

13. You can view the results in the CLI by using the following command:

hdfs dfs -cat hdfs:///data/credit_card/transactions_per_item/*

```
(airflow) [hadoop@ip-172-31-58-178 ~]$ hdfs dfs -cat hdfs:///data/credit_card/transactions_per_item/*
{"item_id":"I305","count":2}
{"item_id":"I307","count":1}
{"item_id":"I303","count":1}
{"item_id":"I302","count":3}
{"item_id":"I301","count":1}
(airflow) [hadoop@ip-172-31-58-178 ~]$ |
```

14. You can switch off your DAG if you don't want it to run anymore.