Creating a Redshift Cluster

1. Choose the Redshift option from the Amazon Dashboard.

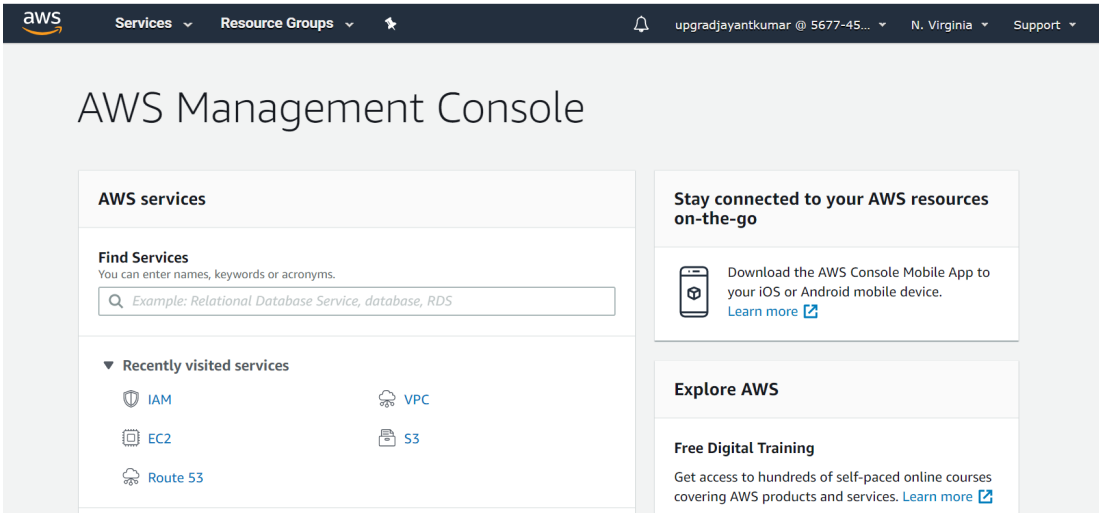

2. You will be redirected to the **AWS Redshift Console**.

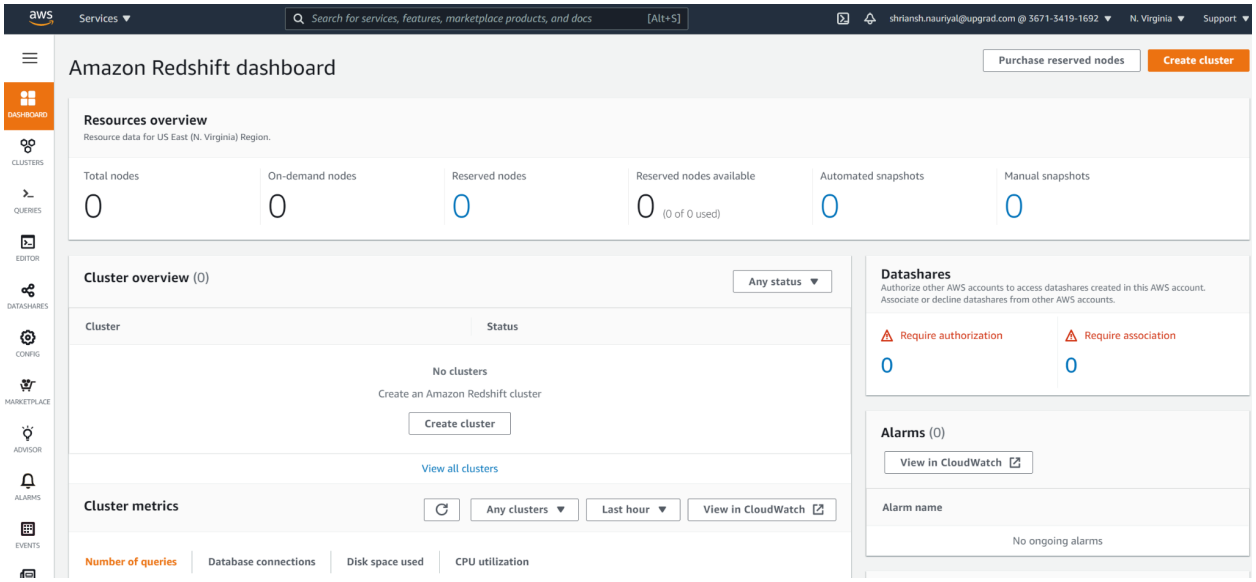

3. Go to **Config**.

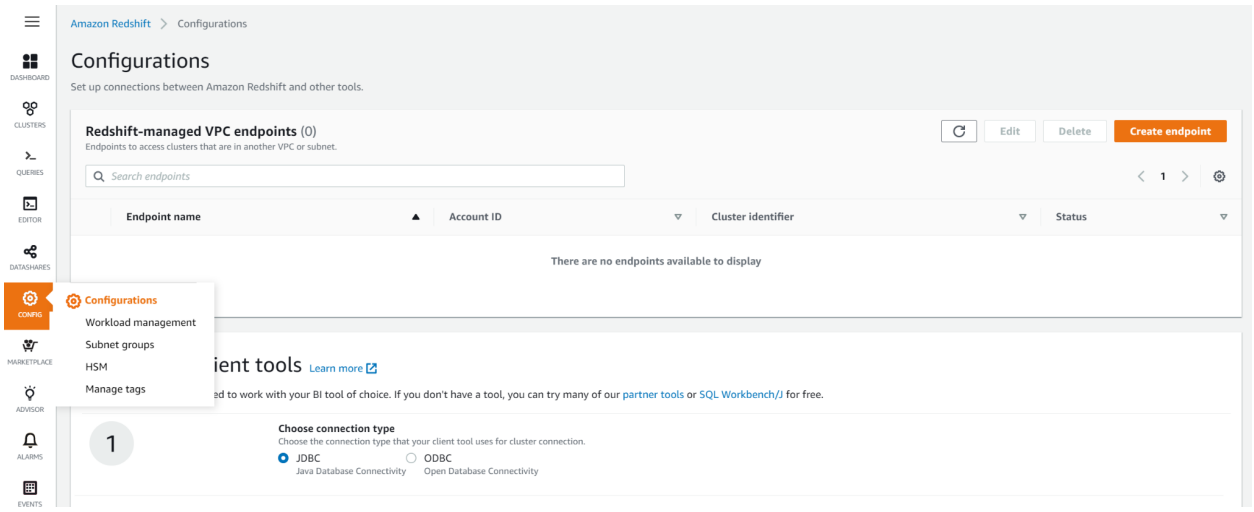

4. In the **Config** dashboard, scroll down to the **Subnet Groups.** Click on '**Manage subnet groups'**.

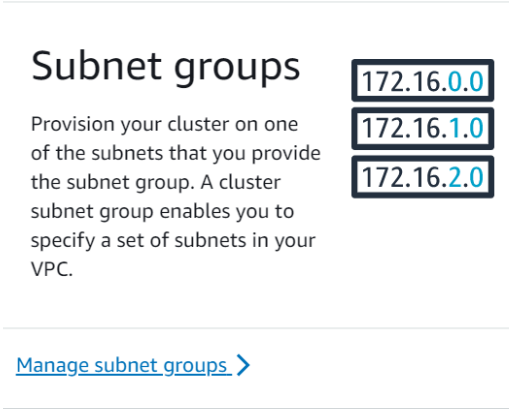

5. Click on the '**Create cluster subnet group'** button.

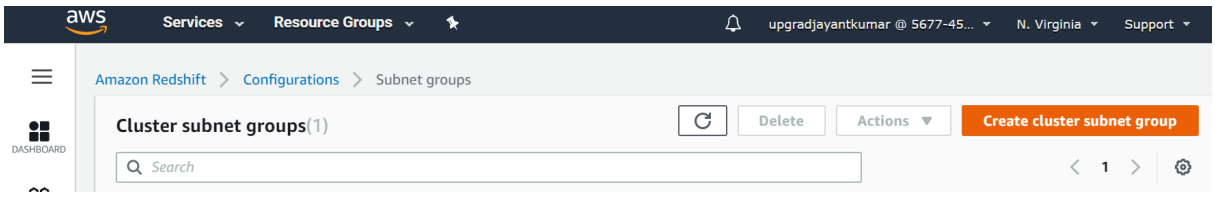

6. Assign a name for the Redshift cluster along with a brief description.

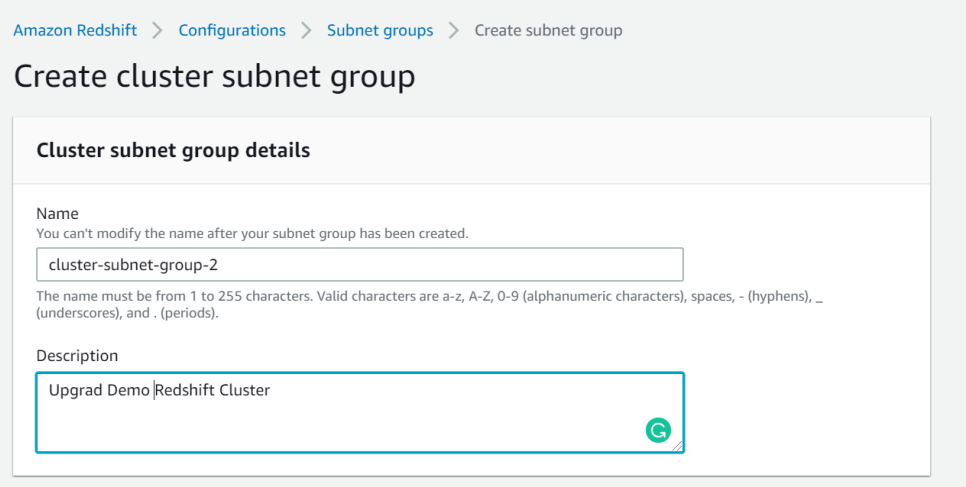

7. Scroll down until you reach the **Add subnets** option. Select your VPC (that you have created previously) from the drop-down list and then click on the **Add all the subnets for this VPC** button.

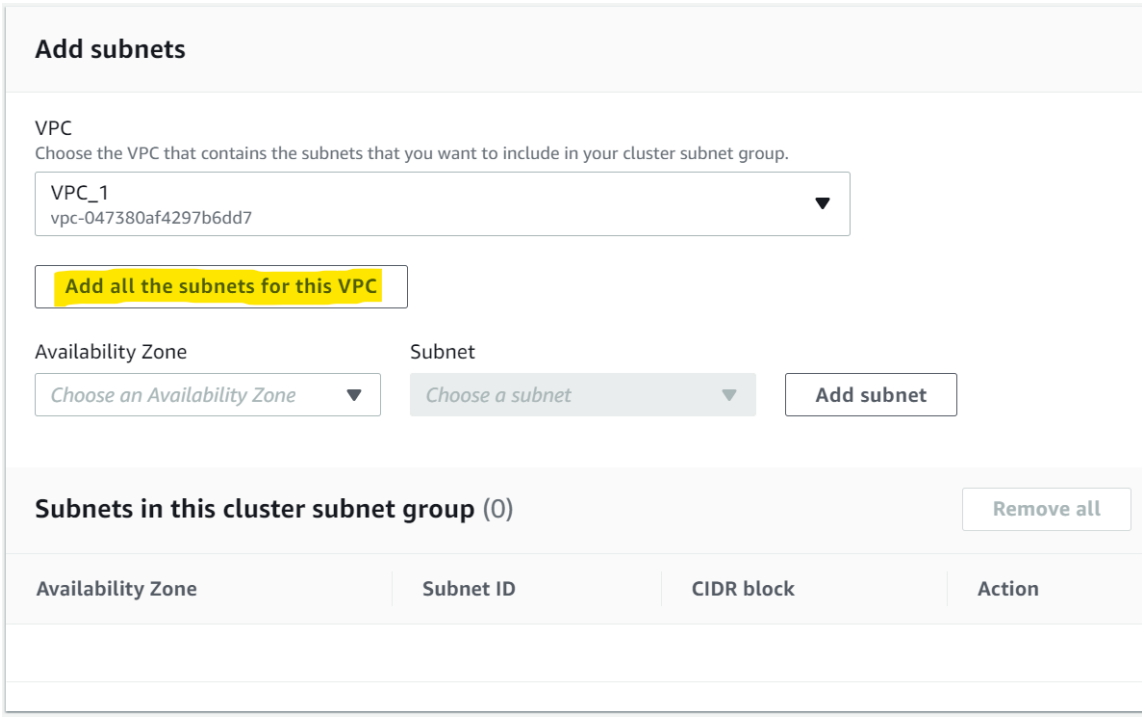

8. After that an available **Availability Zone** appears; select any one. Then, select the **Subnet** drop-down list that appears. Lastly, click on the **Create cluster subnet group**.

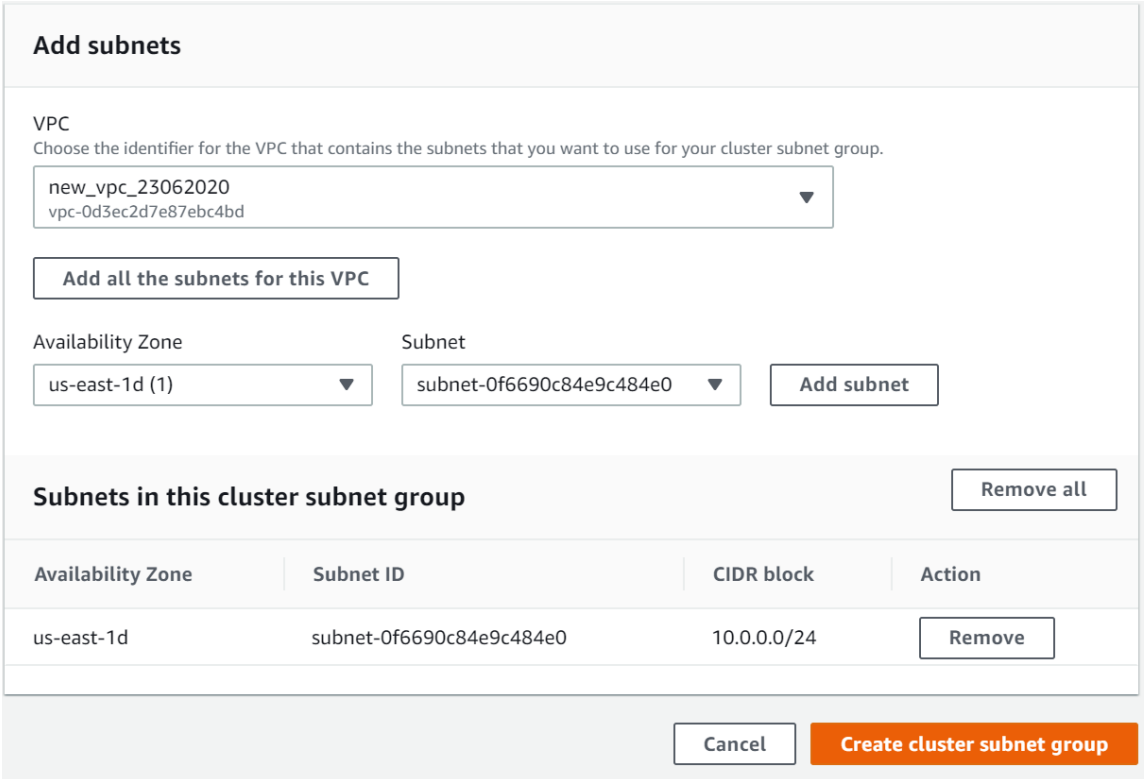

9. The cluster subnet has been created successfully.

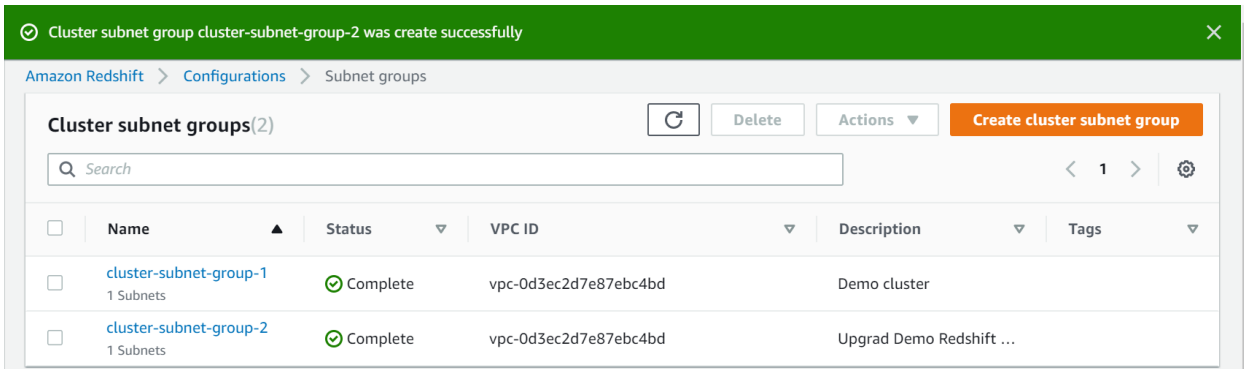

10. Now, you need to go back to the Redshift dashboard. After that, select the **Cluster** option.

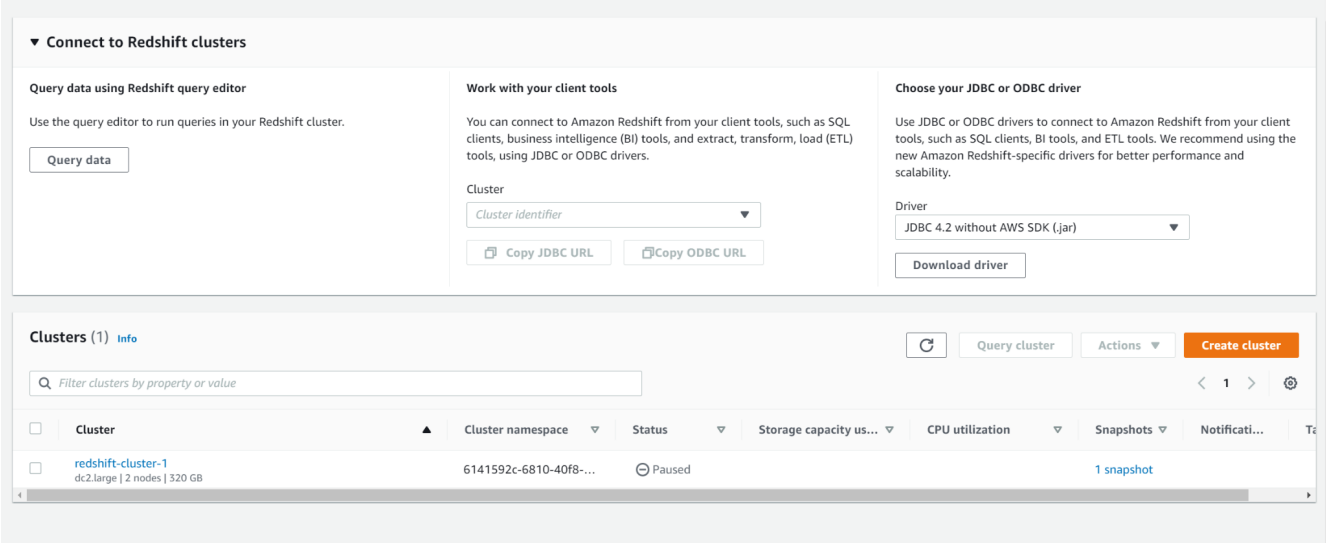

11. Click on **Create Cluster,** assign a name for your cluster and keep the next part as **Production** and then scroll down.

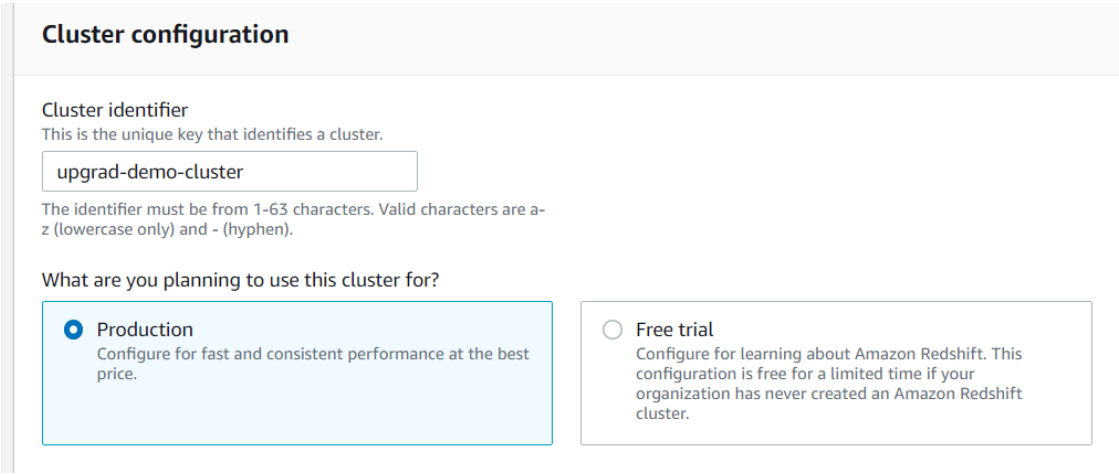

12. Choose a **Node type.** Here, you need to select the **dc2.large** cluster. After that, set the number of nodes to 2. If, by default, it is set to 2, then do not change it.

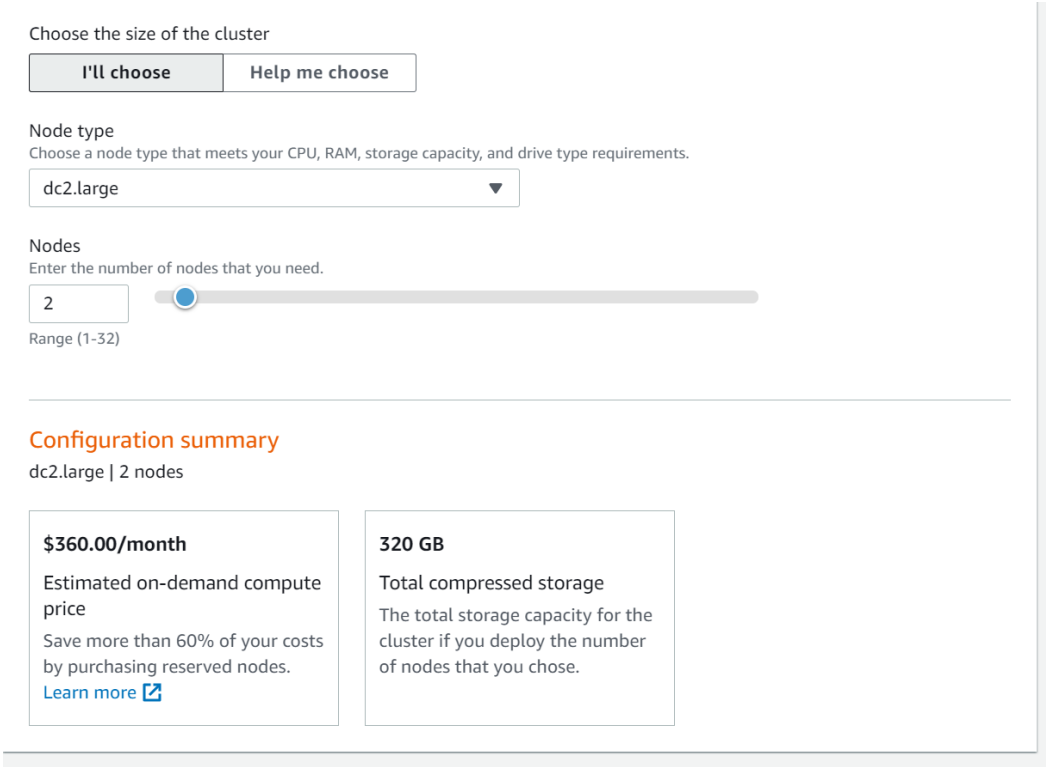

13. Scroll down, and perform **Database configuration.** Set a password with the combination of **number**, **lowercase** letter and **uppercase** letter whose length should be **8–64** characters. You can reset the password once your cluster is created successfully.

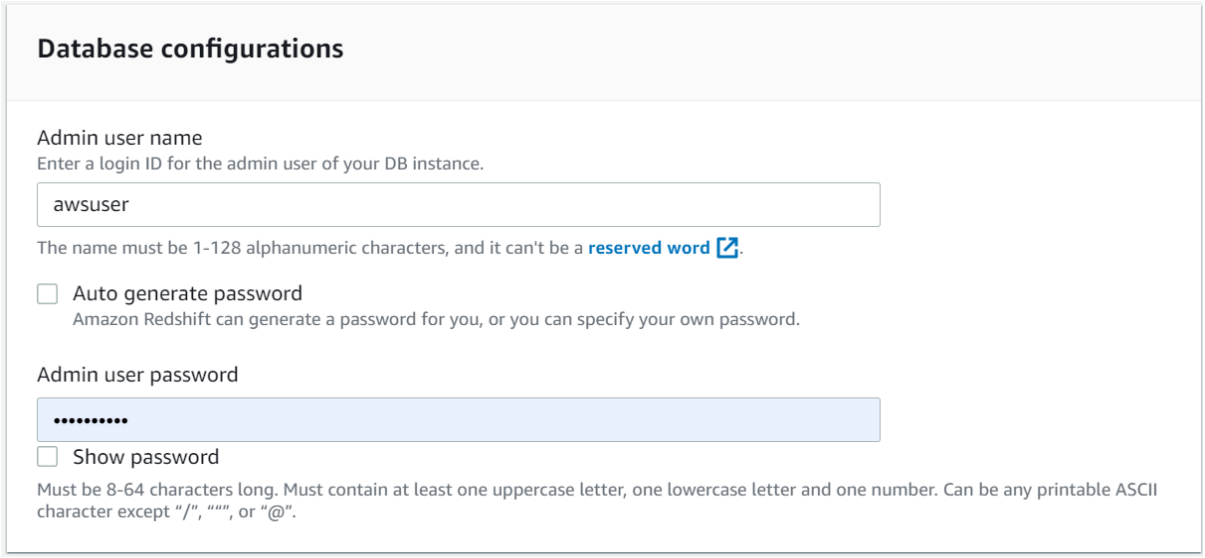

14. Scroll down to the Additional configurations section and click on the button to **deselect Use defaults**.

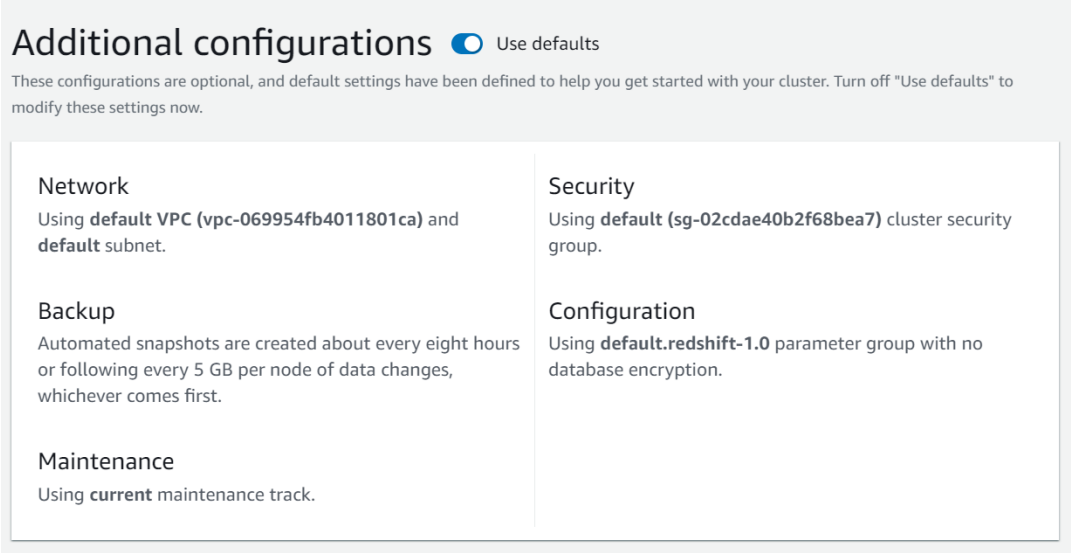

15. Click on **Network and Security** under Additional configurations**.** Here, you need to select the **VPC**, followed by the **Security Group** and then the **subnet group** that you created in step 8**.**

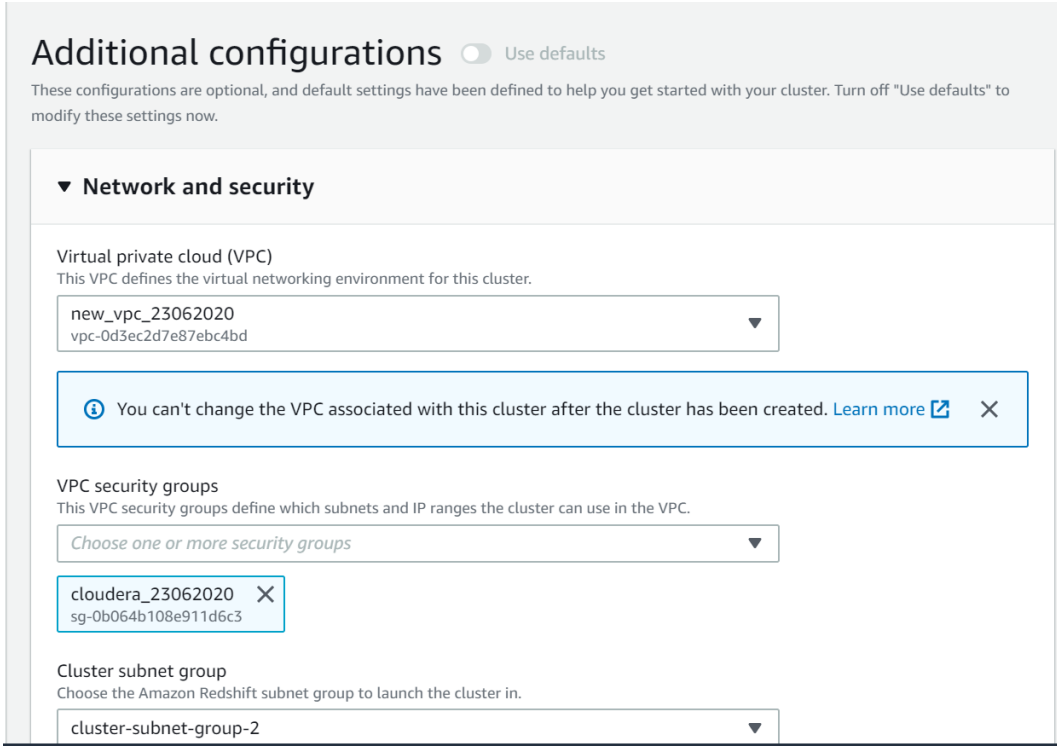

16. Next, click on the Database configurations under additional configurations and here, set a database name. The default port is **5439,** which is known globally. So, the best practice is to change the port number and set it to any one between **1150** and **65535**. This prevents any foreign attack or stops the entrance of intruders to your cluster. Keep all the parameters as the default ones. After this you can click on the **Create Cluster button.**

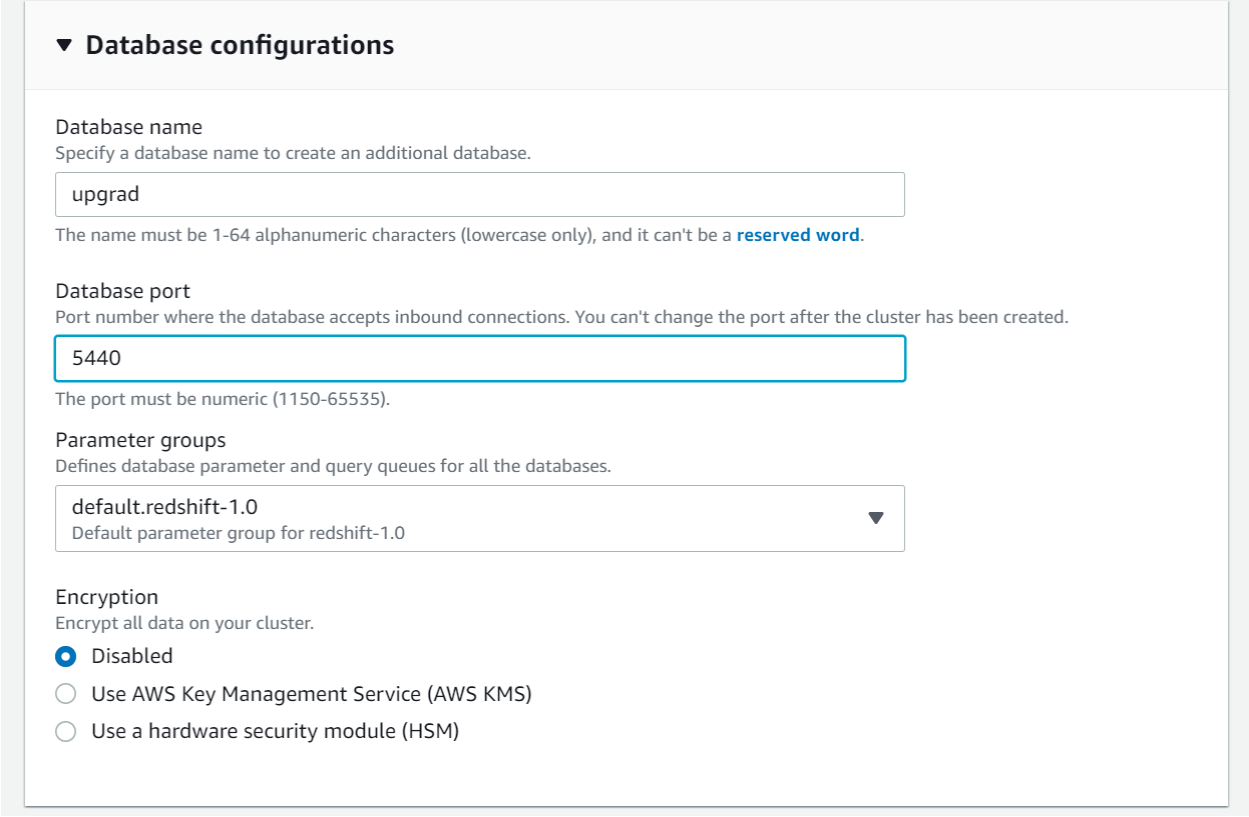

17. The creation of the cluster will take some time

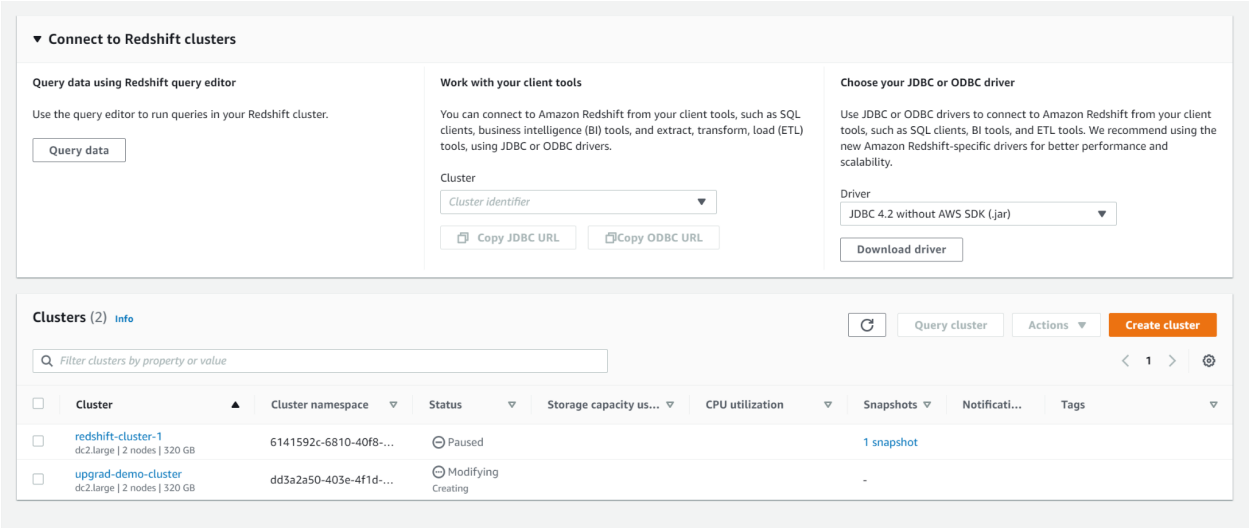

18. While configuring the node type **dc2.large**, you saw that the pricing tag is for per hour. So, as a best practice and for cost-cutting, pause your cluster if you are not using it.

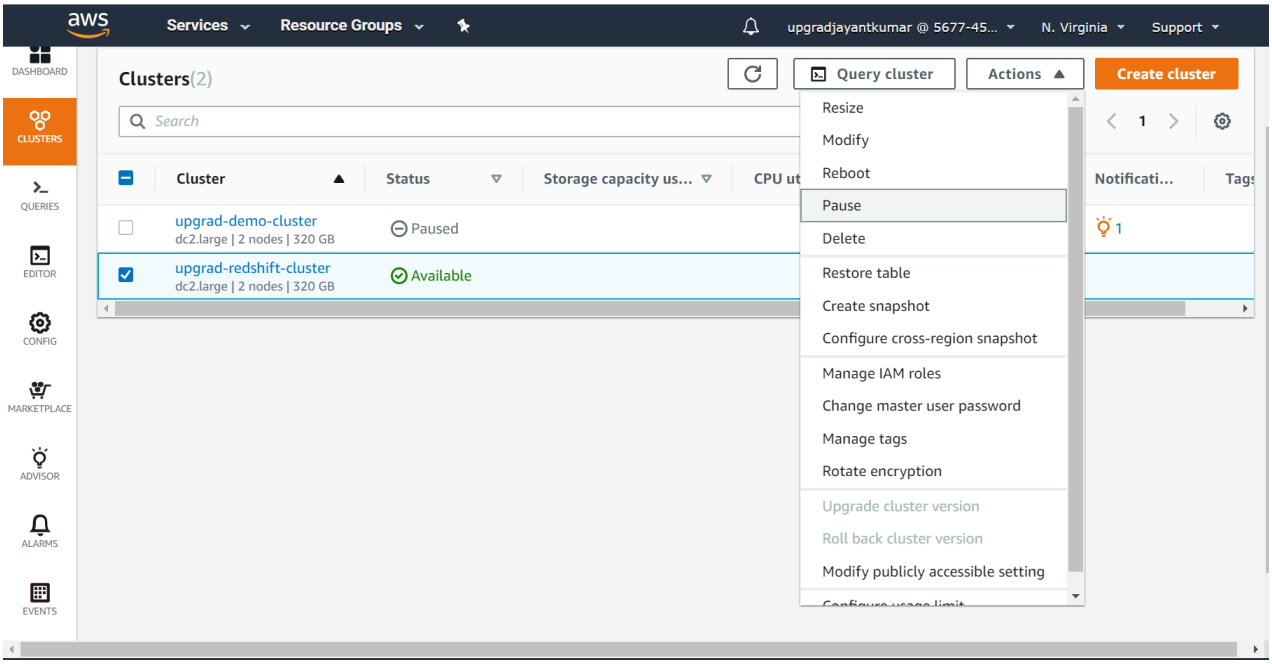

19. Click on your cluster and Go to the **Properties** tab. There you can scroll down a little and you will get the Attach **IAM Role** button.

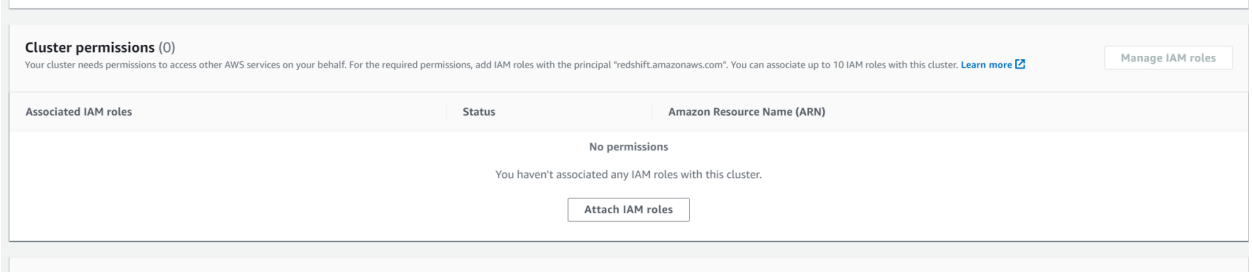

20. You will redirect to the IAM Role dashboard. Select the IAM role that you have created for S3 access from the drop-down list of **Available IAM roles** and click on **Associate IAM Role.** Lastly, click on **Save Changes.**

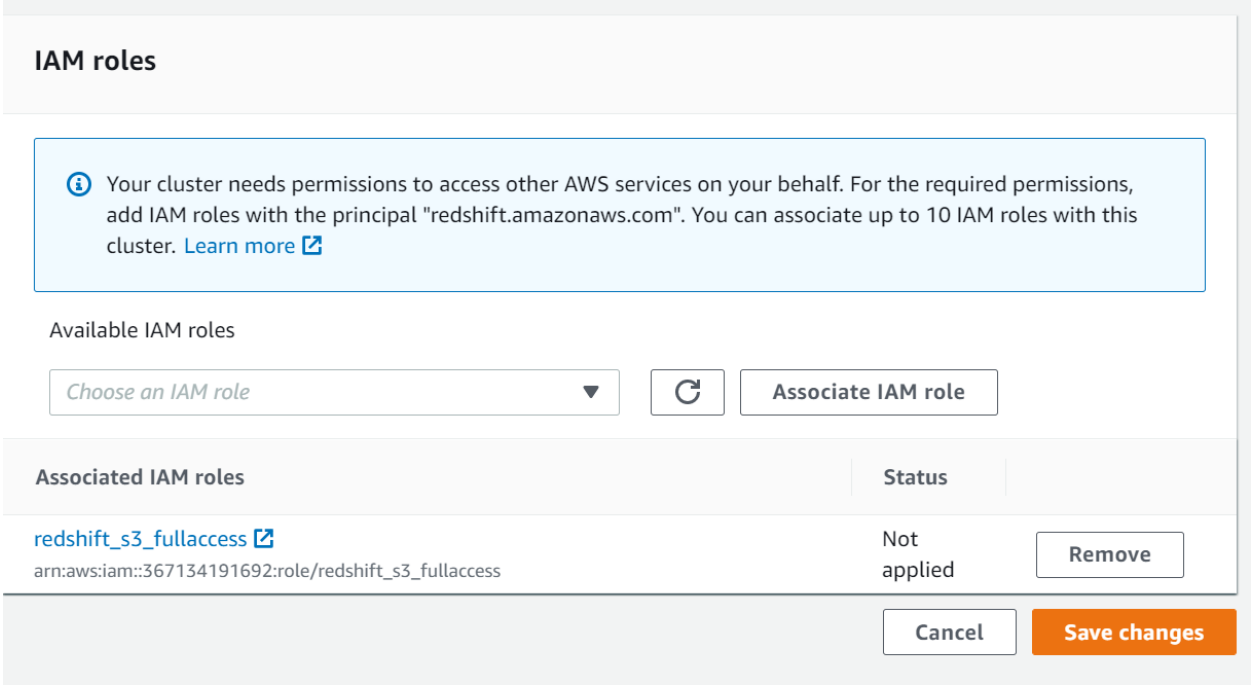

21. If the IAM role is added successfully then you will get status **in-sync**, as shown in below screenshot. This will take a few minutes till the role has been added successfully to the Redshift cluster.