upGrad

Data Ingestion with Apache Sqoop and Apache Flume - Session 1

upGrad

Course: DS - DE

Lecture On: Introduction to Data Ingestion

Instructor: Hitesh Hasija

upGrad

Subject Matter Expert

Hitesh Hasija

Senior Data Engineer: Intuit

Skilled in softwares such as Hadoop, Hive, Sqoop, Spark, Kafka, Flume, Cassandra and MongoDB. I am working in the Big Data domain from last five years.

upGrad

Segment - 01
Module Introduction

Module Introduction

upGrad

Session 1

- What is data ingestion?
- Challenges faced in data ingestion
- Key steps in data ingestion
- Tools used for data ingestion
- Types of data and file formats

Session 2

- Introduction to Sqoop and its advantages and architecture
- Case study introduction
- Setup of Apache Sqoop and Database
- Sqoop export and import
- Various arguments of the Sqoop import command

Module Introduction

upGrad

Session 3

- Additional arguments and options of Apache Sqoop import commands
- Support of SQL queries in Sqoop
- Incremental import in Sqoop
- Sqoop Jobs
- Tuning Sqoop

Session 4

- Introduction to Apache Flume as well as its components and characteristics
- Case study and installation of Flume
- Flume Configuration files and Flume flows
- Tuning Flume and Sqoop vs Flume

upGrad

Segment - 02
Session Overview

Session Overview

upGrad

1. What is data ingestion and why it is needed?
2. Challenges faced while designing systems for data ingestion
3. Key steps of data ingestion
4. Tools used for data ingestion in the industry
5. Types of data and file formats dealt with in data ingestion

upGrad

Segment - 03
What is Data Ingestion?

Learning Objectives

upGrad

1. What is Data Ingestion?
2. Why is Data Ingestion needed?

What is Data Ingestion?

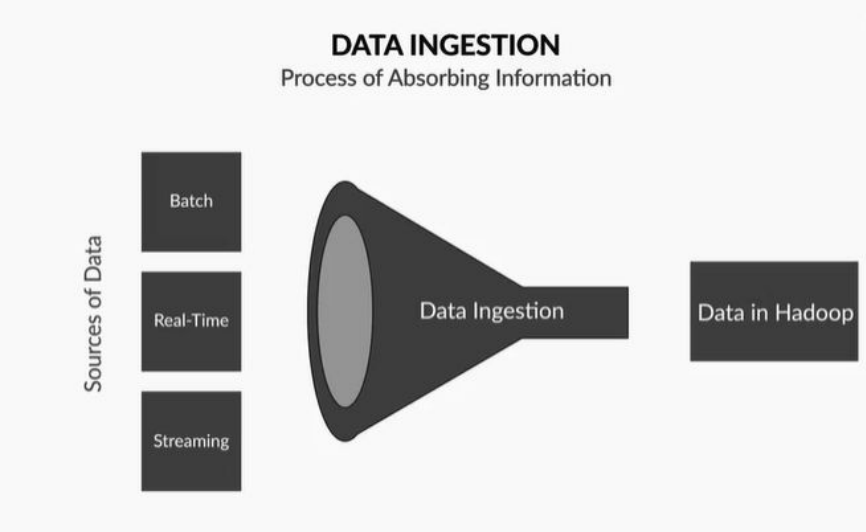

- How is data transferred to the systems in which it can be used in the first place?
- **Data ingestion** is **the process of absorbing data for immediate use or storage.**
- It acts as a bridge between the **source** and the **destination such as the Hadoop Distributed File System (HDFS)**, where it can be used efficiently.
- Data can be of one of the following types:
  - Batch
  - Real-time
  - Streaming

Segment Summary

upGrad

1. Learnt in brief about the process of Data Ingestion

2. Learnt why Data Ingestion is important at the industry level

upGrad

Segment - 04
Challenges in Data Ingestion

Learning Objectives

upGrad

1. Challenges faced in the process of Data Ingestion
2. Massive growth of data in today's era

Challenges in Data Ingestion

upGrad

Challenges

Multiple Data Sources

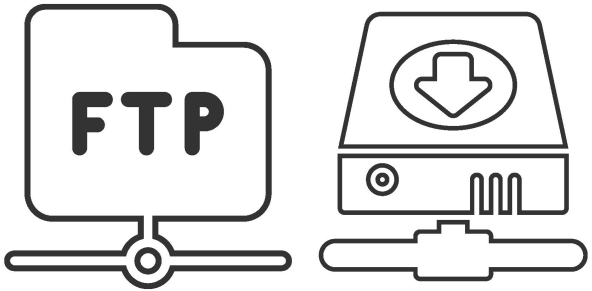

Numerous Data Types and File Formats of Data

Challenges in Data Ingestion

upGrad

Challenges

Processing Time

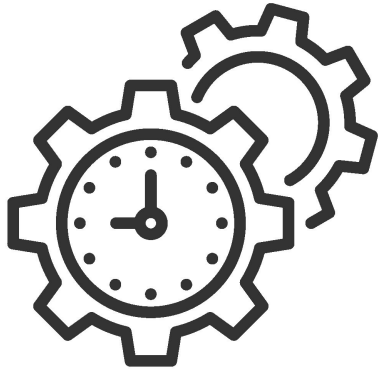

Data Generated at a High Rate and a Huge Scale

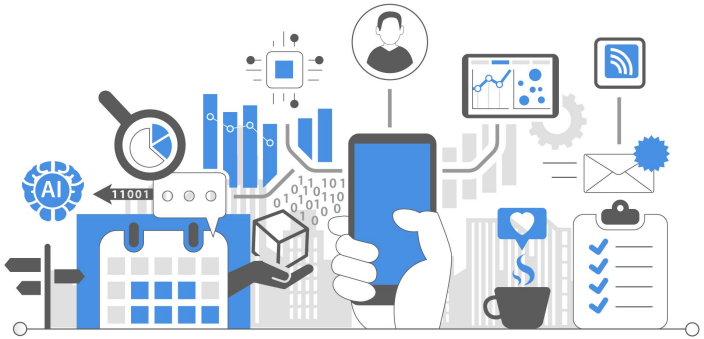

Challenges in Data Ingestion

upGrad

Challenges

Network Performance

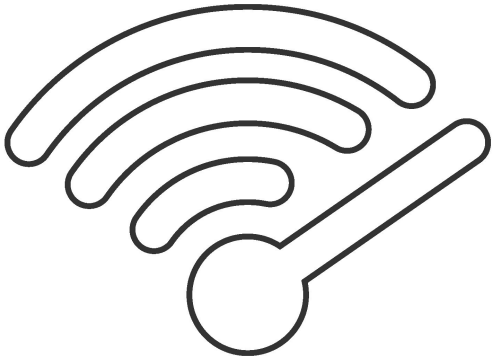

Network Security

Challenges in Data Ingestion

upGrad

Facts about data generated worldwide

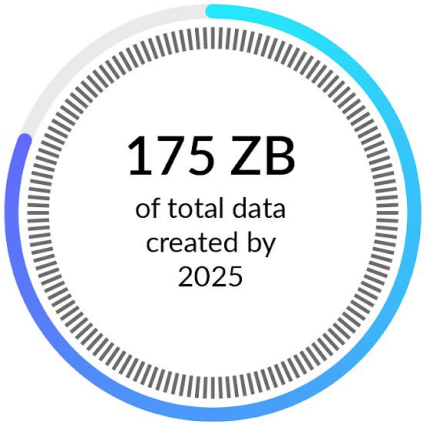

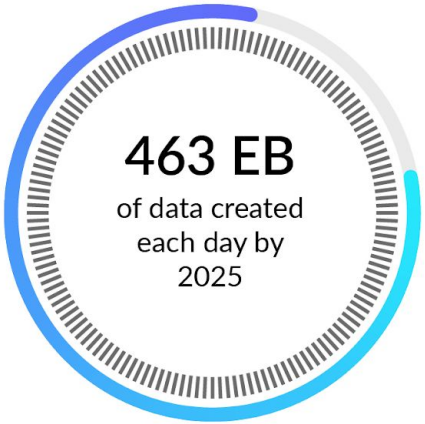

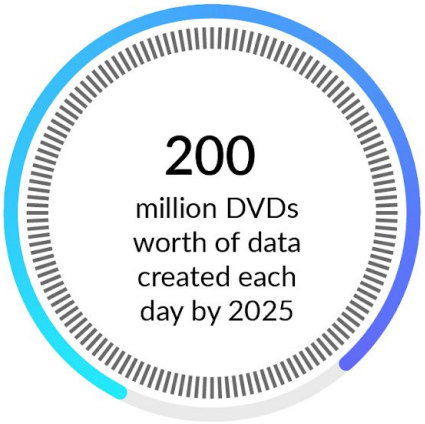

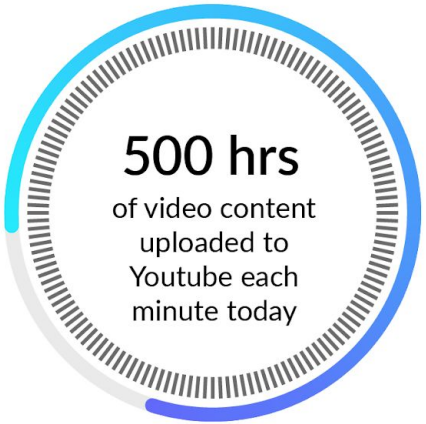

Challenges in Data Ingestion

upGrad

- Today, many tools for data ingestion such as Sqoop, Flume, Kafka and Gobblin have been developed that help in managing data ingestion tasks.
- The following infographic from the ‘Data Age 2025’ white paper shows the predicted data growth until 2025.

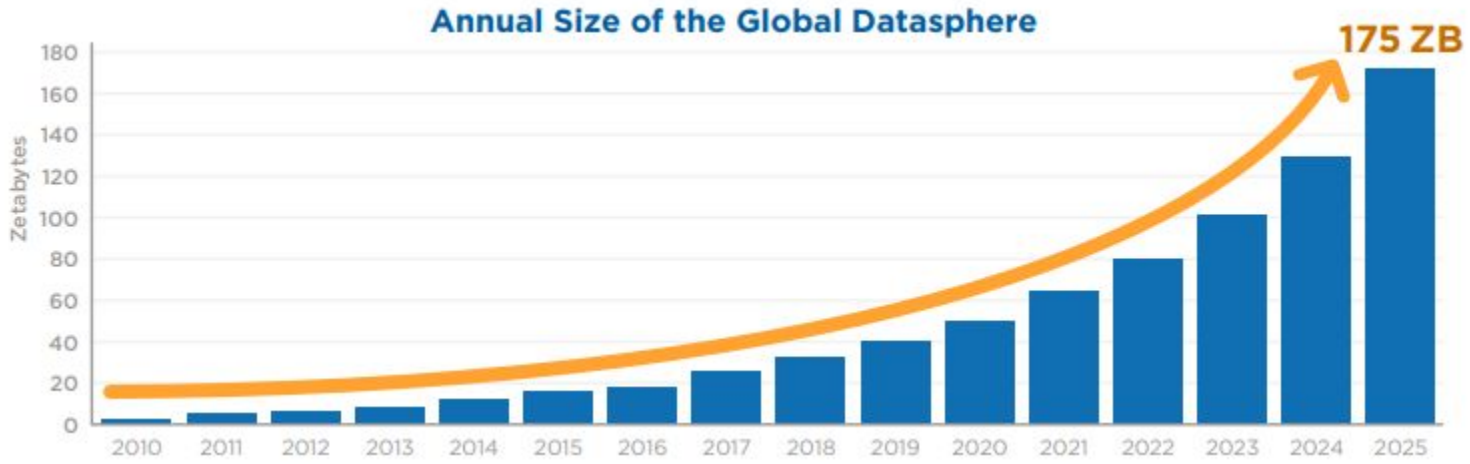

Source: Data Age 2025 sponsored by Seagate with data from the IDC Global DataSphere, November 2018

Segment Summary

upGrad

1. Learnt about the different challenges faced in Data Ingestion
2. Learnt about the massive growth of data in today's era

upGrad

Segment - 05
Key Steps of Data Ingestion

Learning Objectives

upGrad

1. Key steps of Data Ingestion
2. Demonstration of these steps using an example

Key Steps of Data Ingestion

upGrad

Main Steps Followed in Data Ingestion

Let’s understand how data ingestion takes place in a scenario in which you have to design a recommender system for Netflix:

- Which data sources to use?
  - User profile information
  - User browsing history
  - Survey data from emails and forms
- Which data is useless for our system?
  - Data could be duplicate.
  - Data, such as the user name and the user ID, could be useless.
- How does data reach its destination?
  - Many databases and other systems such as HBase, Hive and HDFS

Key Steps of Data Ingestion

upGrad

Broadly, three main steps are being carried out, which are as follows:

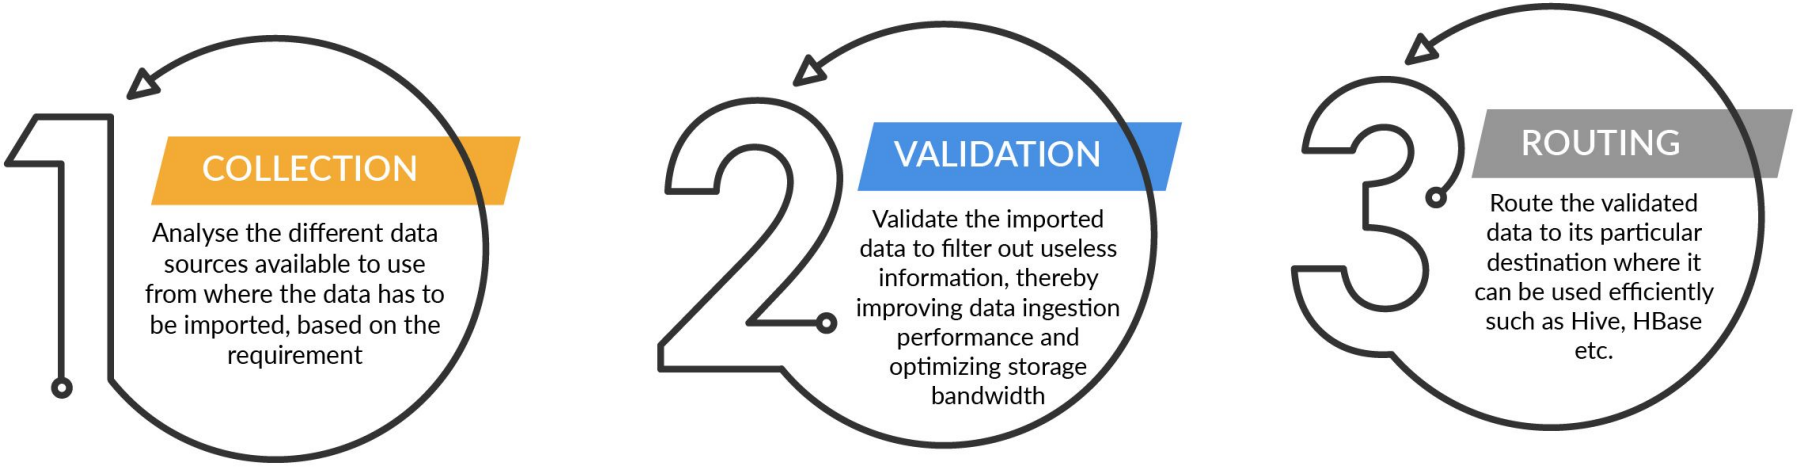

Segment Summary

upGrad

1. Learnt the key steps of Data Ingestion
2. Steps were demonstrated with the help of an example

upGrad

Segment - 06
Tools for Data Ingestion

Learning Objectives

upGrad

1. Brief overview of the different data sources
2. Introduction to some of the tools used for Data Ingestion

Tools for Data Ingestion

upGrad

Common Data Sources at the Industry Level

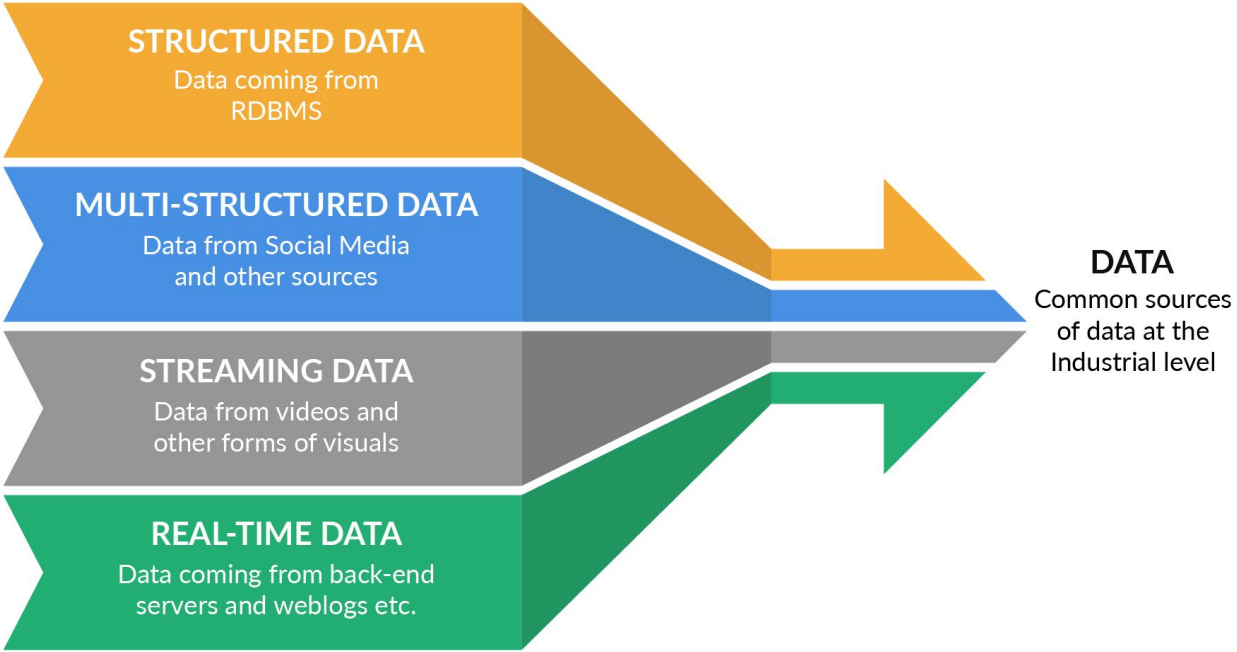

Tools for Data Ingestion

upGrad

File transfer using commands

- ‘distcp’: copy large data sets between two clusters
- ‘put’ and ‘get’: copy files from the local file system to HDFS and vice versa, respectively

```
[root@ip-10-0-0-14 ~]# hadoop fs -put
test.txt /user/root/
```

```
[root@ip-10-0-0-14 ~]# hadoop fs -get
/user/root/test.txt /root/testing
```

Apache Sqoop

- Short for SQL to Hadoop
- Used for importing data from RDBMS to a Big Data Ecosystem (Hive, HBase, etc.) and exporting data back to RDBMS after it is processed.

Tools for Data Ingestion

upGrad

Apache Flume

- Distributed data collection service for collecting, aggregating and transporting large amounts of real-time data from various sources to a centralised place, where it can be processed

Apache Kafka

- Kafka is a fast, scalable distributed system that can handle a high volume of data.
- It enables programmers to pass messages from one point to another.

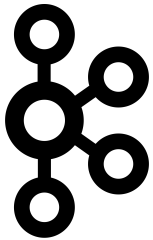

Tools for Data Ingestion

upGrad

Apache Gobblin

- Gobblin is an open source data ingestion framework for extracting, transforming and loading a large volume of data from different data sources. It supports both streaming and batch data ecosystems.

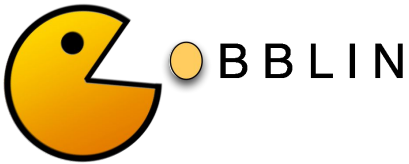

Segment Summary

upGrad

1. Introduced the various sources of data
2. Introduced some tools used for Data Ingestion

upGrad

Segment - 07
Types of Data and File Formats

Learning Objectives

upGrad

1. Different types of data handled in Data Ingestion
2. Different types of file formats handled in Data Ingestion

Types of Data and File Formats

upGrad

Types of Data

- **Structured data**:
  - Organised data is generally stored in databases
  - Can be easily stored, entered, queried and analysed efficiently using SQL
  - Can be easily read by machines
  - Examples: Financial data, user identification data, etc.

Types of Data and File Formats

upGrad

Types of Data

- **Unstructured data**:
  - Opposite of structured data: Cannot be easily stored and organised in databases
  - NoSQL databases can be used for this type of data
  - Approximately 80% of data being created today is unstructured in nature.
  - Examples: Images, audio, video, chat messages, etc.

Types of Data and File Formats

upGrad

Types of Data

- **Semi-structured data**:
  - No predefined scheme unlike structured data
  - May have an internal structure and markings to identify separate data elements, but its schema does not constrain the data as in an RDBMS such as SQL tables.
  - Example: XML and JSON files

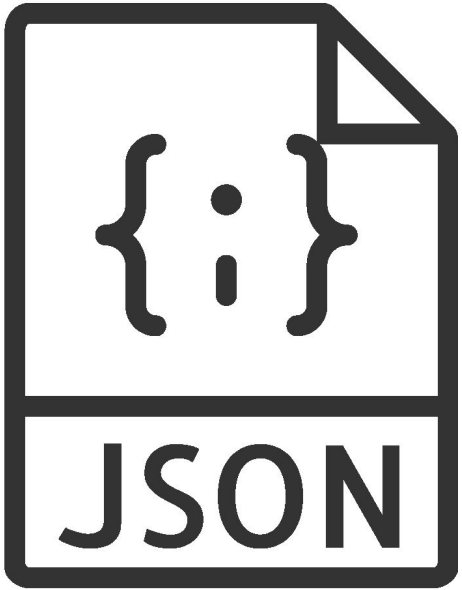

upGrad

File Formats

Factors of data ingestion that vary depending on the file format:

- Processing power
- Network bandwidth
- Available storage

The file formats that are commonly dealt with in Data Ingestion are as follows:

- **Text/CSV**:
  - CSV: Comma-separated values
  - The most commonly used file format for exchanging large datasets between Hadoop and external systems
  - Limited support for schema evolution
  - Does not support block compression

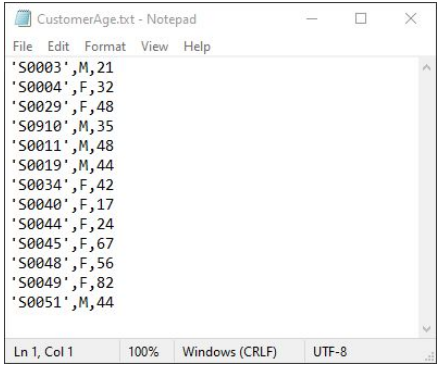

Figure: Sample text/CSV file

Types of Data and File Formats

upGrad

- **XML and JSON**:
  - **XML**: Extensible Markup Language
    JSON: JavaScript Object Notation
  - **XML**: It defines a set of rules, using which documents can be encoded in a machine- and human-readable format.
  - **JSON**: Open-standard file format consisting key-value pairs
  - **Essentially text files**: Do not support block compression and are not compact
  - Splitting is hard and cannot be easily processed parallely because no in-built InputFormat is present for either of the two formats in Hadoop.

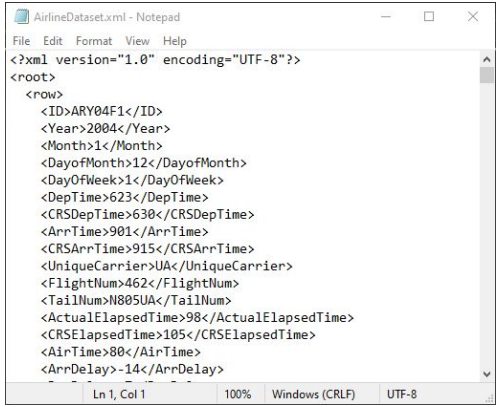

Figure: AirLine data set used as an example for showcasing XML and JSON files

Types of Data and File Formats

upGrad

- **Sequence file**:
  - Store data as binary key-value pairs in a binary format
  - More compact than text files
  - Supports block compression and can be easily processed parallelly
- **Avro**:
  - A language-neutral data serialisation system developed with Apache's Hadoop project
  - Can be easily read after creation, even in a language different from the one used to write the file.
  - Compact: A type of binary file
  - Self-describing, compressible and splittable and, hence, suitable for MapReduce Jobs
  - Supports scheme evolution

Figure: First, Sequence File Format
Second, Avro Logo

Segment Summary

upGrad

1. Discussed the different types of data handled in Data Ingestion
2. Discussed various file formats handled in Data Ingestion

Session Summary

upGrad

1	Discussed about data ingestion and why it is needed
2	Discussed the various challenges faced in Data Ingestion
3	Learnt about the three key steps of Data Ingestion, i.e., data collection, validation and routing
4	Looked at various tools used for Data Ingestion such as Apache Sqoop and Apache Flume
5	Discussed the types of data and file formats that are usually dealt with in Data Ingestion

upGrad

Thank You