

# Hive and Querying

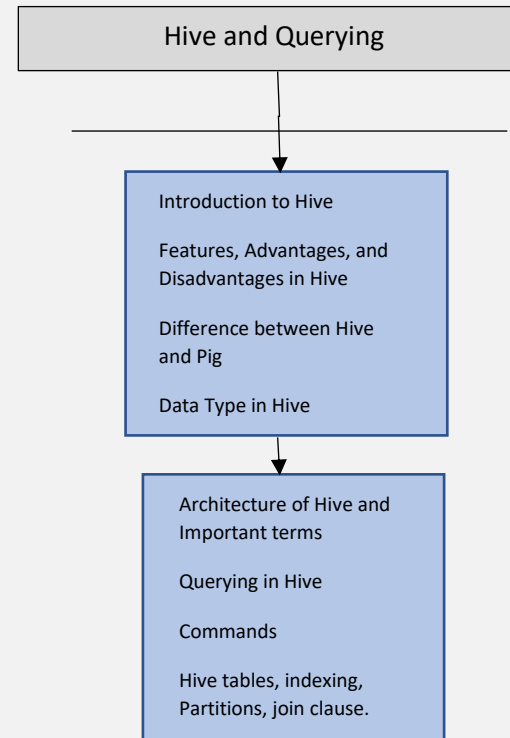
Hive is a data warehouse system which is used to analyse structured data. It runs SQL like queries called HQL (Hive query language) which gets internally converted to MapReduce jobs.

As a part of Hive and Querying, you covered:

- Introduction to Hive, Features, limitation, and advantages
- Difference between Hive and Pig, Hive Architecture, and Important terms
- Hive tables, Partition, Indexing and join clause in Hive
- Data type in Hive and Querying of Data in Hive

## Common Interview Questions:

1. What applications are supported by Hive?
2. What are the different tables available in Hive?
3. What is the difference between external and managed tables?
4. Where does the data of a Hive table get stored?
5. Can Hive be used in OLTP systems?
6. Can a table name be changed in Hive?
7. Where is Hive table data stored?
8. Can we use the LOAD or INSERT command to view?
9. What is a Hive Metastore?



# Hive and Querying

## Apache Hive:

**Hive** is a data warehouse system which is used to analyse structured data. It is built on the top of Hadoop. It was developed by Facebook.

Hive provides the functionality of reading, writing, and managing large datasets residing in distributed storage. It runs SQL like queries called HQL (Hive query language) which gets internally converted to MapReduce jobs.

## Features of Hive:

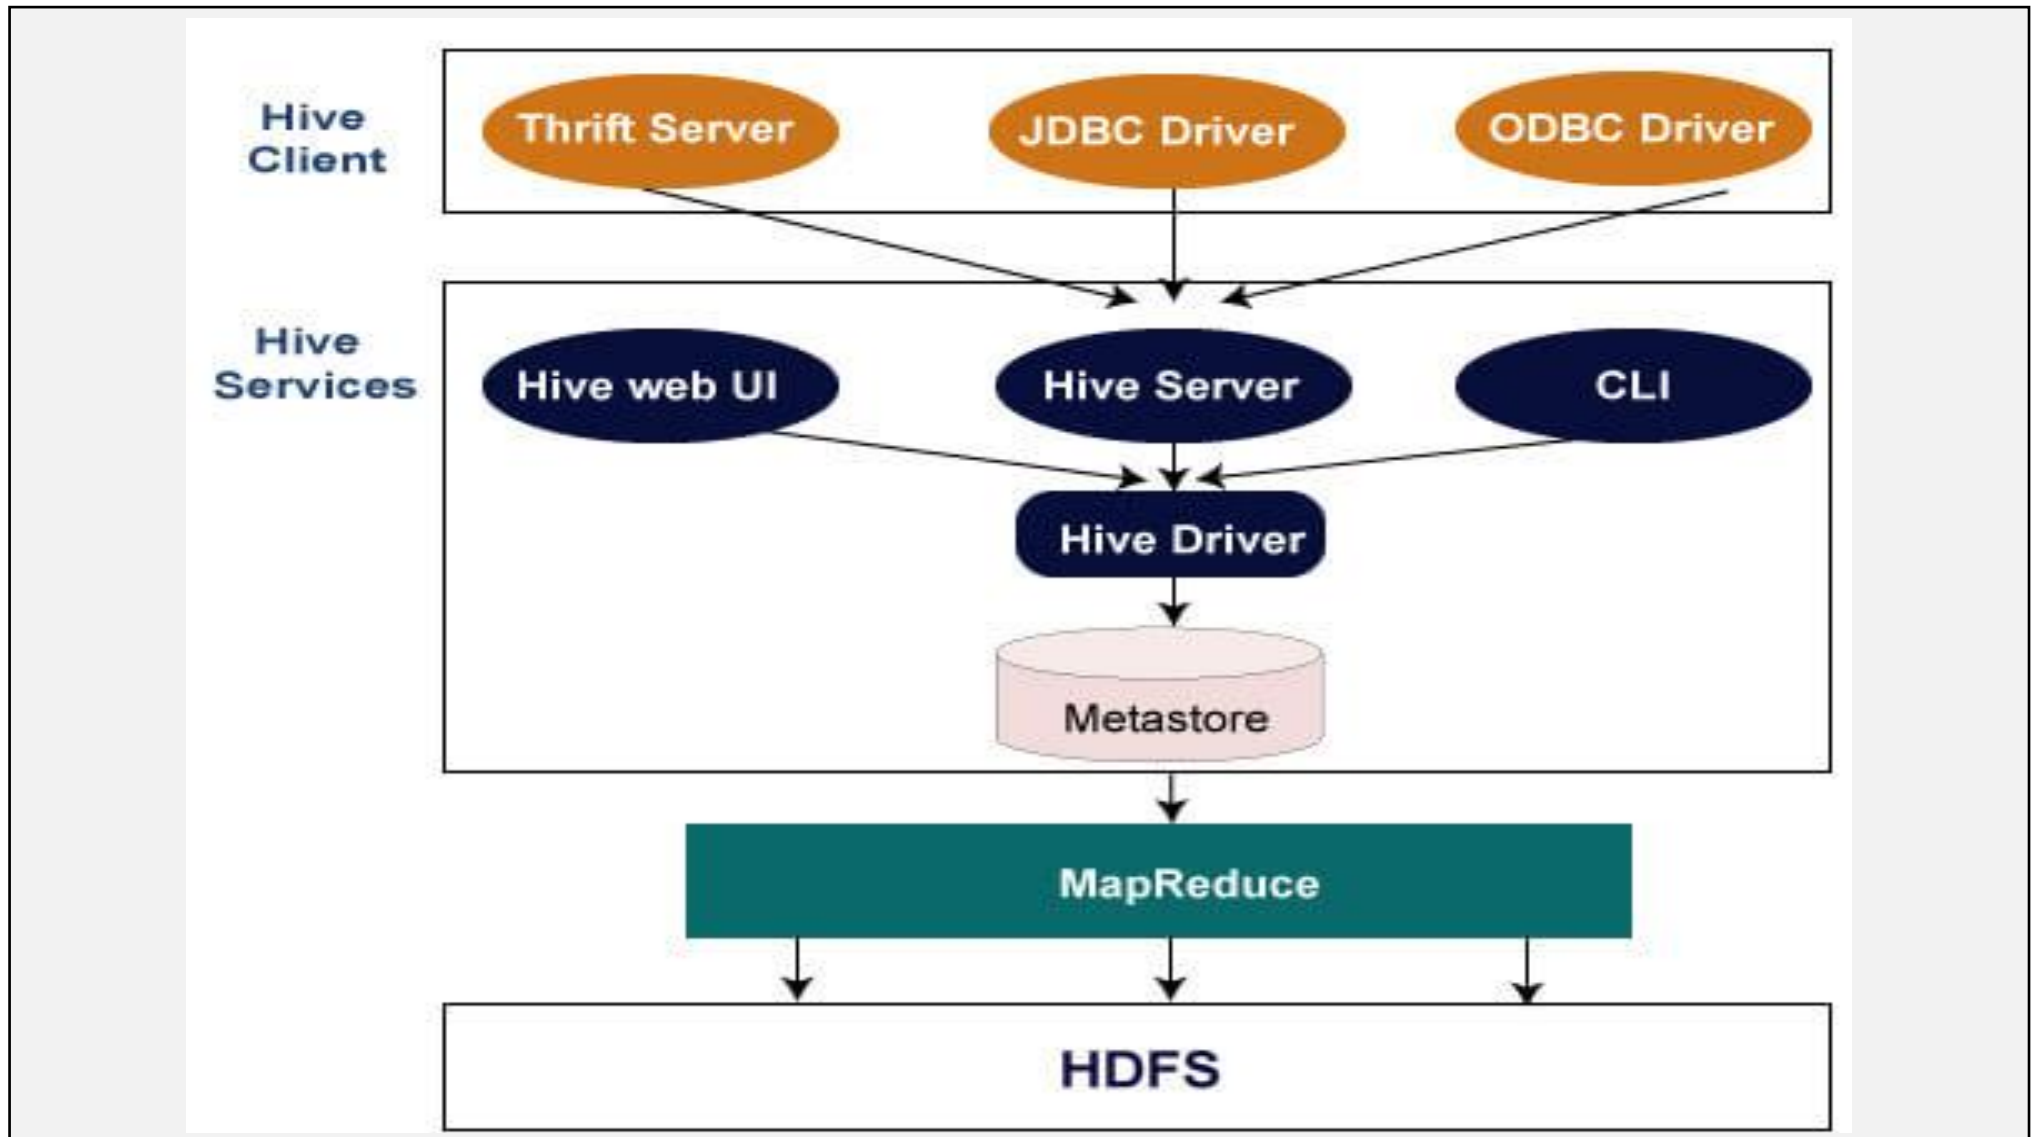


## Apache Hive VS Pig:

Based on	Hive	Pig
Language	Hive uses a declarative language called HiveQL	With Pig Latin, a procedural data flow language is used.
Type of data	Hive works on structured data. Does not work on other types of data	Pig works on structured, semi-structured and unstructured data
Operates on	Works on the server-side of the cluster	Works on the client-side of the cluster
Avro File Format	Hive does not support Avro	Pig supports Avro
Schema	Hive supports schema	Creating schema is not required to store data in Pig
Data Processing	Hive is used for batch processing	Pig is a high-level data-flow language
JDBC/ ODBC	Supported, but limited	Unsupported
Loading Speed	Hive takes time to load but executes quickly	Pig loads data quickly

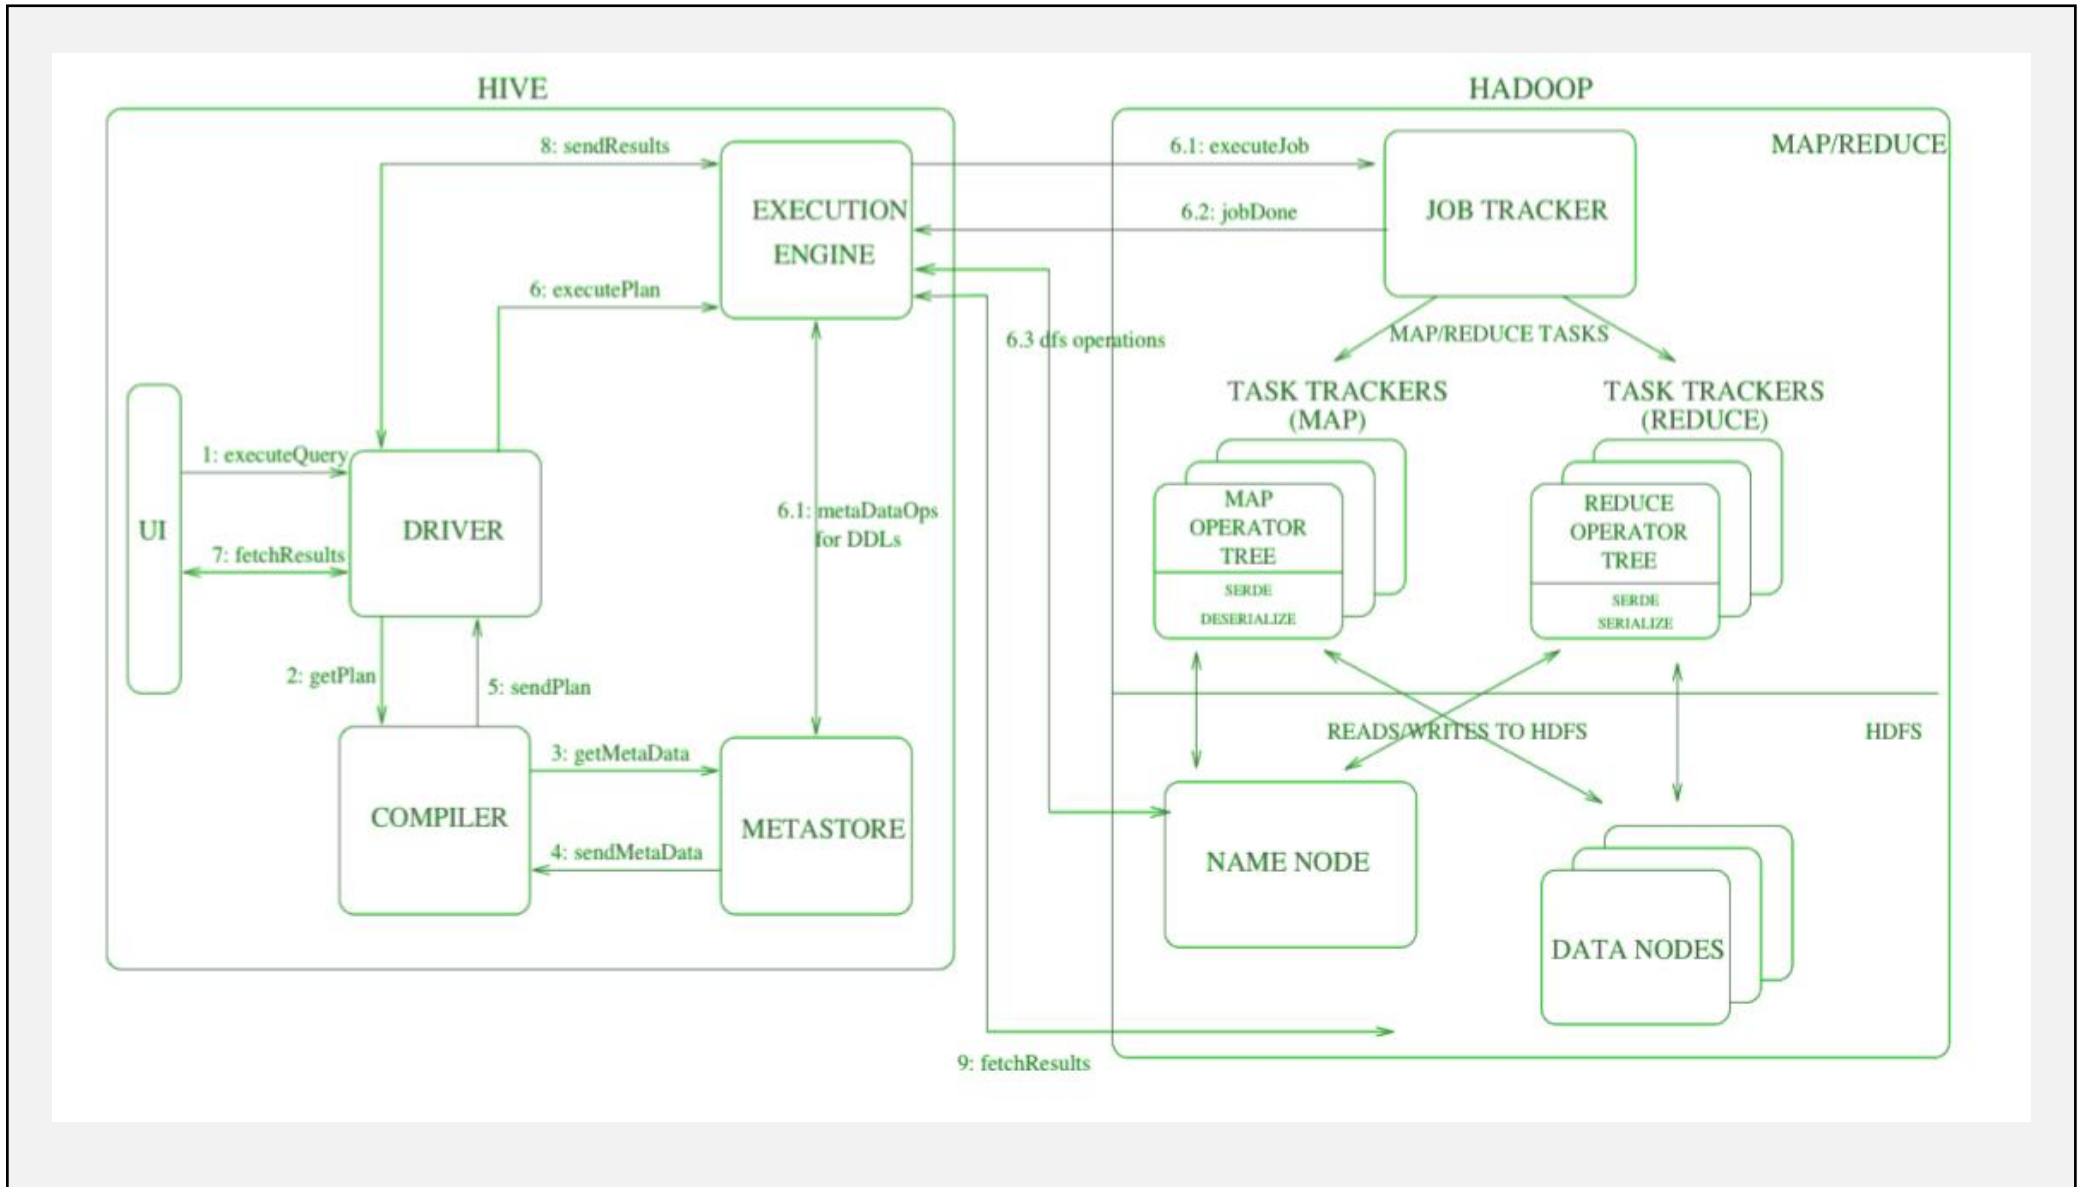
## Hive and Querying

### Hive Architecture:



## Hive and Querying

### Architecture of Hive that is built on the top of Hadoop:



## Hive and Querying

### Important Terminologies in Hive Architecture:

#### Hive Client:

**Thrift Server** - It is a cross-language service provider platform that serves the request from all those programming languages that supports Thrift.

**JDBC Driver** - It is used to establish a connection between hive and Java applications. The JDBC Driver is present in the class `org.apache.hadoop.hive.jdbc.HiveDriver`.

**ODBC Driver** - It allows the applications that support the ODBC protocol to connect to Hive.

#### Hive Services:

**Hive CLI** - The Hive CLI is a shell where we can execute Hive queries and commands.

**Hive Web User Interface** - The Hive Web UI is just an alternative of Hive CLI. It provides a web-based GUI for executing Hive queries and commands.

**Hive MetaStore** - It is a central repository that stores all the structure information of various tables and partitions in the warehouse. It also includes metadata of column and its type information, the serializers and deserializers which is used to read and write data and the corresponding HDFS files where the data is stored.

**Hive Server** - It is referred to as Apache Thrift Server. It accepts the request from different clients and provides it to Hive Driver.

**Hive Driver** - It receives queries from different sources like web UI, CLI, Thrift, and JDBC/ODBC driver. It transfers the queries to the compiler.

**Hive Compiler** - The purpose of the compiler is to parse the query and perform semantic analysis on the different query blocks and expressions. It converts HiveQL statements into MapReduce jobs.

**Hive Execution Engine** - Optimizer generates the logical plan in the form of DAG of map-reduce tasks and HDFS tasks. In the end, the execution engine executes the incoming tasks in the order of their dependencies.

## Hive and Querying

### Advantages and Disadvantages of using Hive:

#### Advantages of using Hive:

- Hive is built on Hadoop, which supports and handles all the capabilities of Hadoop provides like reliable, highly available, node failure, commodity hardware
- Database developers need not learn Java programming for writing map-reduce programs for retrieving data from the Hadoop system.
- Data stored in HDFS so you will have features of scalability, redundancy over hive SQL language
- Querying data using hive is simple and easy to use

#### Disadvantages of using Hive:

- Hive is not for OLAP processing, only supports OLTP processing
- Subqueries are not supported.
- It has a high latency.
- Hive tables don't support delete or update operations.
- ACID transaction limitation on Hive
- indexing limitations,
- schema operation

## Hive and Querying

### Hive Data Type:

#### Integer Types:

Type	Size	Range
<b>TINTINT</b>	1-byte signed integer	-128 to 127
<b>SMALLINT</b>	2-byte signed integer	32,768 to 32,767
<b>INT</b>	4-byte signed integer	2,147,483,648 to 2,147,483,647
<b>BIGINT</b>	8-byte signed integer	-9,223,372,036,854,775,808 to 9,223,372,036,854,775,807

#### Complex Types:

Type	Range
<b>Struct</b>	struct('James','Roy')
<b>Map</b>	map('first','James','last','Roy')
<b>Array</b>	array('James','Roy')

#### String Types:

##### **STRING**

The string is a sequence of characters. Its values can be enclosed within single quotes (') or double quotes (").

##### **Varchar**

The varchar is a variable length type whose range lies between 1 and 65535.

##### **CHAR**

The char is a fixed-length type whose maximum length is fixed at 255.

#### Date/Time Types:

##### **TIMESTAMP:**

Timestamp format "YYYY-MM-DD HH:MM:SS.fffffffff" (9 decimal place precision)

**DATES:** The Date value is used to specify a particular year, month and day, in the form YYYY--MM--DD.

#### Decimal Types:

Type	Range
<b>Float</b>	Single precision floating point number
<b>Double</b>	Double precision floating point number



## Hive and Querying

### Types of Hive Table:



## Hive Internal Tables VS External Tables

### Internal Table

1. Hive moves table data to warehouse directory
2. Dropping deletes table data and metadata
3. Support TRUNCATE
4. Support ACID transaction
5. Query Result Caching works



### External Table

1. Hive does not moves table data to warehouse directory
2. Dropping deletes table's metadata
3. No TRUNCATE support
4. No ACID transaction support
5. Query Result Caching does not works



# Hive and Querying

## Data Querying with Hive:

### Data Querying with Hive:

Hive provides SQL type querying language for the ETL purpose on top of Hadoop file system. Hive Query language (HiveQL) provides SQL type environment in Hive to work with tables, databases, queries. We can have a different type of Clauses associated with Hive to perform different type data manipulations and querying.

### Hive DML Commands:



### Hive DDL Commands:



# Hive and Querying

## Order By:

The ORDER BY syntax in HiveQL uses the “SELECT” statement to help sort data. The query will only pick the column name mentioned in the Order by clause, and display the matching column values in ascending or descending order.

### Example:

```
hive> SELECT * FROM employees_guru ORDER BY Department;
Query ID = hduser_20151117170229_6e8af657-126b-470f-8b88
Total jobs = 1
Launching Job 1 out of 1
Number of reduce tasks determined by job: 1
In order to change the average size of the reduce tasks:
  set hive.exec.reducers.bytes-per-row=<number>
In order to limit the maximum number of reducers:
  set hive.exec.reducers.max=<number>
In order to set a constant number of reducers:
  set mapred.reduce.tasks=<number>
Starting Job = job_201511171701_0001, Tracking URL = http://localhost:5003/jobdetails.jsp?jobid=job_201511171701_0001
Kill Command = /usr/local/hadoop-1.2.1/libexec/./bin/hadoop job -kill job_201511171701_0001
Hadoop job information for Stage-1: number of mappers: 1; number of reducers: 1
2015-11-17 17:02:39,820 Stage-1 map = 0%, reduce = 0%
2015-11-17 17:02:44,008 Stage-1 map = 100%, reduce = 0%
2015-11-17 17:02:52,078 Stage-1 map = 100%, reduce = 33%
2015-11-17 17:02:53,085 Stage-1 map = 100%, reduce = 100%
MapReduce Total cumulative CPU time: 2 seconds 40 msec
Ended Job = job_201511171701_0001
MapReduce Jobs Launched:
  Stage-Stage-1: Map: 1 Reduce: 1 Cumulative CPU: 2.04 sec HDFS Read: 7.0 MB HDFS Write: 380 SUCCESS
Total MapReduce CPU Time Spent: 2 seconds 40 msec
OK
103    Animesh  26      Bangalore  25000.0  ADMIN
104    Anirudh  27      bangalore  27000.0  ADMIN
106    Ramesh   30      Goa        24000.0  FINANCE
101    Rajesh   27      Bangalore  20000.0  HR
102    Rajiv    28      Delhi      30000.0  HR
107    Sravanthi 32      Chennai    20000.0  IT
108    Sravan   31      Mumbai     60000.0  IT
109    Suresh   32      Kolkata    20000.0  IT
110    Ravi     28      bangalore  20000.0  IT
105    Santosh  33      bangalore  25000.0  PR
111    Syam     33      bangalore  25000.0  PR
Time taken: 25.224 seconds, Fetched: 11 row(s)
```

order by query on "employees\_guru" table

order by query output

## Group by:

The query will only look at the columns under the name defined as “group by” clause, and it will show the results by grouping the specific and matching column values.

### Example:

```
hive> SELECT Department, count(*) FROM employees_guru GROUP BY Department;
Query ID = hduser_20151105155307_1574cd2b-866e-437a-8d14-637e02b7e515
Total jobs = 1
Launching Job 1 out of 1
Number of reduce tasks determined by job: 1
In order to change the average size of the reduce tasks:
  set hive.exec.reducers.bytes-per-row=<number>
In order to limit the maximum number of reducers:
  set hive.exec.reducers.max=<number>
In order to set a constant number of reducers:
  set mapred.reduce.tasks=<number>
Starting Job = job_201511051442_0005, Tracking URL = http://localhost:5003/jobdetails.jsp?jobid=job_201511051442_0005
Kill Command = /usr/local/hadoop-1.2.1/libexec/./bin/hadoop job -kill job_201511051442_0005
Hadoop job information for Stage-1: number of mappers: 1; number of reducers: 1
2015-11-05 15:53:19,229 Stage-1 map = 0%, reduce = 0%
2015-11-05 15:53:21,235 Stage-1 map = 100%, reduce = 0%, Cumulative CPU 1.00 sec
2015-11-05 15:53:28,277 Stage-1 map = 100%, reduce = 0%, Cumulative CPU 1.00 sec
2015-11-05 15:53:29,281 Stage-1 map = 100%, reduce = 0%, Cumulative CPU 1.00 sec
MapReduce Total cumulative CPU time: 2 seconds 130 msec
Ended Job = job_201511051442_0005
MapReduce Jobs Launched:
  Stage-Stage-1: Map: 1 Reduce: 1 Cumulative CPU: 2.013 sec HDFS Read: 7.0 MB HDFS Write: 380 SUCCESS
Total MapReduce CPU Time Spent: 2 seconds 130 msec
OK
ADMIN  2
FINANCE 1
HR      2
IT      4
PR      2
Time taken: 23.057 seconds, Fetched: 5 row(s)
```

Groupby query on "employees\_guru"

Group by query output

# Hive and Querying

## Sort By:

When a Hive query comes with a “Sort by” clause, it goes through the columns under the name defined by the query. Once executed, the query explores columns of Hive tables to sort the output.

## Example:

```
hive> Select * from employees guru SORT BY id DESC;
Query ID = hdus_1_20151105164027_55bf2f64-6f5b-4764-b94e-e03f
Total jobs = 1
Launching Job 1 out of 1
Number of reduce tasks not specified, assuming 1 from input
In order to change the average size of the reduce tasks: (in bytes)
  set hive.exec.reducers.bytesinreduce=<number>
In order to limit the maximum number of reducers:
  set hive.exec.reducers.max=<number>
In order to set a constant number of reducers:
  set mapred.reduce.tasks=<number>
Starting Job = job_201511051442_0007, Tracking URL = http://10.10.10.10:8080/
Kill Command = /usr/local/hadoop-1.2.1/libexec/./bin/hadoop
Hadoop job information for Stage-1: number of mappers: 1; number of reducers: 1;
2015-11-05 16:40:34,093 Stage-1 map = 0% reduce = 0%
2015-11-05 16:40:36,098 Stage-1 map = 0% reduce = 0%, Cumulative CPU: 1.62 sec
2015-11-05 16:40:44,145 Stage-1 map = 100% reduce = 100%, Cumulative CPU: 1.62 sec
MapReduce Total cumulative CPU time: 1 seconds 620 msec
Ended Job = job_201511051442_0007
MapReduce Jobs Launched:
Stage-Stage-1: Map: 1 Reduce: 1 Cumulative CPU: 1.62 sec Total MapReduce CPU Time Spent: 1 seconds 620 msec
OK
111      Syam      33      bangalore      25000.0 PR
110      Ravi       28      bangalore      20000.0 IT
109      Suresh     32      Kolkata 20000.0 IT
108      Sravan     31      Mumbai  60000.0 IT
107      Sravanthi   32      Chennai 20000.0 IT
106      Ramesh     30      Goa       24000.0 FINANCE
105      Santosh    33      bangalore 25000.0 PR
104      Anirudh    27      bangalore 27000.0 ADMIN
103      Animesh    26      Bangalore 25000.0 ADMIN
```

sort by query

sort by output on "employees\_guru" table

## Cluster By:

Hive queries with a CLUSTER BY clause or command are typically deployed in queries to perform the functions of both DISTRIBUTE BY and SORT BY together.

## Example:

```
hive> Select Id,Name from employees guru CLUSTER BY Id;
Query ID = hdus_1_20151105165000_72cedc06-a797-48b1-a120-
Total jobs = 1
Launching Job 1 out of 1
Number of reduce tasks not specified, assuming 1 from input
In order to change the average size of the reduce tasks: (in bytes)
  set hive.exec.reducers.bytesinreduce=<number>
In order to limit the maximum number of reducers:
  set hive.exec.reducers.max=<number>
In order to set a constant number of reducers:
  set mapred.reduce.tasks=<number>
Starting Job = job_201511051442_0009, Tracking URL = http://10.10.10.10:8080/
Kill Command = /usr/local/hadoop-1.2.1/libexec/./bin/hadoop
Hadoop job information for Stage-1: number of mappers: 1; number of reducers: 1;
2015-11-05 16:50:08,541 Stage-1 map = 0% reduce = 0%
2015-11-05 16:50:10,546 Stage-1 map = 0% reduce = 0%
2015-11-05 16:50:17,563 Stage-1 map = 100% reduce = 100%, Cumulative CPU: 1.62 sec
MapReduce Total cumulative CPU time: 1 seconds 600 msec
Ended Job = job_201511051442_0009
MapReduce Jobs Launched:
Stage-Stage-1: Map: 1 Reduce: 1 Cumulative CPU: 1.62 sec Total MapReduce CPU Time Spent: 1 seconds 600 msec
OK
101      Rajesh
102      Rajiv
103      Animesh
104      Anirudh
105      Santosh
106      Ramesh
107      Sravanthi
108      Sravan
109      Suresh
```

cluster by query

cluster by query output

# Hive and Querying

## Distribute By:

The DISTRIBUTE BY instruction determines how the output is divided among reducers in a MapReduce job.

## Example:

```
hive> Select Id,Name from employees guru DISTRIBUTE BY Id;
Query ID = hdfs_20151105165433_65088a93-2ec2-4878-985d-cb
Total jobs = 1
Launching Job 1 out of 1
Number of reduce tasks not assigned yet: 1 from input
In order to change the average size of the reduce (in bytes)
  set hive.exec.reducers.bytesinreduce=<number>
In order to limit the maximum number of reducers
  set hive.exec.reducers.max=<number>
In order to set a constant number of reducers
  set mapred.reduce.tasks=<number>
Starting Job = job_201511051442_0010, Tracking URL = http://
Kill Command = /usr/local/hadoop-1.2.1/libexec/./bin/hadoop
Hadoop job information for Stage-1: number of mappers: 1; n
2015-11-05 16:54:42,456 Stage-1 map = 0%, reduce = 0%
2015-11-05 16:54:43,456 Stage-1 map = 100%, reduce = 0%, C
2015-11-05 16:54:51,456 Stage-1 map = 100%, reduce = 100%, C
MapReduce Total cumulative CPU time: 1 seconds 790 msec
Ended Job = job_201511051442_0010
MapReduce Jobs Launched = 1
Stage-Stage-1: Map: 1 Reduce: 1 Cumulative CPU: 1.79 sec
Total MapReduce Time Spent: 1 seconds 790 msec
OK
101 Rajesh
102 Rajiv
103 Animesh
104 Anirudh
105 Santosh
106 Ramesh
107 Sravanthi
108 Sravan
109 Suresh
110 Ravi
111
```



## Some effective ways to optimize Hive queries:

**Partition keys** present an opportunity to target a subset of the table data, rather than scan data you don't need for your operations.

**Bucketing** allows you to target a subset of data. It improves join performance specifically by scanning fewer data.

**Compression** minimizes the amount of data that traverses each of those steps and decreases the time spent moving through the query states.

**Eliminating** small file operations from your query promotes a healthy Hive ecosystem. Each file is tracked by the Hive metastore and stored in HDFS, which are each performance-optimized to handle larger files over many smaller files.

# Hive and Querying

## Indexes in Hive:

### Indexes in Hive:

Indexes are a pointer or reference to a record in a table as in relational databases. In Hive, the index table is different than the main table. Indexes facilitate in making query execution or search operation faster.

There are two types of indexes in Hive:

**Bitmap Indexing:** This is used with columns having a few distinct values. It is known to store both the indexed column's value and the list of rows as a bitmap.

**Compact Indexing:** This type of indexing is known to store the column value and storage blockid.

## Hive Partitions:

### Hive Partitions:

Apache Hive organizes tables into partitions. Partitioning is a way of dividing a table into related parts based on the values of particular columns like date, city, and department.

There are two types of Partitioning in Hive:

**Static Partitioning:** Insert input data files individually into a partition table is Static Partition.

**Dynamic Partitioning:** Single insert to partition table is known as a dynamic partition.

## JOIN Clause in Hive:

### JOIN Clause in Hive:

JOIN is a clause that is used for combining specific fields from two tables by using values common to each one. It is used to combine records from two or more tables in the database.

### SYNTAX:

```
join_table:

table_reference JOIN table_factor [join_condition]
| table_reference {LEFT|RIGHT|FULL} [OUTER] JOIN table_reference
join_condition
| table_reference LEFT SEMI JOIN table_reference join_condition
| table_reference CROSS JOIN table_reference [join_condition]
```