

NoSQL Databases and Apache HBase

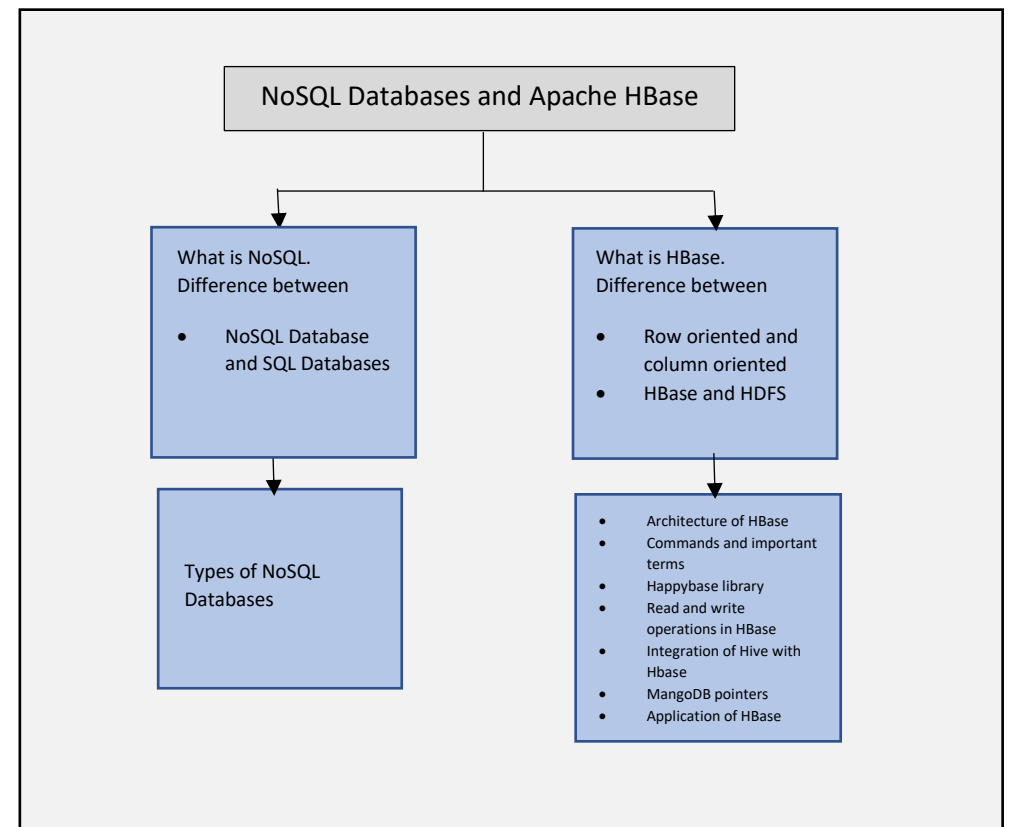
NoSQL Database is a non-relational Data Management System, that does not require a fixed schema. HBase is an essential part of our Hadoop ecosystem. It can store massive amounts of data from terabytes to petabytes.

As a part of NoSQL Databases and Apache HBase, you covered:

- Introduction to NoSQL
- Types of NoSQL
- Introduction to HBase
- HBase Architecture and commands

Common Interview Questions:

1. Mention the difference between HBase and Relational Database?
2. Explain why to use HBase?
3. HBase support syntax structure like SQL yes or No?
4. Mention what are the key components of HBase?
5. Explain what is the row key?
6. What are the pros and cons of a graph database under NoSQL databases?
7. List the different kinds of NoSQL data stores?
8. List some of the features of NoSQL?
9. Explain NoSQL Databases



NoSQL Databases and Apache HBase


SQL and NoSQL Databases:

NoSQL Databases:


NoSQL Database is a non-relational Data Management System, that does not require a fixed schema. It avoids joins, and is easy to scale. The major purpose of using a NoSQL database is for distributed data stores with humongous data storage needs. NoSQL is used for Big data and real-time web apps.

Types of NoSQL Databases:


Document




Graph



Key-Value



Wide-column



Based on	SQL	NoSQL
Data Storage Model	Tables with fixed rows and columns	Document: JSON documents, Key-value: key-value pairs, Wide-column: tables with rows and dynamic columns, Graph: nodes and edges
Schemas	Rigid	Flexible
Scaling	Vertical (scale-up with a larger server)	Vertical (scale-up with a larger server)
Vertical	Supported	Most do not support multi-record ACID transactions
Joins	Typically required	Typically, not required
Examples	Oracle, MySQL, Microsoft SQL Server, and PostgreSQL	Document: MongoDB and CouchDB, Key-value: Redis and DynamoDB, Wide-column: Cassandra and HBase, Graph: Neo4j and Amazon Neptune

NoSQL Databases and Apache HBase

HBase:

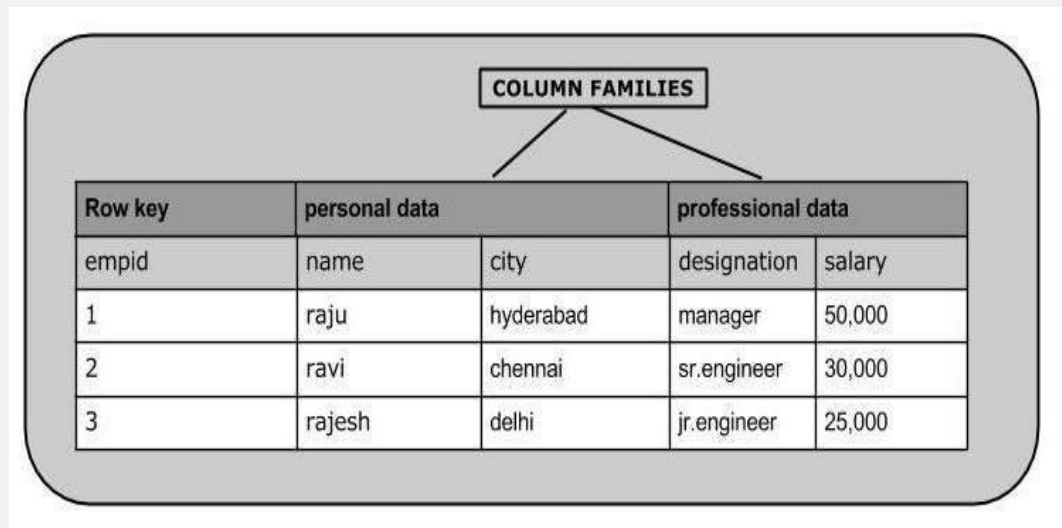
HBase is an open source, distributed database developed by Apache software foundation written in Java. HBase is an essential part of our Hadoop ecosystem. HBase runs on top of HDFS (Hadoop Distributed File System). It can store massive amounts of data from terabytes to petabytes. It is column oriented and horizontally scalable.

HDFS and HBase:

HDFS	HBase
It provides high latency batch processing; no concept of batch processing.	It provides low latency access to single rows from billions of records (Random access).
It provides only sequential access of data.	HBase internally uses Hash tables and provides random access, and it stores the data in indexed HDFS files for faster lookups.

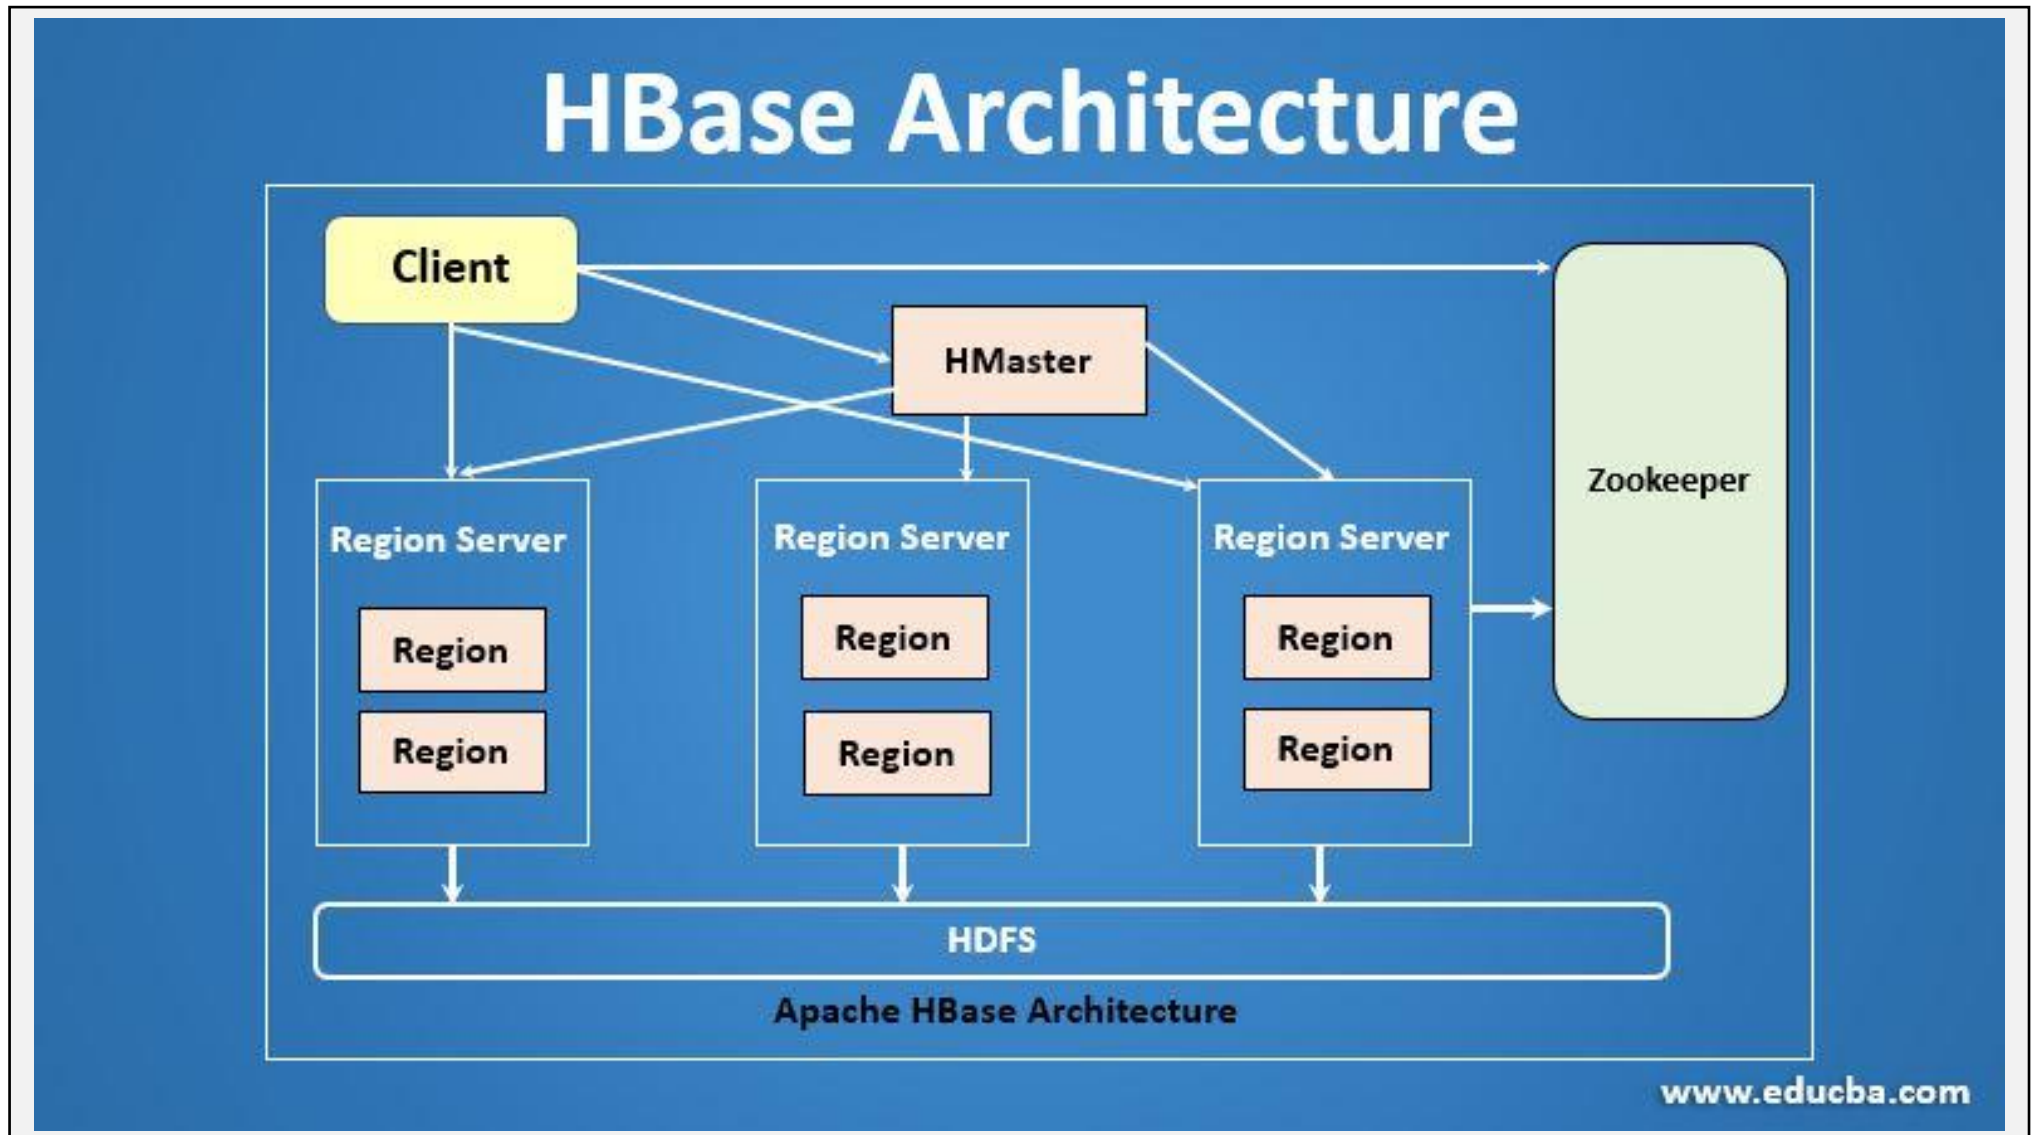
Row-oriented and Column-Oriented Databases:

Row-Oriented Database	Column-Oriented Database
It is suitable for Online Transaction Process (OLTP).	It is suitable for Online Analytical Processing (OLAP).
Such databases are designed for small number of rows and columns.	Column-oriented databases are designed for huge tables.



NoSQL Databases and Apache HBase

HBase Architecture:



NoSQL Databases and Apache HBase

Important Terminologies:

MasterServer: Assigns regions to the region servers and takes the help of Apache ZooKeeper for this task. Handles load balancing of the regions across region servers. It unloads the busy servers and shifts the regions to less occupied servers. Is responsible for schema changes and other metadata operations such as creation of tables and column families.

Region: Regions are nothing but tables that are split up and spread across the region servers.

Zookeeper: Zookeeper is an open-source project that provides services like maintaining configuration information, naming, providing distributed synchronization, etc. Zookeeper has ephemeral nodes representing different region servers. Master servers use these nodes to discover available servers.

General Commands in HBase:

Command	Summary
Status hbase(main):009:0> status	This command returns the status of the system including the details of the servers running on the system
Version hbase(main):010:0> version	This command returns the version of HBase used in your system.
Table_help hbase(main):02:0> table_help	This command guides you what and how to use table-referenced commands
Whoami hbase(main):008:0> whoami	This command returns the user details of HBase.

NoSQL Databases and Apache HBase

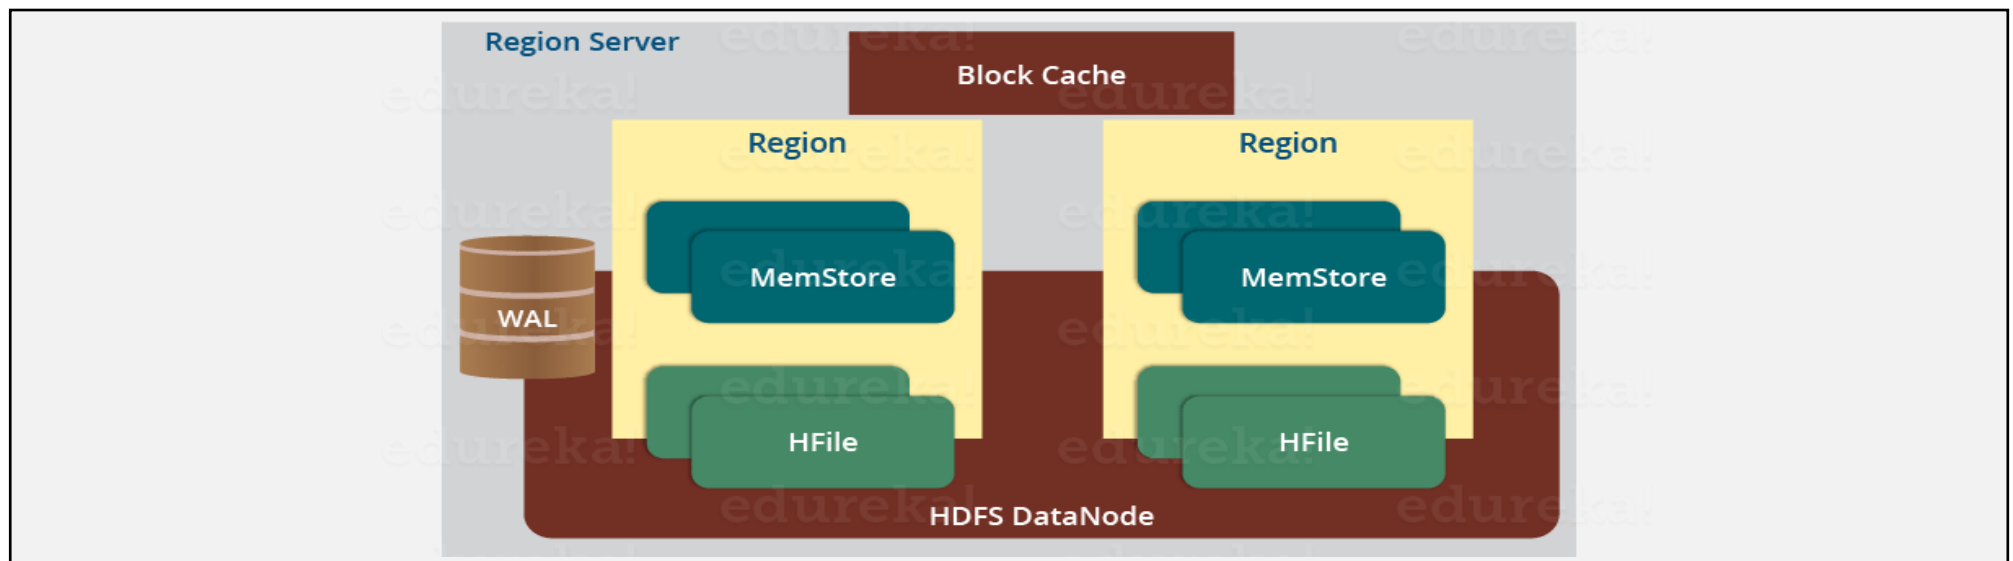
Region Server Components:

WAL: Write Ahead Log (WAL) is a file attached to every Region Server inside the distributed environment. The WAL stores the new data that hasn't been persisted or committed to the permanent storage. It is used in case of failure to recover the data sets.

Block Cache: The Block Cache resides in the top of Region Server. It stores the frequently read data in the memory. If the data in BlockCache is least recently used, then that data is removed from BlockCache.

MemStore: It is the write cache. It stores all the incoming data before committing it to the disk or permanent memory. There is one MemStore for each column family in a region. As you can see in the image, there are multiple MemStores for a region because each region contains multiple column families. The data is sorted in lexicographical order before committing it to the disk.

HFile: HFile is stored on HDFS. Thus it stores the actual cells on the disk. MemStore commits the data to HFile when the size of MemStore exceeds.



NoSQL Databases and Apache HBase

Read and Write Operations in HBase:

Steps for Read Operation:

Step 1: The client gets the information about region servers from META Table.

Step 2: The Client gets the information about row keys and the META table location.

Step 3: After getting information from META table Client communicates to Region Servers

Step 4: The client will retrieve the rows from the corresponding region servers.

Steps for Write operation:

Step 1: Each data is first written into Write Ahead Log (Wal). WAL is used as a backup or recovery process.

Step 2: After the data is written to WAL, Data is written to Memstore.

Step 3: Memstore is the intermediate store before the final commit to Hfile. The data is sorted before finally persisted in the disk. There is one Memstore per column family.

Step 4: When there is enough data in Memstore, the intermediate data is written into Hfile.

Step 5: Data is stored as key-value pair in Hfile.

NoSQL Databases and Apache HBase

HappyBase:

HappyBase is a developer-friendly Python library to interact with Apache HBase. HappyBase is designed for use in standard HBase setups, and offers application developers a Pythonic API to interact with HBase. HappyBase uses the Python Thrift library to connect to HBase using its Thrift gateway.

Application of HappyBase:

Pre requisites:

- 1) Linux virtual environment setup
- 2) HBase installed
- 3) A table created in HBase

Step 1: Install Python 3 and HappyBase package:

```
apk add python3  
pip install happybase
```

Step 2: Start Trift server:

```
alpine-hbase:/home/downloads/hbase-2.2.3/bin# ./hbase-daemon.sh start thrift  
running thrift, logging to /home/downloads/hbase-2.2.3/bin/../logs/hbase-root-thrift-alpine-hbase.out
```

Step 3: Start HBase server:

```
alpine-hbase:/home/downloads/hbase-2.2.3/bin# ./start-hbase.sh  
running master, logging to /home/downloads/hbase-2.2.3/bin/../logs/hbase-root-master-alpine-hbase.out  
alpine-hbase:/home/downloads/hbase-2.2.3/bin# _
```

Step 4: Now we open an Interactive Python session and make a connection to HBase to see that everything is working fine using Happybase APIs:

```
alpine-hbase:/home/downloads/hbase-2.2.3/bin# python3  
Python 3.8.1 (default, Dec 30 2019, 15:43:37)  
[GCC 9.2.0] on linux  
Type "help", "copyright", "credits" or "license" for more information.  
>>> import happybase as hb  
>>> conn = hb.Connection('127.0.0.1',9090)  
>>> conn.tables()  
[b'employers']  
>>> conn.table('employers').row('row1')  
{b'empAnagrafical:name': b'andrea', b'empAnagrafical:surname': b'guidi'}  
>>> conn.table('employers').row('row2')  
{b'empAnagrafical:name': b'george', b'empAnagrafical:surname': b'clooney'}  
>>>
```


NoSQL Databases and Apache HBase

Integration of HBase with Hive:

Example: Migrate the data from Hive to HBase table.

For setting up of HBase Integration with Hive, we mainly require a few jar files. The required jar files are:

```
zookeeper-*.jar          //This will be present in $HIVE_HOME/lib directory
hive-hbase-handler-*.jar //This will be present in $HIVE_HOME/lib directory
guava-*.jar              //This will be present in $HIVE_HOME/lib director
hbase-*.jar files        //This will be present in $HBASE_HOME/lib directory
```

Step 1: Create Hive table

```
hive> CREATE TABLE hive_table(
>   empno int,
>   ename string,
>   designation string,
>   manager int,
>   hire_date string,
>   sal int,
>   deptno int)
> ROW FORMAT DELIMITED
> FIELDS TERMINATED BY ','
> LINES TERMINATED BY '\n'
> STORED AS TEXTFILE;
OK
Time taken: 3.03 seconds
```

Step 2: Load data into Hive

```
hive> LOAD DATA LOCAL INPATH '/root/bdp/hbase/data/emp_data.csv' INTO TABLE hive_table;
Loading data to table db_bdp.hive_table
Table db_bdp.hive_table stats: [numFiles=1, totalSize=598]
OK
Time taken: 4.922 seconds
hive> select * from hive_table;
OK
7369 SMITH CLERK 7902 12/17/1980 800 20
7499 ALLEN SALESMAN 7698 2/20/1981 1600 30
7521 WARD SALESMAN 7698 2/22/1981 1250 30
7566 TURNER MANAGER 7839 4/2/1981 2975 20
7654 MARTIN SALESMAN 7698 9/28/1981 1250 30
7698 MILLER MANAGER 7839 5/1/1981 2850 30
7782 CLARK MANAGER 7839 6/9/1981 2450 10
7788 SCOTT ANALYST 7566 12/9/1982 3000 20
7839 KING PRESIDENT NULL 11/17/1981 5000 10
7844 TURNER SALESMAN 7698 9/8/1981 1500 30
7876 ADAMS CLERK 7788 1/12/1983 1100 20
7900 JAMES CLERK 7698 12/3/1981 950 30
7902 FORD ANALYST 7566 12/3/1981 3000 20
7934 MILLER CLERK 7782 1/23/1982 1300 10
Time taken: 5.041 seconds, Fetched: 14 row(s)
```

Step 3: Create HBase-Hive Mapping table

```
hive> CREATE TABLE hbase_table_employee
> (
>   empno INT,
>   ename STRING,
>   designation STRING,
>   manager INT,
>   hire_date STRING,
>   sal INT,
>   deptno INT
> )
> STORED BY 'org.apache.hadoop.hive.hbase.HBaseStorageHandler'
> WITH SERDEPROPERTIES ("hbase.columns.mapping" = ":key,cf:ename,
no")
> TBLPROPERTIES ("hbase.table.name" = "employee_hbase");
OK
Time taken: 17.237 seconds
```

Step 4: Load data into HBase from Hive

```
hive> INSERT INTO TABLE hbase_table_employee SELECT * FROM hive_table;
Query ID = root_20171228155155_22b3lbc-b-d741-4126-8d05-c64602e26312
Total jobs = 1
Launching Job 1 out of 1
Per session was closed. Recreating...
Session re-established.

Status: Running (Executing on YARN cluster with App id application_1514453878963_0009)

-----
VERTICES    STATUS  TOTAL  COMPLETED  RUNNING  PENDING  FAILED  KILLED
Map 1 ..... SUCCEEDED  1      1      0      0      0      0
VERTICES: 01/01 [=====] 100% ELAPSED TIME: 15.08 s
-----
OK
Time taken: 44.236 seconds
```

Step 5: Scan HBase Table

```
hbase(main):00810> scan 'employee_hbase'
ROW COLUMN+CELL
7369 column=cf:deptno, timestamp=1514476352028, value=20
7369 column=cf:designation, timestamp=1514476352028, value=CLERK
7369 column=cf:hire_date, timestamp=1514476352028, value=12/17/1980
7369 column=cf:manager, timestamp=1514476352028, value=7902
7369 column=cf:sal, timestamp=1514476352028, value=800
7499 column=cf:deptno, timestamp=1514476352028, value=30
7499 column=cf:designation, timestamp=1514476352028, value=SALESMAN
7499 column=cf:ename, timestamp=1514476352028, value=ALLEN
7499 column=cf:hire_date, timestamp=1514476352028, value=2/20/1981
7499 column=cf:manager, timestamp=1514476352028, value=7698
7499 column=cf:sal, timestamp=1514476352028, value=1600
7521 column=cf:deptno, timestamp=1514476352028, value=30
7521 column=cf:designation, timestamp=1514476352028, value=SALESMAN
7521 column=cf:hire_date, timestamp=1514476352028, value=2/22/1981
7521 column=cf:manager, timestamp=1514476352028, value=7698
7521 column=cf:sal, timestamp=1514476352028, value=1250
7566 column=cf:deptno, timestamp=1514476352028, value=20
7566 column=cf:designation, timestamp=1514476352028, value=MANAGER
7566 column=cf:ename, timestamp=1514476352028, value=TURNER
7566 column=cf:hire_date, timestamp=1514476352028, value=4/2/1981
7566 column=cf:manager, timestamp=1514476352028, value=7839
7566 column=cf:sal, timestamp=1514476352028, value=2975
7654 column=cf:deptno, timestamp=1514476352028, value=30
7654 column=cf:designation, timestamp=1514476352028, value=SALESMAN
```

NoSQL Databases and Apache HBase

MongoDB:

MongoDB is an open-source document database and leading NoSQL database. MongoDB is written in C++. A record in MongoDB is a document, which is a data structure composed of key value pairs similar to the structure of JSON objects.

Features of MongoDB:

- 1) Support ad hoc queries
- 2) Can index any field in a document
- 3) Duplication of data
- 4) Load balancing
- 5) Supports map reduce and aggregation tools.
- 6) Uses JavaScript instead of Procedures.
- 7) It is a schema-less database written in C++.
- 8) Provides high performance.
- 9) Stores files of any size easily without complicating your stack.

Application of HBase:

