

upGrad

NoSQL

**Databases and
Apache HBase**

Course: Data Engineering- I

Lecture On: NoSQL Databases
& Apache HBase

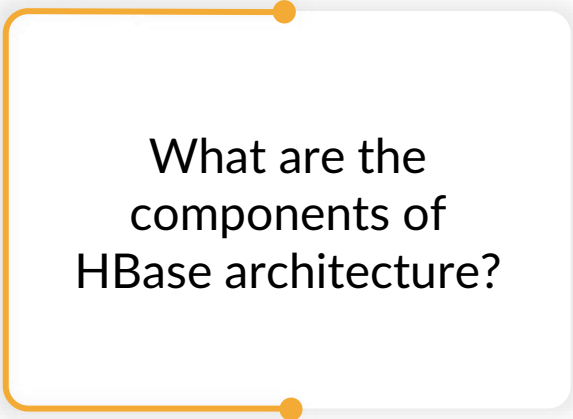
Session 3

How HBase works?

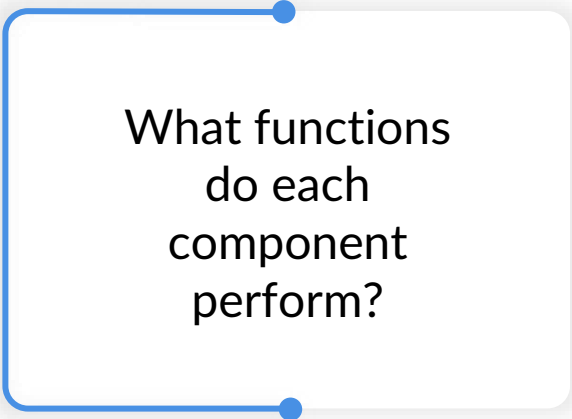
Segment 2

HBase Architecture

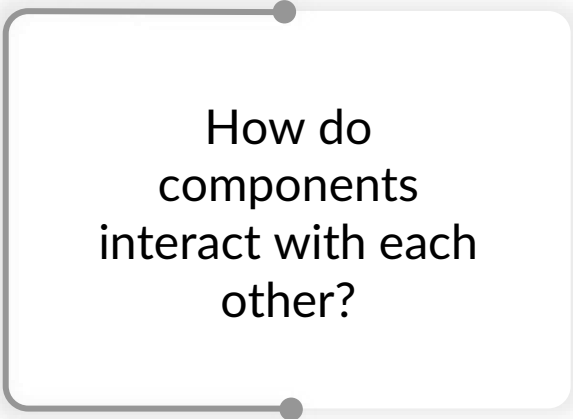
LEARNING OBJECTIVES

A white rounded rectangle with an orange border. The border is composed of a vertical line on the left, a horizontal line on the top, and a horizontal line on the bottom, with small orange circles at the top-right and bottom-right corners.

What are the
components of
HBase architecture?


A white rounded rectangle with a blue border. The border is composed of a vertical line on the left, a horizontal line on the top, and a horizontal line on the bottom, with small blue circles at the top-right and bottom-right corners.

What functions
do each
component
perform?

A white rounded rectangle with a grey border. The border is composed of a vertical line on the left, a horizontal line on the top, and a horizontal line on the bottom, with small grey circles at the top-right and bottom-right corners.

How do
components
interact with each
other?


HBASE ARCHITECTURE



HBase Client

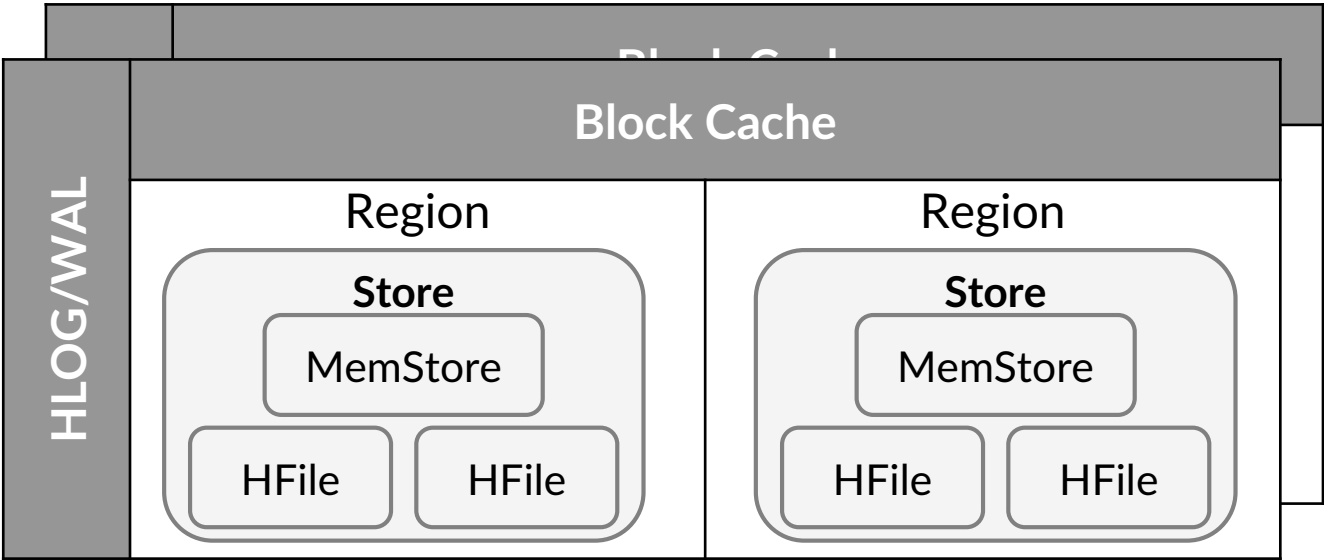


Zookeeper




HMaster

Hmaster
Servers



RegionServer



HFile

...

DataNode(HDFS)

HBASE ARCHITECTURE: HMASTER

- Acts as the master server
- Manages multiple region servers
- An HBase cluster may have one or more master nodes
- Only one HMaster is active at a given time

HBASE ARCHITECTURE: HMASTER


This active master node is responsible for:

- Performing the admin functionalities
- Coordinating with region server
 - **Load Balancing:** When there is a high load on any region server, HMaster unloads the busy server and reassigns this load to less occupied servers.
 - **Recovery:** HMaster also handles region server failures by reassigning the regions or load of the failed region server to another functioning region server.


HBASE ARCHITECTURE: REGION SERVERS



HBase Client

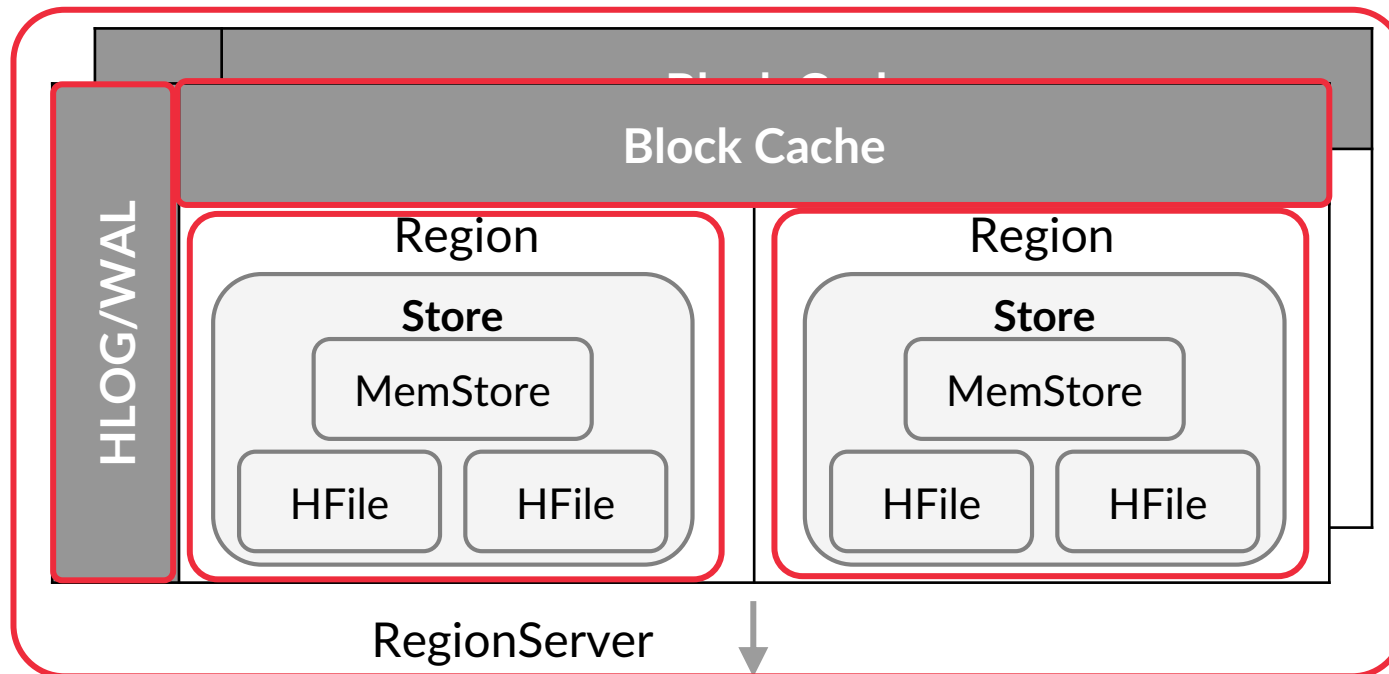


Zookeeper



HMaster

Hmaster
Servers



DataNode(HDFS)

HBASE ARCHITECTURE: REGION SERVERS

- It is the Slave server.
- There are multiple region servers in a cluster.
- Each region server contains:
 - **WAL** (Write Ahead Log), which stores new or updated data that has not been written to HDFS. The region server stores the WAL file in HDFS.

HBASE ARCHITECTURE: REGION SERVERS

- It is the Slave server.
- There are multiple region servers in a cluster.
- Each region server contains:
 - **WAL** (Write Ahead Log), which stores new or updated data that has not been written to HDFS. The region server stores the WAL file in HDFS.
 - **Block Cache**, which stores frequently read data from HDFS.

HBASE ARCHITECTURE: REGION SERVERS

- It is the Slave server.
- There are multiple region servers in a cluster.
- Each region server contains:
 - **WAL** (Write Ahead Log), which stores new or updated data that has not been written to HDFS. The region server stores the WAL file in HDFS.
 - **Block Cache**, which stores frequently read data from HDFS.
 - **Region**, which contains a Store per column family. The Store consists of a Memstore.

HBASE TABLE


Lexicographic
Order

Row key	A		B	
	Col-1	Col-2	Col-1	Col-2
Row1	2	4	6	4
Row11	1	6	5	8
Row2	0	9	3	5
	5	1	0	8
Row30	6	2	4	0
Row4	2	5	7	9
Row6	6	3	1	0

REGIONS

Region1	[Row key	A: Col-1	A: Col-2	B: Col-1	B: Col-2
		Row1	2	4	6	4
		Row11	1	6	5	8
Region2	[Row2	0	9	3	5
			5	1	0	8
		Row30	6	2	4	0
Region3	[Row4	2	5	7	9
		Row6	6	3	1	0

HBASE ARCHITECTURE



HBase Client

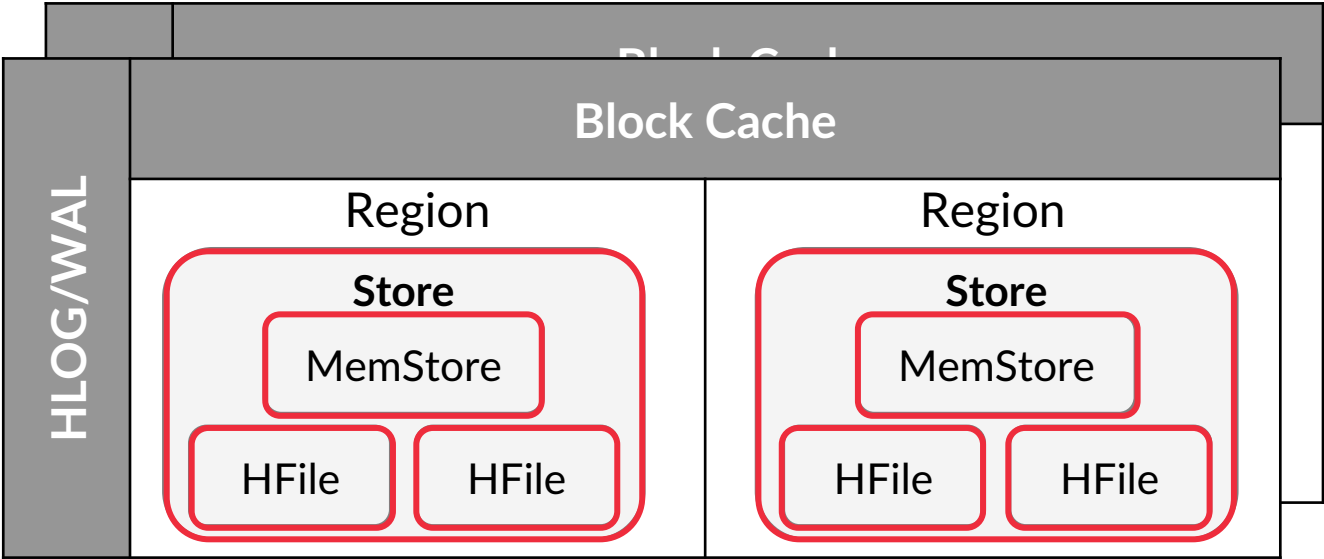


Zookeeper




HMaster

Hmaster
Servers

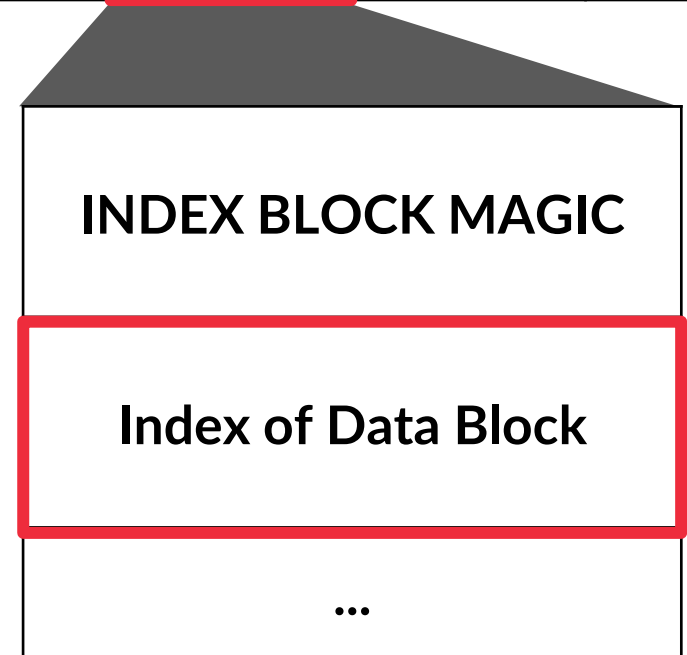
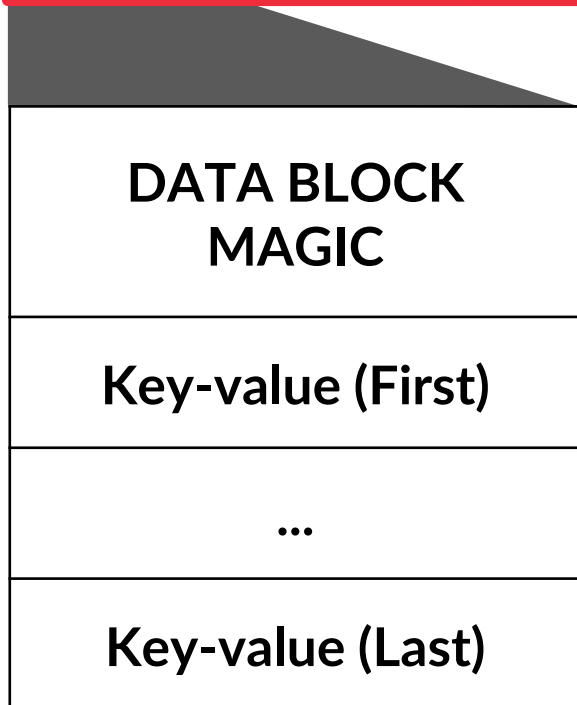



RegionServer



DataNode(HDFS)

HFILE STRUCTURE




B+ TREES

- These are data structures used to store large amounts of data which cannot be stored in the main memory.
- The Internal nodes are stored in the main memory.
- The Leaf nodes are stored in the secondary memory.


HBASE ARCHITECTURE: REGION SERVERS

- It is the Slave server.
- It manages regions. There are multiple region servers in a cluster.
- Each region server contains:
 - **WAL** (Write Ahead Log), which stores new or updated data that has not been written to HDFS. The region server stores the WAL file in HDFS.
 - **Block Cache**, which stores frequently read data from HDFS.
 - **Region**, which contains one Store per column family. The Store consists of a Memstore.
- A client talks directly with the region server to perform read/write operations.
- This region server assigns the request to a specific region.

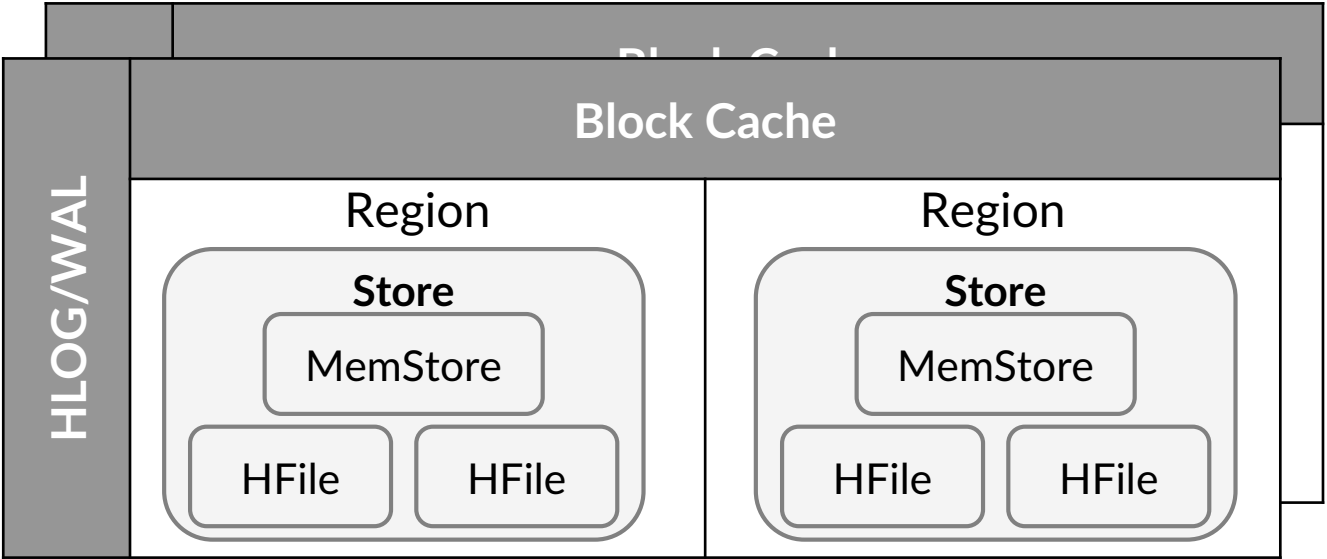
HBASE ARCHITECTURE: ZOOKEEPER




HBase Client



Hmaster
Servers




RegionServer




DataNode(HDFS)

HBASE ARCHITECTURE: ZOOKEEPER

- It is a distributed open-source coordinating service for distributed applications.
- It maintains the configuration information of clusters and naming services of clusters, in addition to providing distributed synchronisation, group services, etc.
- HBase uses Zookeeper to maintain the live server state in the cluster.




ZOOKEEPER IN HBASE CLUSTER



HBASE ARCHITECTURE: ZOOKEEPER

- It is a distributed open-source coordinating service for distributed applications.
- It maintains the configuration information of clusters and naming services of clusters in addition to providing distributed synchronisation, group services, etc.
- HBase uses Zookeeper to maintain the live server state in the cluster.
- It provides server failure notifications.
- It stores the location of the META table.



KEY TAKEAWAYS

- HBase architecture is based on the concept of master-slave architecture.
- There are three servers: HMaster, Region Servers and Zookeeper.
- HMaster is the master server and manages the region servers.
- A region server acts as the slave server and manages multiple regions.
- A region has multiple stores, a Memstore and HFiles where actual data is present.
- Zookeeper is responsible for the communication between HMaster and region servers.

Thank You!

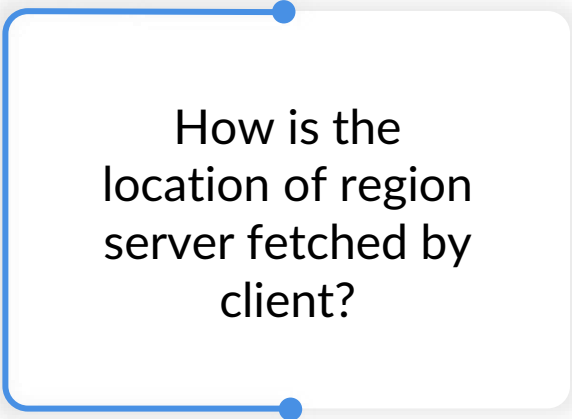
Segment 3

Read Operation in HBase

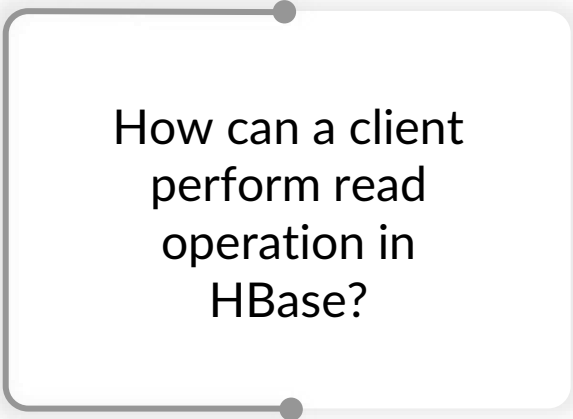
LEARNING OBJECTIVES



What is Meta Table
and its structure?



How is the
location of region
server fetched by
client?



How can a client
perform read
operation in
HBase?

META TABLE

Table

Region Server




META TABLE

Zookeeper

META Location

Region Server Location

Row

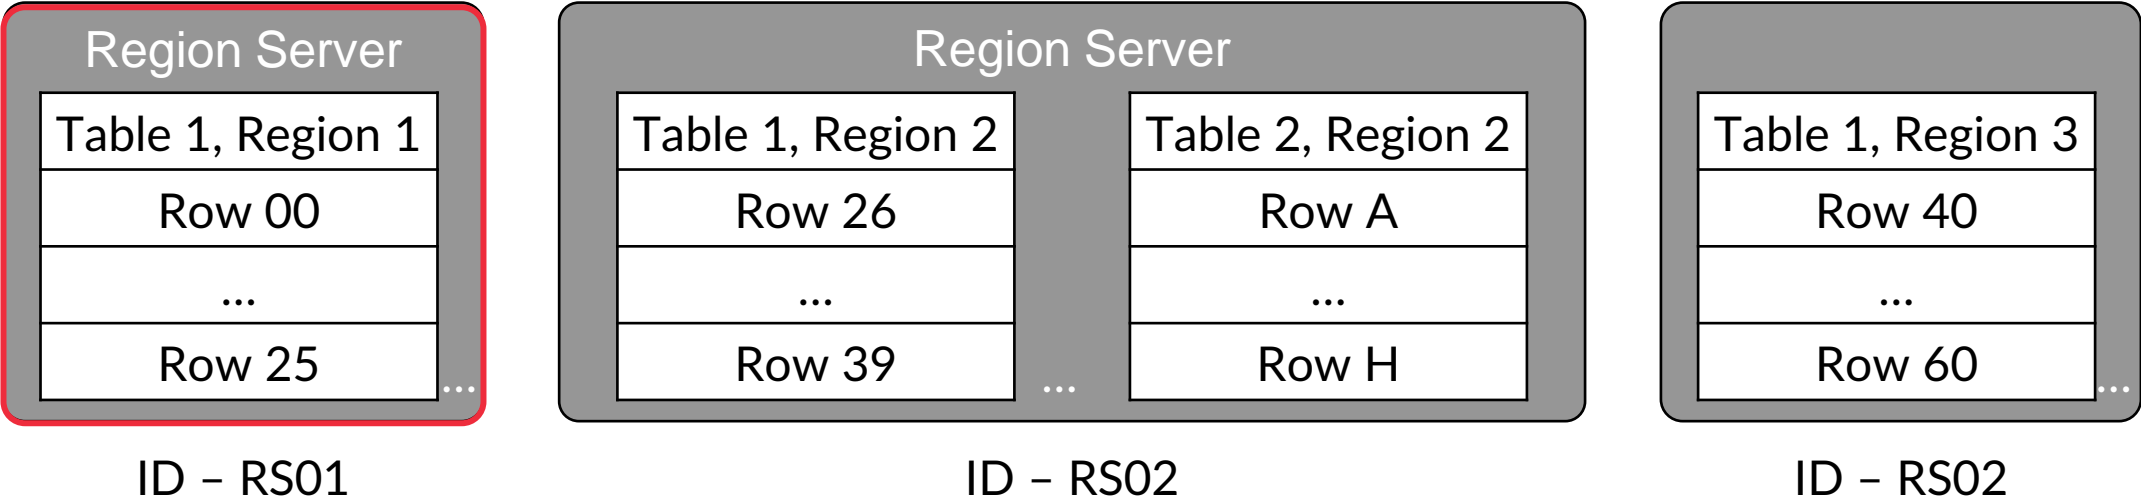


META TABLE

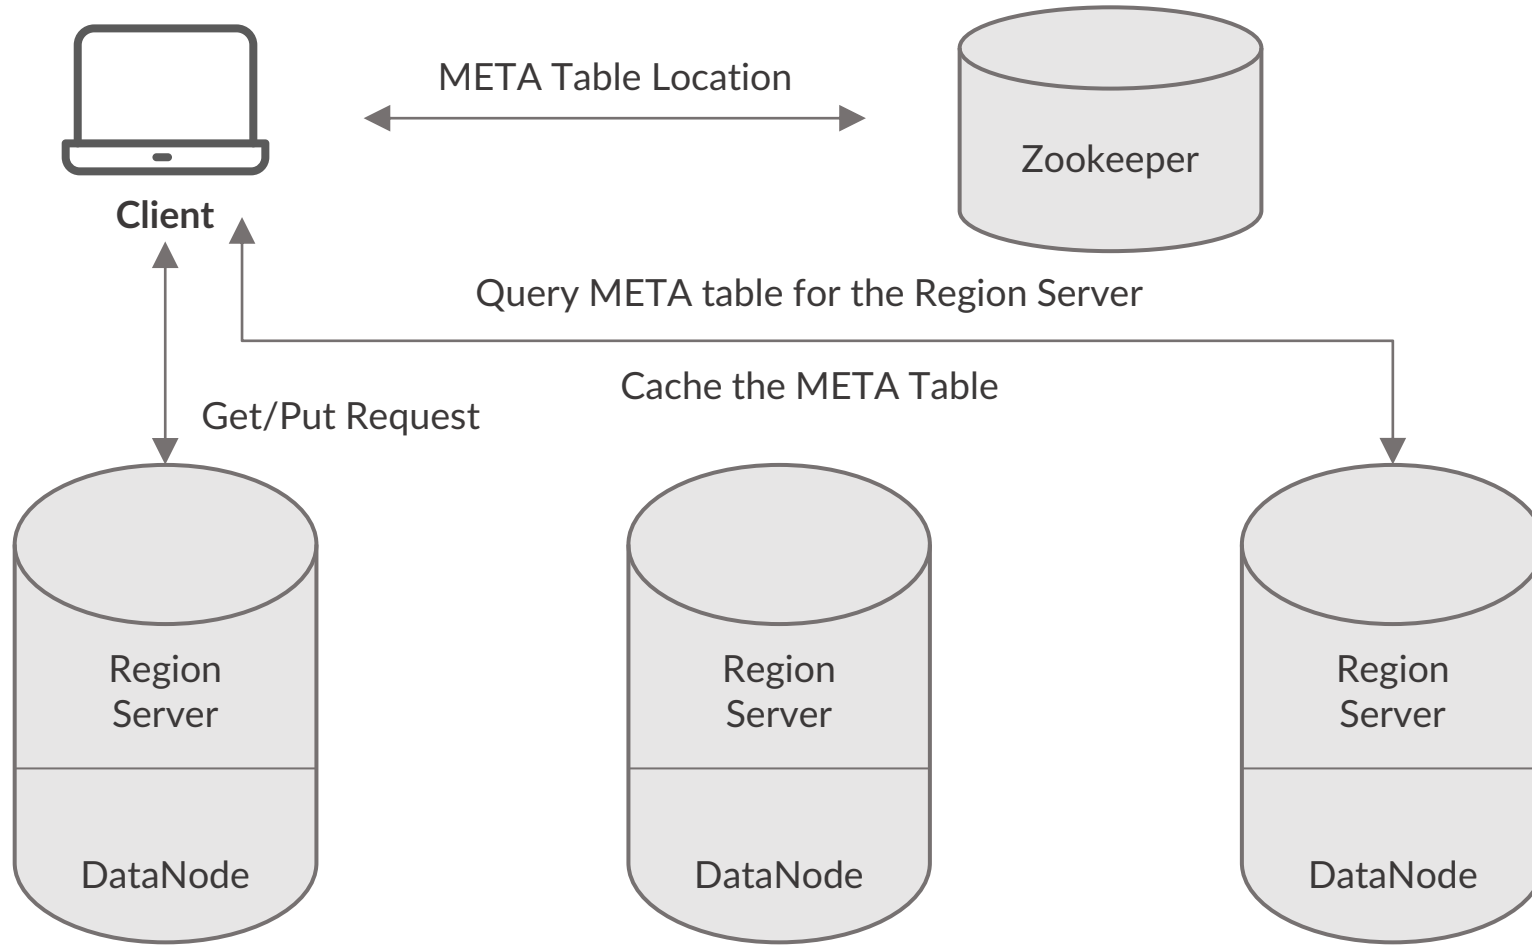
Key		Value	
Table Name	Start Key	Region ID	Region ServerID
Table 1	Row 00	1	RS01
Table 1	Row 26	2	RS02
Table 1	Row 40	3	RS03
...			
Table 2	RS01	RS01	RS01
...			

Key: <table Name, Start row Key, regionID>

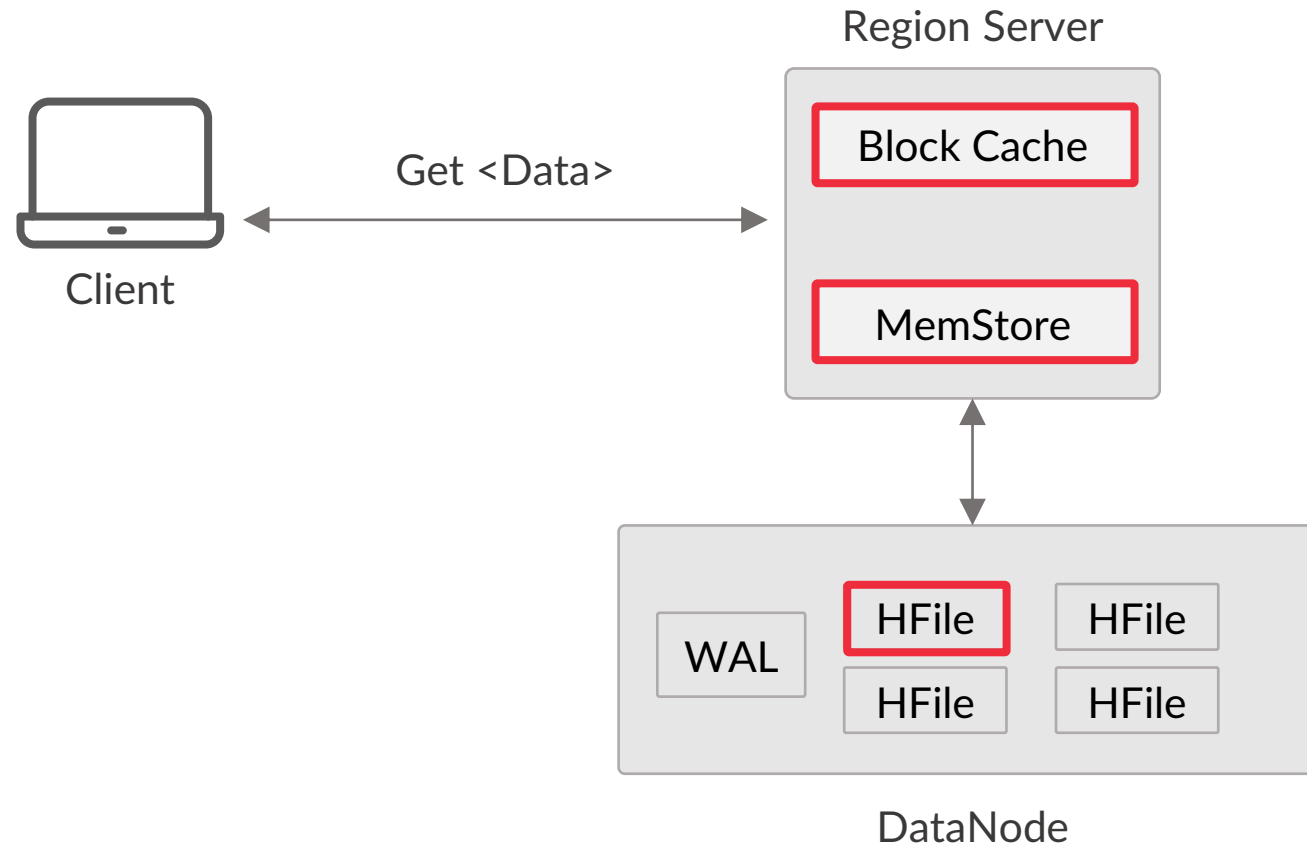
Value: <Region Server ID>



FETCHING THE REGION SERVER



HBASE READ OPERATION



HBASE READ OPERATION

- The region server first checks its **Block Cache**, which stores the recently accessed data.

HBASE READ OPERATION

- The region server first checks its **Block Cache**, which stores the recently accessed data.
- If the data is not available in the Block Cache, it checks the required data in the in-memory store, i.e., **MemStore**.

HBASE READ OPERATION

- The region server first checks its **Block Cache**, which stores the recently accessed data.
- If the data is not available in the Block Cache, it checks the required data in the in-memory store, i.e., **MemStore**.
- If the MemStore does not contain that particular key-value, the region server uses **Bloom Filters** to find the HFile(StoreFile), which contains that particular key-value pair.

BLOOM FILTERS

- It is designed to check whether an element is present in a set or not.
- It has a probabilistic data structure.
- It has a bit vector-based data structure.

HBASE READ OPERATION

- The region server first checks its **Block Cache** which stores the recently accessed data.
- If the data is not available in the Block Cache, it checks the required data in the in-memory store, i.e., **MemStore**.
- If the MemStore does not contain that particular key-value, the region server uses **Bloom Filters** to find the HFile(StoreFile) which contains that particular key-value pair.
- Once the **HFile** is identified, data is read in an efficient manner and returned to the client (if present).

KEY TAKEAWAYS

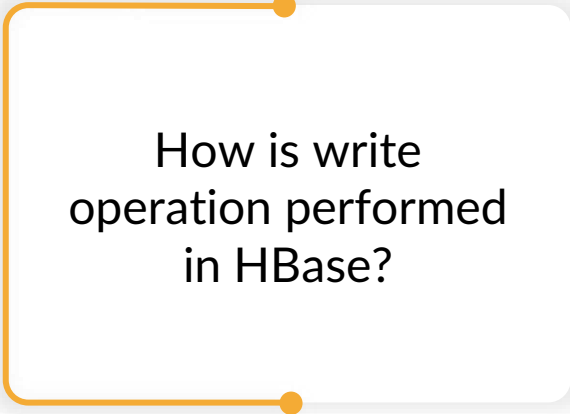
- Meta Table stores the location of region servers.
- The location of Meta Table is stored in Zookeeper.
- The client first fetches the location of region server where the region having the required data is present.
- For reading the data:
 - First, the block cache is checked, and if the data is not found in the block cache, the Memstore is checked for the data.
 - If it is not found in the Memstore, then the HFile, where the data is present, is searched using the bloom filters.
 - The data is read from the HFile.

Thank You!

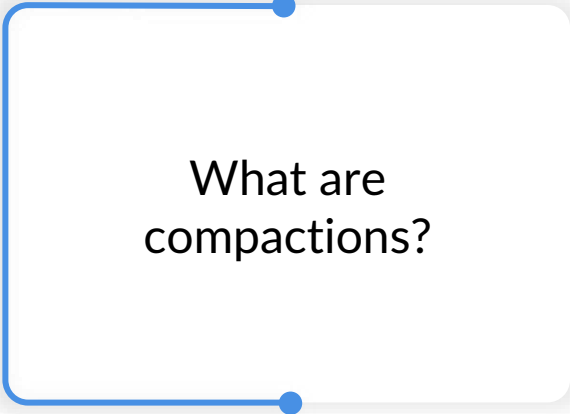
Segment 4

Write Operation in HBase


LEARNING OBJECTIVES

A white rounded rectangle with an orange border. The border is composed of a vertical line on the left, a horizontal line on the top, and a horizontal line on the bottom, with small orange dots at the top-right and bottom-right corners.

How is write
operation performed
in HBase?


A white rounded rectangle with a blue border. The border is composed of a vertical line on the left, a horizontal line on the top, and a horizontal line on the bottom, with small blue dots at the top-right and bottom-right corners.

What are
compactions?

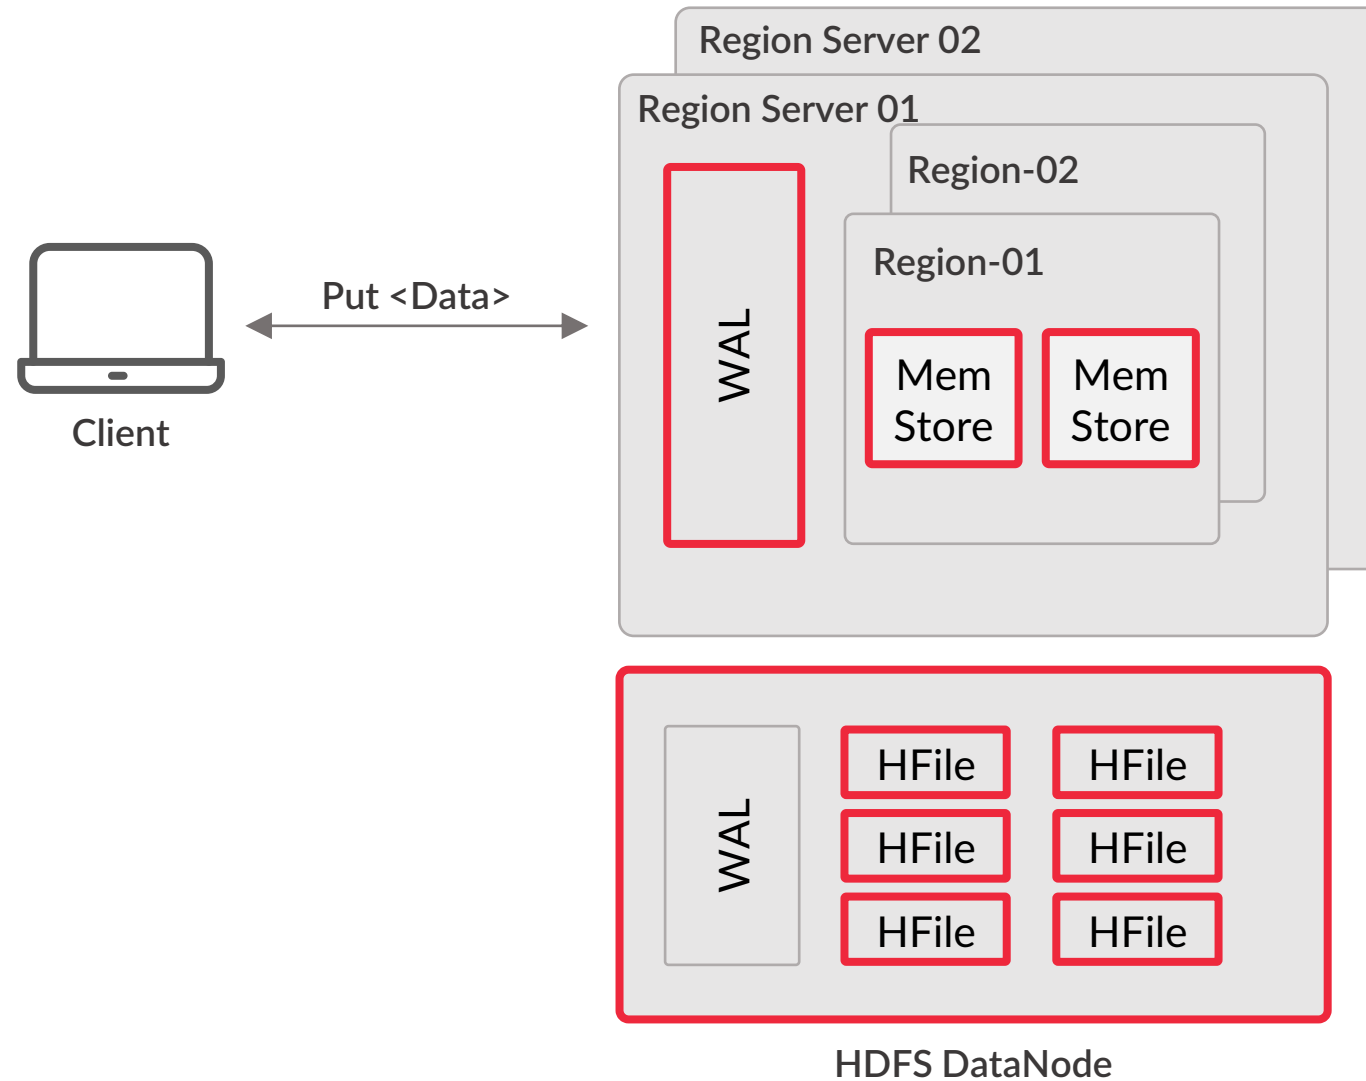
A white rounded rectangle with a grey border. The border is composed of a vertical line on the left, a horizontal line on the top, and a horizontal line on the bottom, with small grey dots at the top-right and bottom-right corners.

What are the
steps in delete
operation?

FETCHING THE REGION SERVER



HBASE WRITE OPERATION



HBASE WRITE OPERATION

- The data first needs to be written to the **WAL**.

HBASE WRITE OPERATION

- The data first needs to be written to the **WAL**.
- After that, it is placed in the **MemStore** of a region and the key-value pairs.

HBASE WRITE OPERATION

- The data first needs to be written to the **WAL**.
- After that, it is placed in the **MemStore** of a region and the key-value pairs.
- When the MemStore becomes full, its contents are flushed to HDFS(DataNode) to form a new HFile.

HBASE WRITE OPERATION

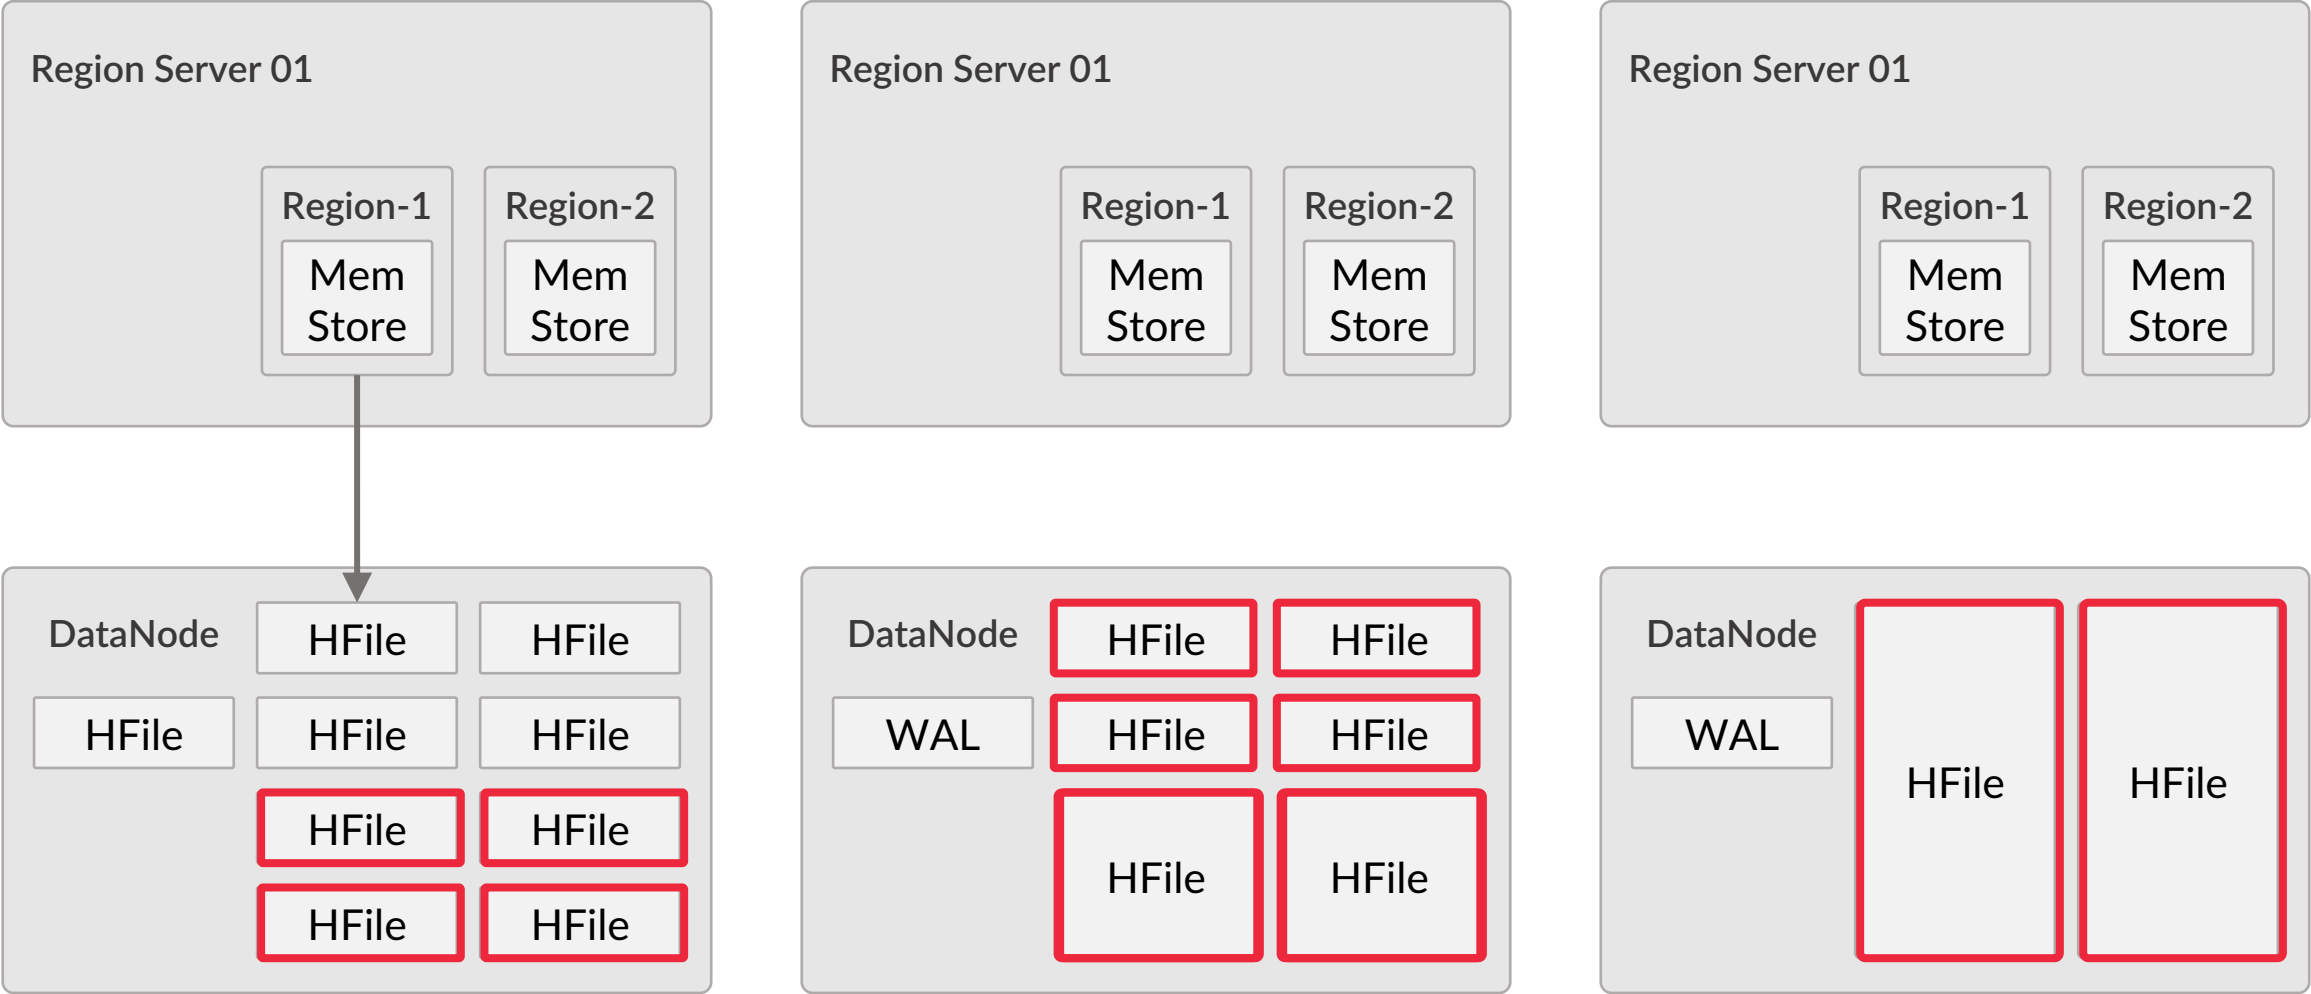
- The data first needs to be written to the **WAL**.
- After that, it is placed in the **MemStore** of a region and the key-value pairs.
- When the MemStore becomes full, its contents are flushed to HDFS(DataNode) to form a new HFile.
- Finally, an acknowledgement is sent back to the client.

COMPACTION

- Heavy write operations lead to two challenges:
 - Low efficiency of read operations
 - Dirty data
- Compaction is the process of combining the smaller HFiles into larger HFiles.
- Two types:
 - **Minor Compactions:** Smaller HFiles are combined to form large HFiles.

COMPACTION

Minor and Major



Minor Compaction

Major Compaction

COMPACTION

- Heavy write operations lead to two challenges:
 - Low efficiency of read operations
 - Dirty data
- Compaction is the process of combining the smaller HFiles into larger Hfiles.
- Two types:
 - **Minor Compactions:** Smaller HFiles are combined to form large HFiles.
 - **Major Compactions:** Store files of a region are combined into a single Store File. It removes all the deleted and expired data from the HFiles.

HBASE DELETE OPERATION

- This is a special type of Update operation.
- The deleted value is marked with a **Tombstone Marker**.
- When the deleted data is accessed, NULL is returned.
- HFiles are immutable.
- Deleted data is removed when Major Compaction is done.
- Three types of Tombstone Markers:
 - Version Delete Marker
 - Column Delete Marker
 - Family Delete Marker

KEY TAKEAWAYS

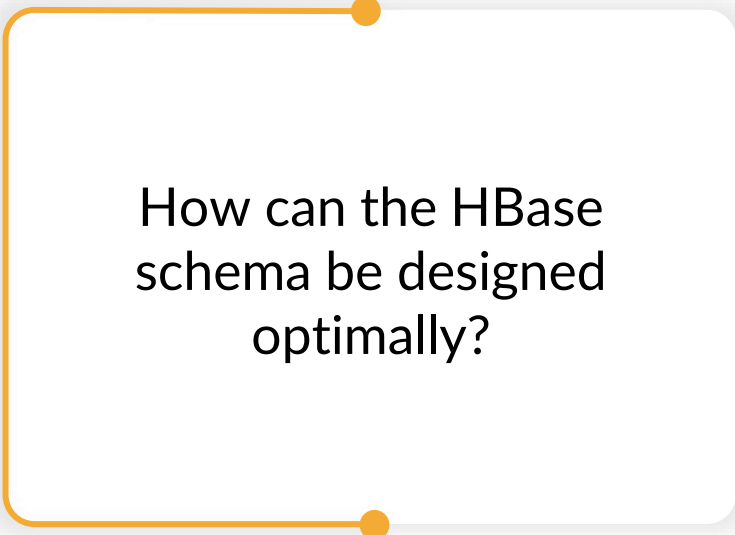
- For write operation:
 - First the data is written to write ahead log (WAL).
 - The data from WAL is moved to Memstore. Once the Memstore is full, the data is flushed into HDFS creating a new HFile.
- Heavy write operations lead to dirty data and reduction in read efficiency.
- Compaction is a technique in which smaller HFiles are combined to form larger HFiles.
- There are two types of compactions: major and minor.
- The delete operation is a special type of update operation.

Thank You!

Segment 5

HBase Schema Design

LEARNING OBJECTIVES



How can the HBase schema be designed optimally?



How can filters be applied while reading data?

SCHEMA DESIGN

- Consider-
 - n- number of rows
 - e- number of key-value pairs in HFile
 - b- number of blocks
 - c- number of columns

SCHEMA DESIGN

- If data is found in **Block Cache**, Complexity- $O(1)$
- If data is found in **MemStore**, Complexity- $O(\log(e))$
- If not found in MemStore-
 - To find the required **block** in HFile, Complexity- $O(\log(b))$
 - If data found in the same block, Complexity- $O(b/e)$
 - If not found in the same block, Complexity- $O(c)$
- If c increases, complexity increases.
- Rows should be equally distributed across region servers.
- Filtering on the basis of column, timestamp and version.
- Denormalise data.

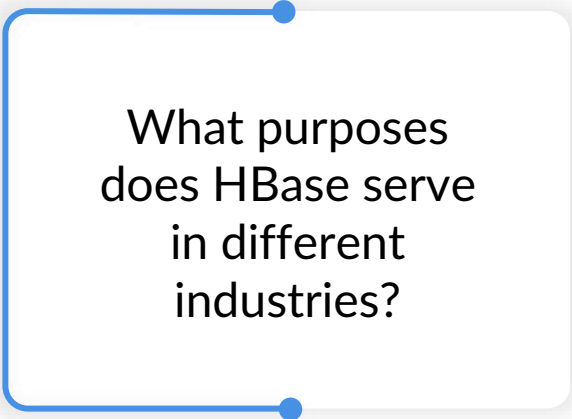
Segment 6

HBase Use Cases

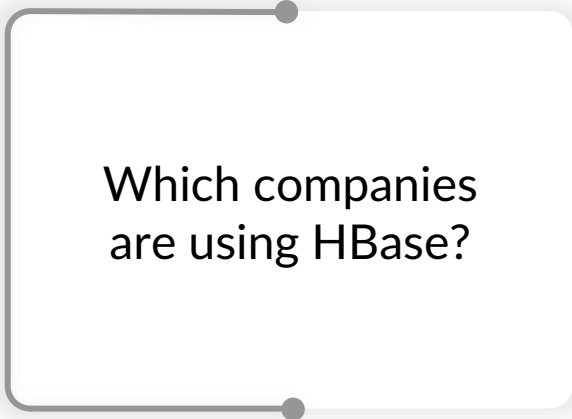
LEARNING OBJECTIVES

A white rounded rectangle with an orange border. The border is composed of a vertical line on the left, a horizontal line on the top, and a horizontal line on the bottom, with small orange dots at the top and bottom corners.

Which industries are using HBase?

A white rounded rectangle with a blue border. The border is composed of a vertical line on the left, a horizontal line on the top, and a horizontal line on the bottom, with small blue dots at the top and bottom corners.

What purposes does HBase serve in different industries?

A white rounded rectangle with a grey border. The border is composed of a vertical line on the left, a horizontal line on the top, and a horizontal line on the bottom, with small grey dots at the top and bottom corners.


Which companies are using HBase?

FIELDS OF USE

Medical



Banking



Telecom



E-Commerce



Sports



Finance

HBASE IN THE INDUSTRY

Bloomberg

FLIPBOARD

photobucket

HORTONWORKS

OPower

Cerner

Optimizely

Dropbox

jive
an Aurea company

appnexus



YAHOO!

FINRA

splice
MACHINE



intel



IBM

TREND
MICRO

explorys

FLURRY

HUAWEI

GROUPON

淘宝网
Taobao.com

airbnb

salesforce

AOL

box

KLOUT

Pinterest



OCLC

cloudera

CyberAgent

ebay

HubSpot

Google

Limelight
NETWORKS

Microsoft

NGDATA

rocketfuel

HBASE IN THE INDUSTRY

- The use cases that HBase serves in the industry are very diverse, for example:
 - **Yahoo** uses HBase for content personalisation and web cache for searching.
 - E-commerce giant **eBay** uses HBase for optimising eBay's search engine performance.
 - **Pinterest** uses HBase for relevant search.
 - **Bloomberg** uses HBase for time series data storage.
 - **Facebook** used to employ HBase for their messaging application.

KEY TAKEAWAYS

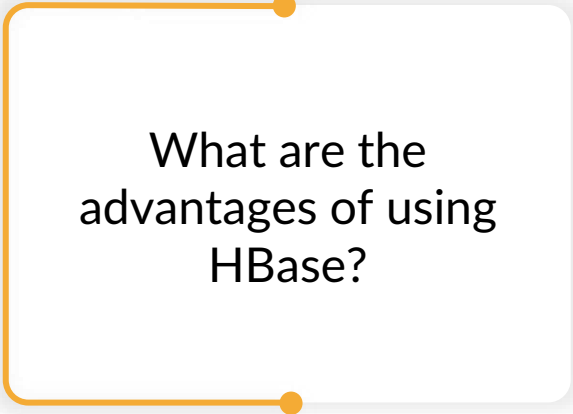
- HBase serves industries like e-commerce, sports, medical, etc.
- Companies like Yahoo, eBay, Airbnb and Intel are using HBase for various purposes.

Thank You!

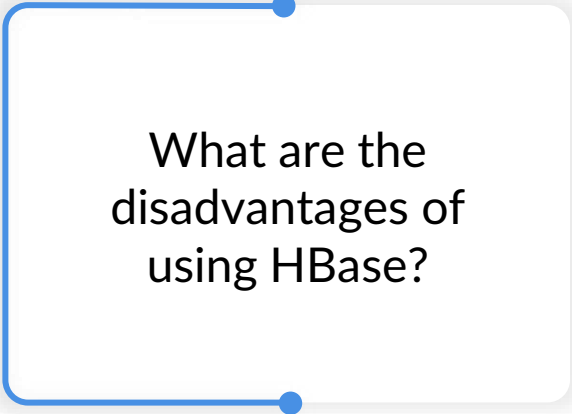
Segment 7

HBase Advantages and Disadvantages

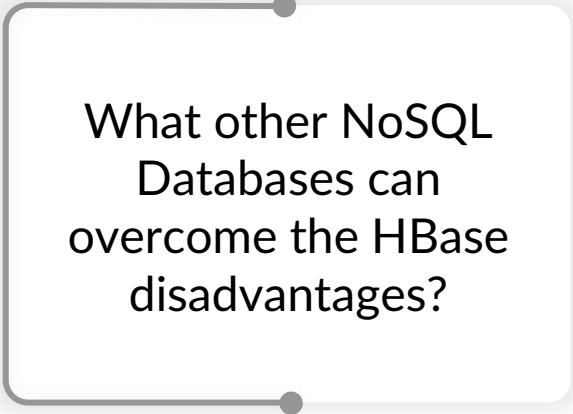
LEARNING OBJECTIVES

A white rounded rectangle with an orange border. The border is composed of a vertical line on the left, a horizontal line at the top, and a horizontal line at the bottom, with small orange circles at the top and bottom corners.

What are the advantages of using HBase?

A white rounded rectangle with a blue border. The border is composed of a vertical line on the left, a horizontal line at the top, and a horizontal line at the bottom, with small blue circles at the top and bottom corners.

What are the disadvantages of using HBase?

A white rounded rectangle with a grey border. The border is composed of a vertical line on the left, a horizontal line at the top, and a horizontal line at the bottom, with small grey circles at the top and bottom corners.


What other NoSQL Databases can overcome the HBase disadvantages?

HBASE: ADVANTAGES

Distributed
Databases



Schema-Less



Query
Optimisation



Read/Write
Consistency



Less Reading and
Processing Time



Horizontal
Scalability

HBASE: DISADVANTAGES

No Built-in
Authentication



Resource
Intensive



Does Not Support
SQL Functions Like
joins, groupby, etc.



Single Sort Key

HBASE: DISADVANTAGES



- Single point of failure
- No query language available
- Key-value pair



- The system is always available.
- Cassandra Query Language (CQL)

SESSION SUMMARY

- HBase architecture has three servers- HMaster, Region Servers and Zookeeper.
- HBase architecture is a master-slave architecture, with HMaster being the master and Region server being the slave.
- Zookeeper is responsible for the communication between HMaster and region server.
- When smaller HFiles are combined to form larger HFiles, it is called compaction.
- Equal distribution of rows can improve the schema design.
- HBase is used in industries like medical, e-commerce, sports etc.

Thank You!