

# EC2 INSTANCE SETUP FOR KAFKA

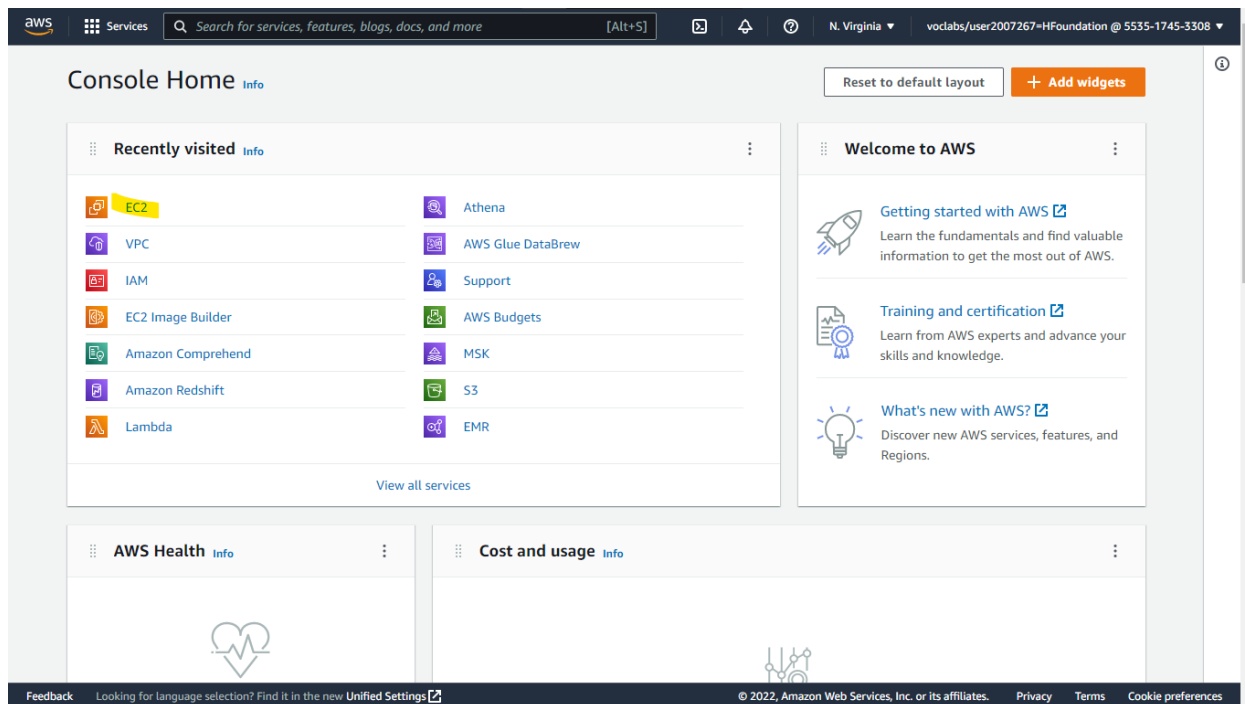
This document guides you through the setup of an EC2 instance which has been set up exclusively for Kafka. The following services are required for working with the Kafka service in the module:

1. Java 1.8
2. Zookeeper
3. Kafka
4. Jupyter Notebook

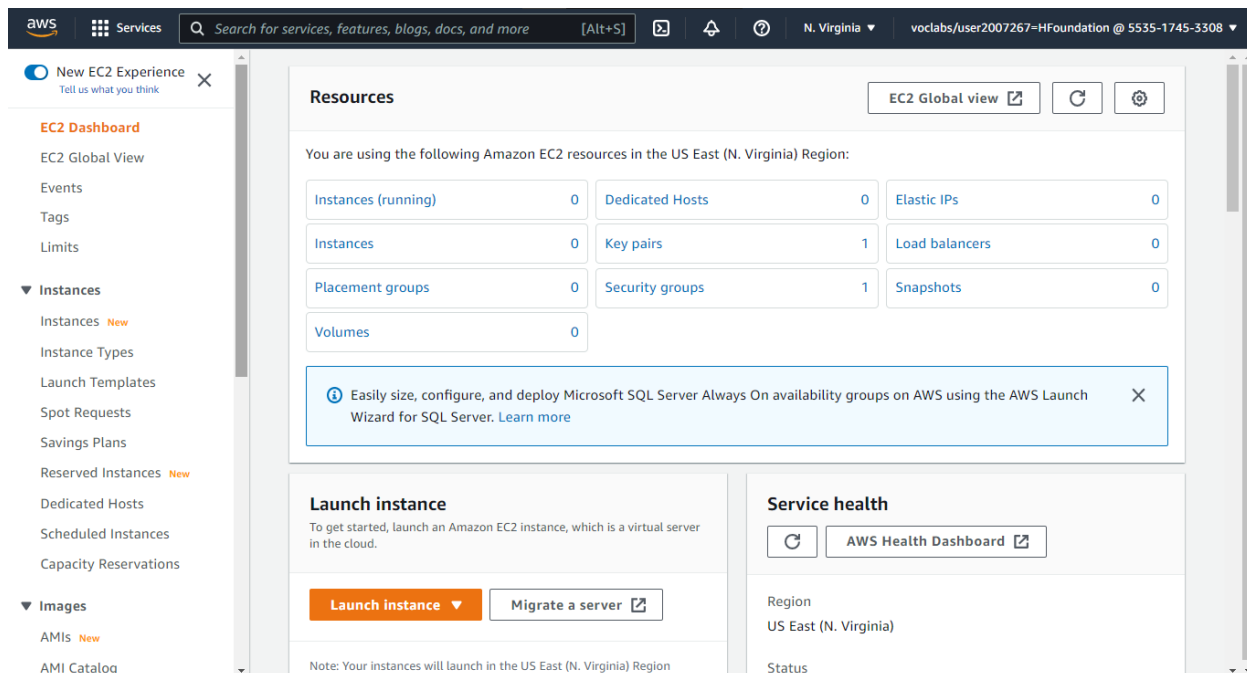
There're OS images published as AMIs present in the AWS Marketplace that contain these services. In such cases you can directly use the AMI in your EC2 instance.

This documentation contains instructions to set up all the required services to an EC2 instance from scratch.

- Go to the AWS Management Console and click on “All Services” and then click on “EC2” under the heading “Compute.”



- Once you click on that, you will be redirected to a new page, shown below. Click on the “Launch Instance” button.

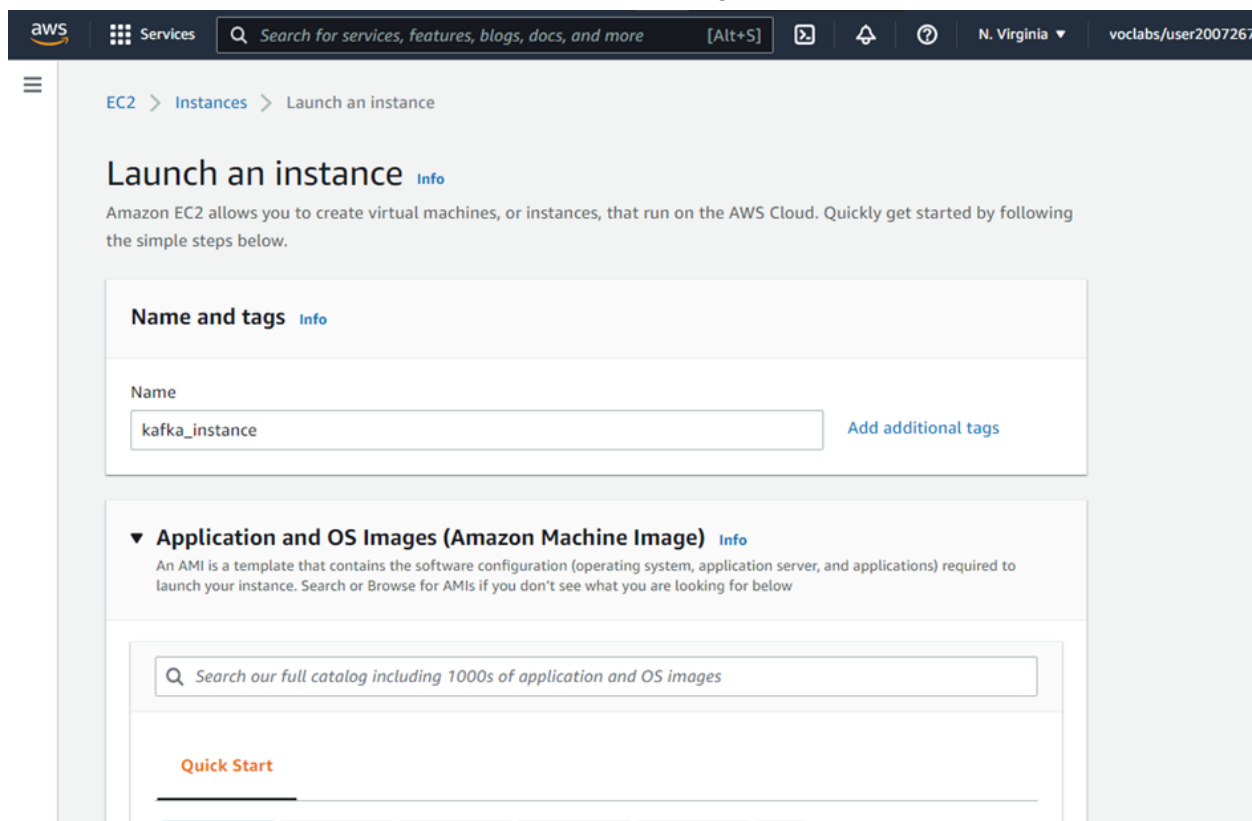


The screenshot shows the AWS Management Console interface. The top navigation bar includes the AWS logo, a search bar, and the current region (N. Virginia). The left sidebar shows the 'EC2 Dashboard' with various navigation options. The main content area is titled 'Resources' and shows a summary of EC2 resources in the US East (N. Virginia) Region. A table lists the following resources:

Resource	Count
Instances (running)	0
Dedicated Hosts	0
Elastic IPs	0
Instances	0
Key pairs	1
Load balancers	0
Placement groups	0
Security groups	1
Snapshots	0
Volumes	0

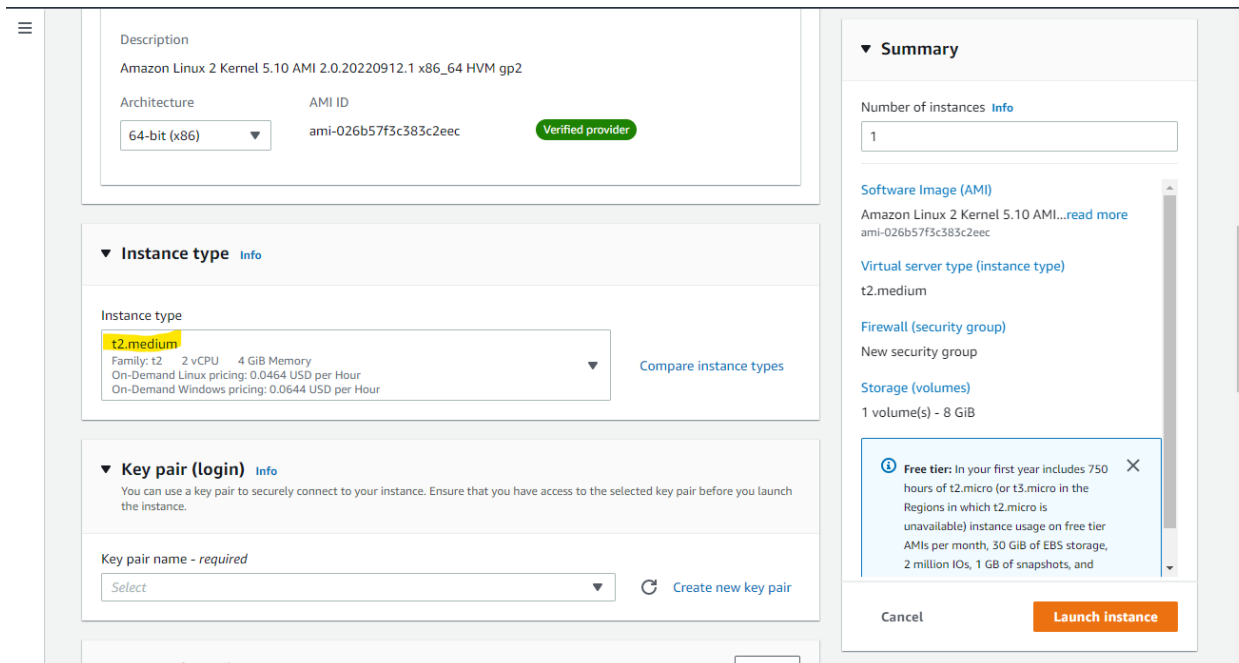
Below the table, there is a 'Launch instance' section with a 'Launch instance' button and a 'Migrate a server' link. A 'Service health' section shows the AWS Health Dashboard link. A note at the bottom states: 'Note: Your instances will launch in the US East (N. Virginia) Region'.

- You'll be redirected the EC2 instance launch page.

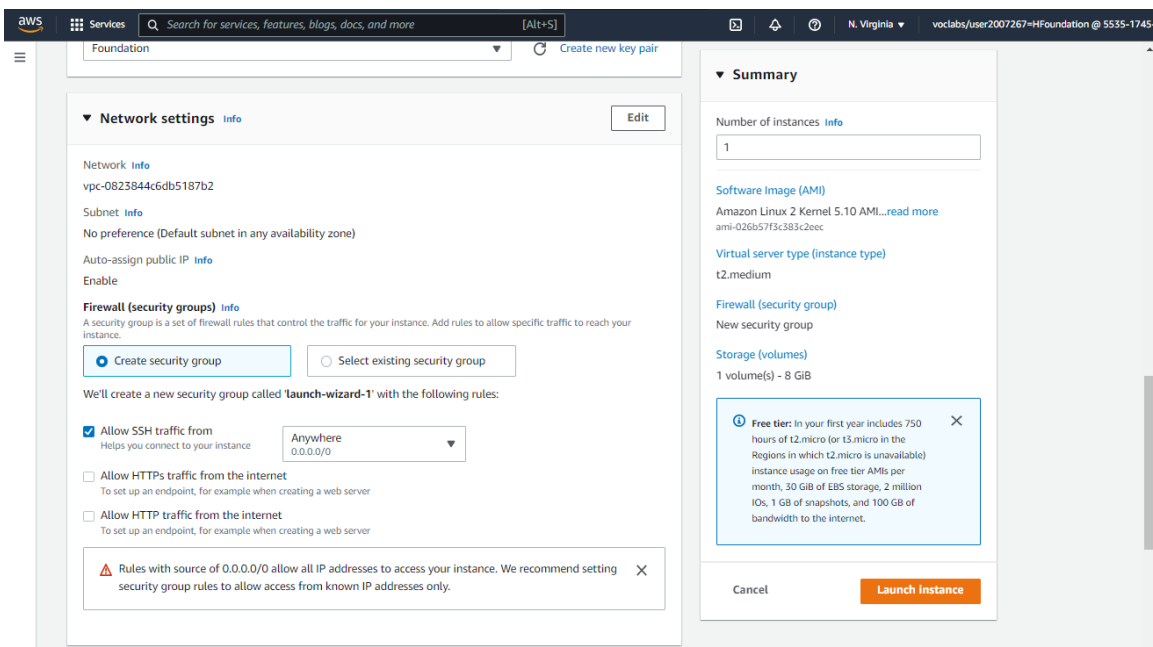


The screenshot shows the 'Launch an instance' page in the AWS Management Console. The page title is 'Launch an instance' with an 'Info' link. The main content area provides instructions on how to launch an instance. Below the instructions, there is a 'Name and tags' section with a text input field containing 'kafka\_instance' and an 'Add additional tags' link. Below this, there is a section for 'Application and OS Images (Amazon Machine Image)' with a search bar and a 'Quick Start' section.

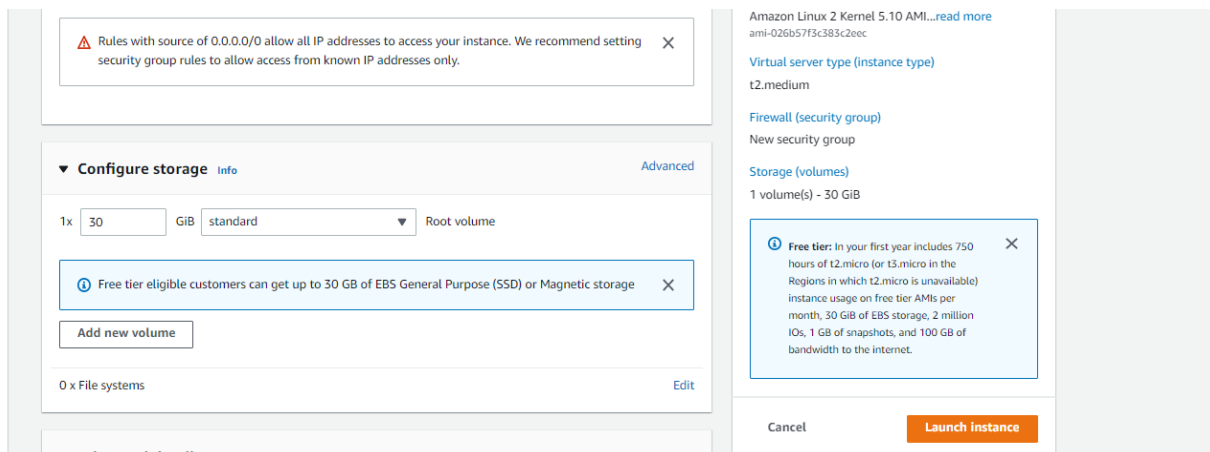
- Select the **General Purpose t2.medium** type EC2 instance, as shown in the image below. Using any other instance type, such as t2.micro which has 1 vCPU and 1 GB memory, may lead to memory errors while working with Kafka. You also need to select the Key Pair to login to the instance via SSH.



- In the “Network Settings” section, go with the default security group. Make sure that the option “Allow SSH traffic from Anywhere” is ticked. This will ensure that you’re able to SSH into the instance from your SSH client. The settings for this will be modified in the later steps.



- In the “Configure storage” settings, you need to enter the volume size as **30 GiB** and volume type as **standard (magnetic)**, as shown in the image below. Once the settings have been updated, click on the “Launch Instance” button to create the instance.



Rules with source of 0.0.0.0/0 allow all IP addresses to access your instance. We recommend setting security group rules to allow access from known IP addresses only.

Amazon Linux 2 Kernel 5.10 AMI...[read more](#)  
ami-026b57f3c383c2e0c

Virtual server type (instance type)  
t2.medium

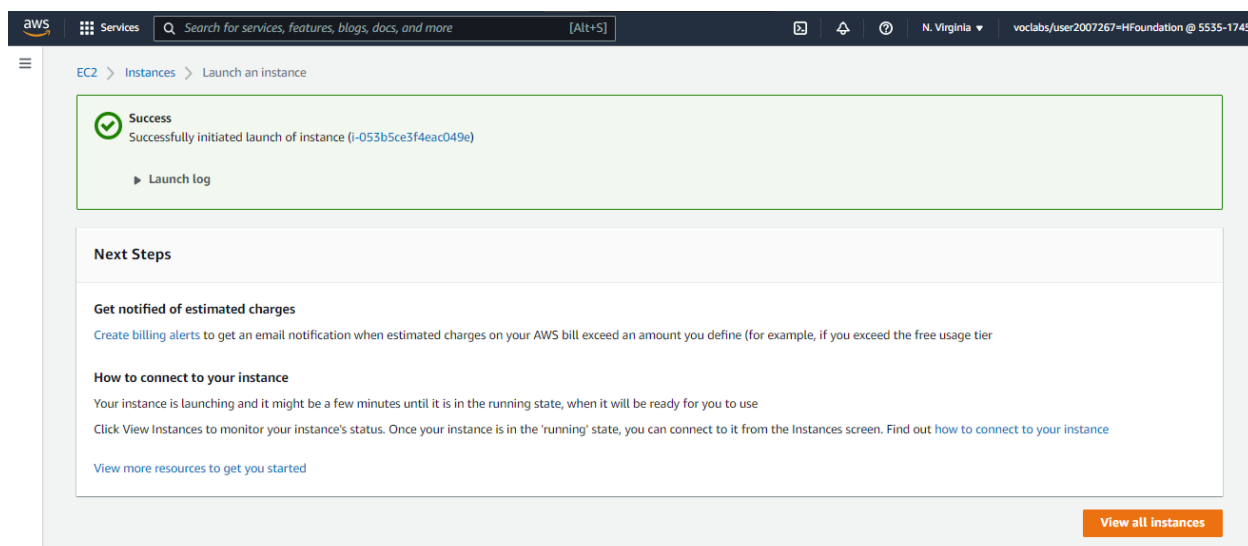
Firewall (security group)  
New security group

Storage (volumes)  
1 volume(s) - 30 GiB

Free tier: In your first year includes 750 hours of t2.micro (or t3.micro in the Regions in which t2.micro is unavailable) instance usage on free tier AMIs per month, 30 GiB of EBS storage, 2 million I/Os, 1 GB of snapshots, and 100 GB of bandwidth to the internet.

Cancel **Launch Instance**

- You'll need to wait for a while until the instance is created. Once finished, click on the “View all Instances” button.



aws Services Search for services, features, blogs, docs, and more [Alt+S] N. Virginia voclabs/user2007267-HFoundation @ 5535-1745-

EC2 > Instances > Launch an instance

**Success**  
Successfully initiated launch of instance (i-053b5ce3f4eac049e)  
[Launch log](#)

**Next Steps**

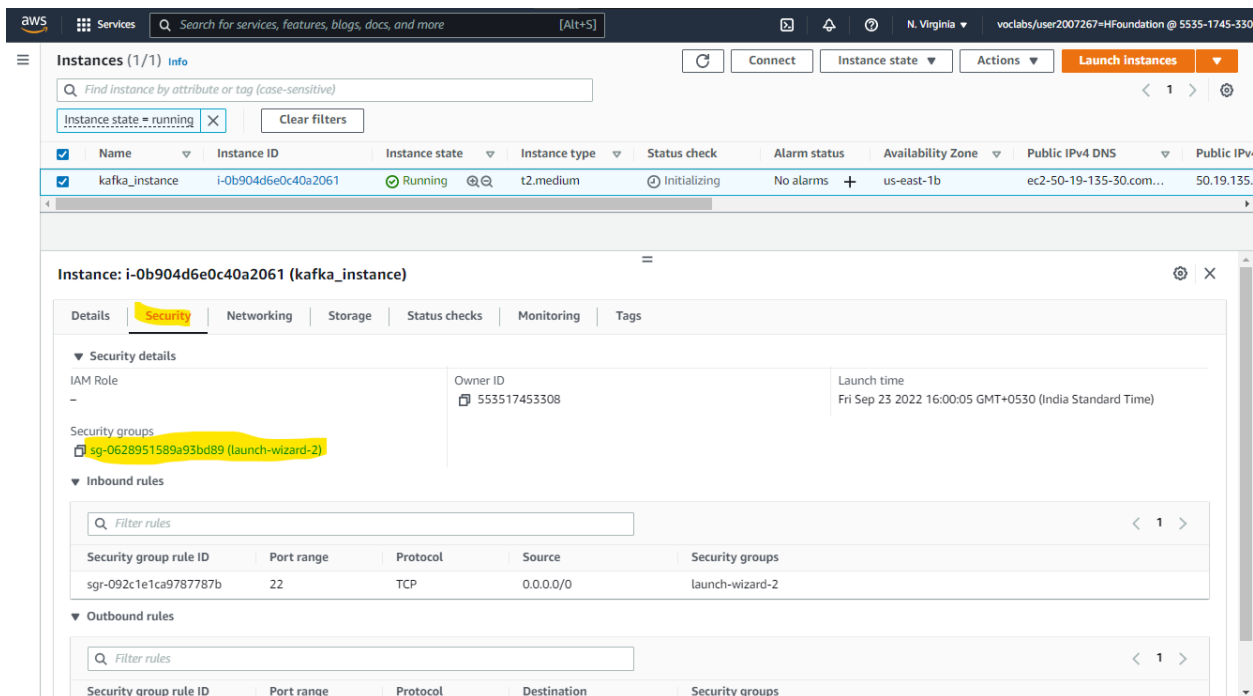
**Get notified of estimated charges**  
[Create billing alerts](#) to get an email notification when estimated charges on your AWS bill exceed an amount you define (for example, if you exceed the free usage tier)

**How to connect to your instance**  
Your instance is launching and it might be a few minutes until it is in the running state, when it will be ready for you to use  
Click [View Instances](#) to monitor your instance's status. Once your instance is in the 'running' state, you can connect to it from the Instances screen. Find out [how to connect to your instance](#)

[View more resources to get you started](#)

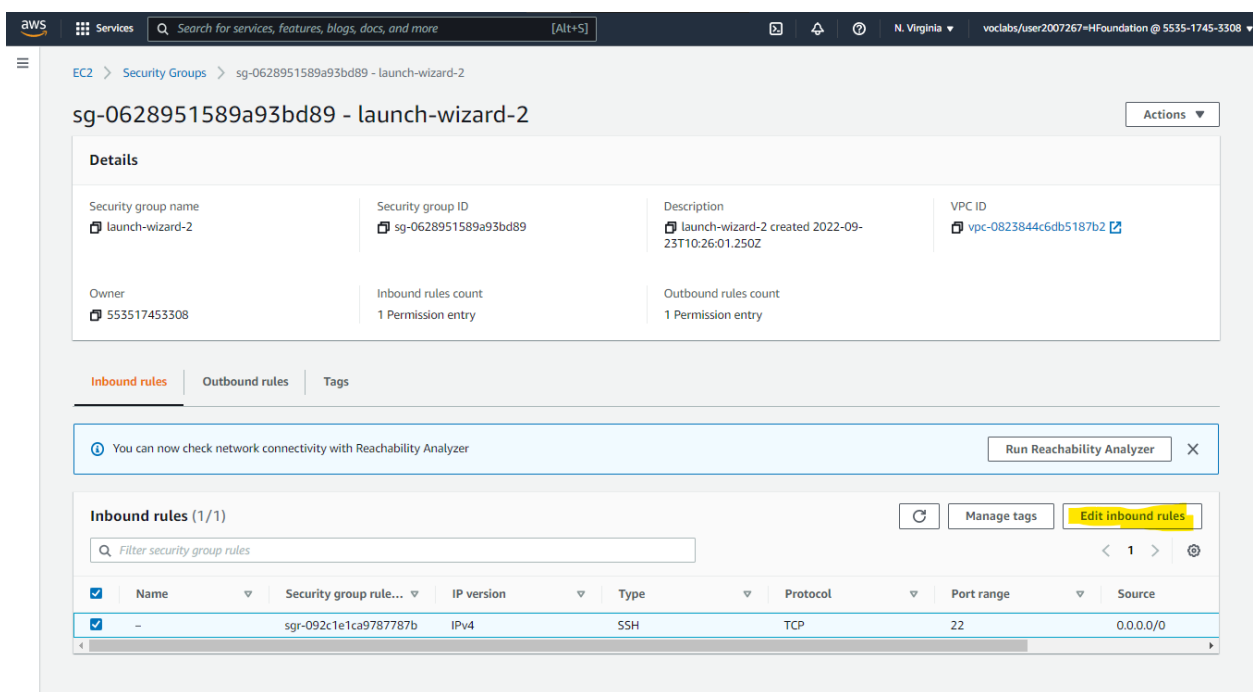
**View all instances**

- You'll be redirected to a new page where you can view all running EC2 instances. Click on the security tab and the security group ID.



The screenshot shows the AWS Management Console. At the top, there's a search bar and navigation tabs. The 'Instances' page is active, showing a list of instances. The 'kafka\_instance' is highlighted. Below the list, the details for 'Instance: i-0b904d6e0c40a2061 (kafka\_instance)' are shown. The 'Security' tab is selected, displaying the security group 'sg-0628951589a93bd89 (launch-wizard-2)'. The 'Inbound rules' section shows a single rule with ID 'sgr-092c1e1ca9787787b' on port 22, allowing TCP traffic from 0.0.0.0/0.

- As you can see, there's only one inbound rule in the security group. Click on “**Edit Inbound Rules**”



The screenshot shows the 'Edit inbound rules' page for the security group 'sg-0628951589a93bd89 - launch-wizard-2'. The 'Inbound rules' tab is selected, showing a table with one rule. The rule has ID 'sgr-092c1e1ca9787787b', is for IPv4, SSH, TCP, on port 22, with source 0.0.0.0/0. The 'Edit inbound rules' button is highlighted in yellow.

- Click on **Add Rule** button and configure the security group as shown below. Enter the following values as shown in the image below

**Type:** Custom TCP

**Port Range:** 8888

**Source:** Anywhere-IPv4

Edit inbound rules [Info](#)

Inbound rules control the incoming traffic that's allowed to reach the instance.

Security group rule ID	Type <a href="#">Info</a>	Protocol <a href="#">Info</a>	Port range <a href="#">Info</a>	Source <a href="#">Info</a>	Description - optional <a href="#">Info</a>	
sgr-092c1e1ca9787787b	SSH	TCP	22	Custom	0.0.0.0/0	Delete
-	Custom TCP	TCP	8888	Anywhere-IPv4	0.0.0.0/0	Delete

Add rule

Cancel Preview changes Save rules

Similarly, you need to add the following port numbers:

**2181, 8888, 9092, 9000, 8080**

Make sure that it looks like this:

Inbound rules [Info](#)

Security group rule ID	Type <a href="#">Info</a>	Protocol <a href="#">Info</a>	Port range <a href="#">Info</a>	Source <a href="#">Info</a>	Description - optional <a href="#">Info</a>	
sgr-0d0ec2fc8f22b8311	Custom TCP	TCP	9000	Anywhere-IPv4	0.0.0.0/0	Delete
sgr-092c1e1ca9787787b	SSH	TCP	22	Anywhere-IPv4	0.0.0.0/0	Delete
sgr-03235df0897c234c5	Custom TCP	TCP	8888	Anywhere-IPv4	0.0.0.0/0	Delete
sgr-05b023b607fd1d6a5	Custom TCP	TCP	9092	Anywhere-IPv4	0.0.0.0/0	Delete
-	Custom TCP	TCP	2181	Anywhere-IPv4	0.0.0.0/0	Delete
-	Custom TCP	TCP	8080	Anywhere-IPv4	0.0.0.0/0	Delete

Add rule

Cancel Preview changes Save rules

**Note: This is a very important step. Double-check that all the port range and source are identical as shown in the image.**

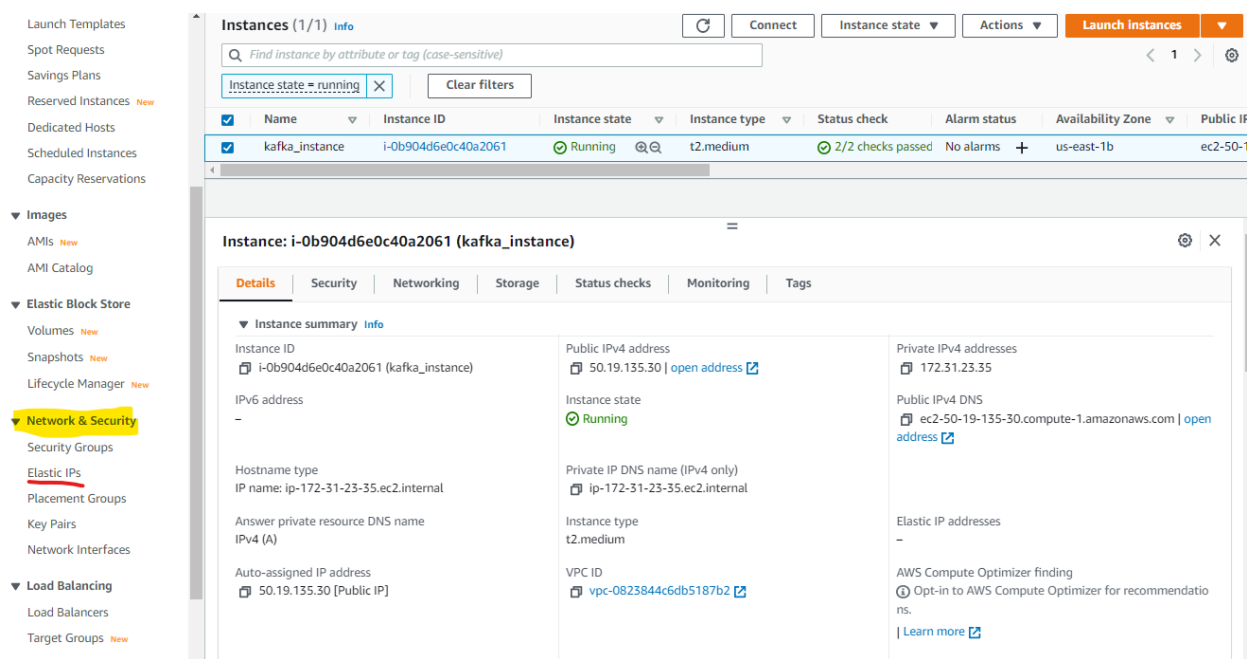
- Next click on the **Save rules**.

## ELASTIC IP CREATION INSTRUCTIONS

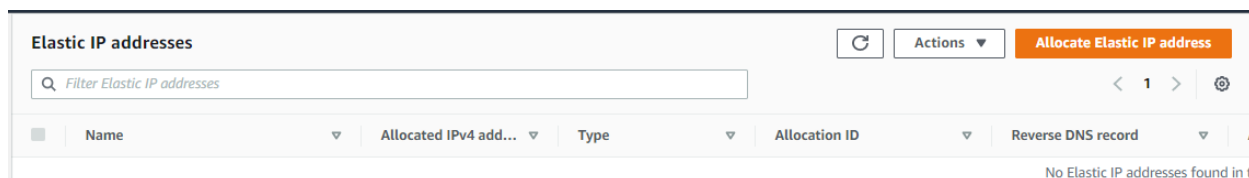
Whenever you start an EC2 instance, it generates a new IPv4 public IP. To keep the public IP constant, an elastic IP is required. A constant public IP is crucial in Kafka as you are required to create Kafka producers and consumers by writing code. You will understand the significance of having a constant IP when you will actually write the codes.

To create an elastic IP, do the following:

- Go to your EC2 instance page and on the left-hand side scroll till you find the **Network and Security** tab, on the left-hand side, under which you will see the **Elastic IPs** link, as shown in the image below.



- Click on the **Elastic IPs** link. You will be redirected to a new page. Click on the **Allocate Elastic IP address** button



Once you do so, you will be redirected to a new page as shown below. Choose the settings and click on the “Allocate” button.

## Allocate Elastic IP address [Info](#)

### Elastic IP address settings [Info](#)

Network Border Group [Info](#)

us-east-1

Public IPv4 address pool

- ☒ Amazon's pool of IPv4 addresses
- ☐ Public IPv4 address that you bring to your AWS account (option disabled because no pools found) [Learn more](#)
- ☐ Customer owned pool of IPv4 addresses (option disabled because no customer owned pools found) [Learn more](#)

Global static IP addresses

AWS Global Accelerator can provide global static IP addresses that are announced worldwide using anycast from AWS edge locations. This can help improve the availability and latency for your user traffic by using the Amazon global network. [Learn more](#)

[Create accelerator](#)

### Tags - optional

A tag is a label that you assign to an AWS resource. Each tag consists of a key and an optional value. You can use tags to search and filter your resources or track your AWS costs.


No tags associated with the resource.

[Add new tag](#)

You can add up to 50 more tag

[Cancel](#)
[Allocate](#)

- Click on **Allocate**. You will be redirected back to the elastic IP screen. Please note that only one Elastic IP address can't be allotted to multiple EC2 instances.

 Elastic IP address allocated successfully.  
Elastic IP address 52.21.15.133

[Associate this Elastic IP address](#)

### Elastic IP addresses (1/1)

Filter Elastic IP addresses

Public IPv4 address: 52.21.15.133 [Clear filters](#)

<input checked="" type="checkbox"/>	Name	Allocated IPv4 add...	Type	Allocation ID	Reverse DNS record	Associated instance ID
<input checked="" type="checkbox"/>	-	52.21.15.133	Public IP	eipalloc-046779924386e333c	-	-

52.21.15.133

[Summary](#) [Tags](#)

- Click on the checkbox against your newly created elastic IP. Then, click on the Actions button, you will get a drop-down list. From this list, click on **Associate Address**.



Elastic IP address allocated successfully.  
Elastic IP address 52.21.15.133

Associate this Elastic IP address

Elastic IP addresses (1/1)

Filter Elastic IP addresses

Public IPv4 address: 52.21.15.133

Clear filters

Actions

- View details
- Release Elastic IP addresses
- Associate Elastic IP address
- Disassociate Elastic IP address
- Update reverse DNS

	Name	Allocated IPv4 address	Type	Allocation ID	Reverse DNS	Instance ID
<input checked="" type="checkbox"/>	-	52.21.15.133	Public IP	eipalloc-046779924386e333c	-	-

- A new page will appear, as shown below. Click inside the Instance text box; you will get a list of available instances, as shown below. From the list, choose the EC2 instance with which you want to attach the elastic IP. If you have followed the steps in this document, then its name would be **kafka\_instance**. The Instance ID of that EC2 instance will appear in the text box, as shown in the example below.

## Associate Elastic IP address

Choose the instance or network interface to associate to this Elastic IP address (52.21.15.133)

**Elastic IP address: 52.21.15.133**

Resource type  
Choose the type of resource with which to associate the Elastic IP address.

☒ Instance

☐ Network interface

**Warning:** If you associate an Elastic IP address to an instance that already has an Elastic IP address associated, this previously associated Elastic IP address will be disassociated but still allocated to your account. [Learn more](#)

Instance

Choose an instance

i-0b904d6e0c40a2061 (kafka\_instance) - running

Private IP address

The private IP address with which to associate the Elastic IP address.

Choose a private IP address

Reassociation

Specify whether the Elastic IP address can be reassociated with a different resource if it already associated with a resource.

☐ Allow this Elastic IP address to be reassociated

Cancel Associate

- Next, click on the text-box next to **Private IP** and select the automatically generated **Private IP address** and then click on Associate.

## Associate Elastic IP address

Choose the instance or network interface to associate to this Elastic IP address (52.21.15.133)

**Elastic IP address: 52.21.15.133**

**Resource type**  
Choose the type of resource with which to associate the Elastic IP address.

☒ Instance  
☐ Network interface

**Warning:** If you associate an Elastic IP address to an instance that already has an Elastic IP address associated, this previously associated Elastic IP address will be disassociated but still allocated to your account. [Learn more](#)

**Instance**  
i-0b904d6e0c40a2061

**Private IP address**  
The private IP address with which to associate the Elastic IP address.

Choose a private IP address  
172.31.23.35

Specify whether the Elastic IP address can be reassocated with a different resource if it already associated with a resource.

☐ Allow this Elastic IP address to be reassocated

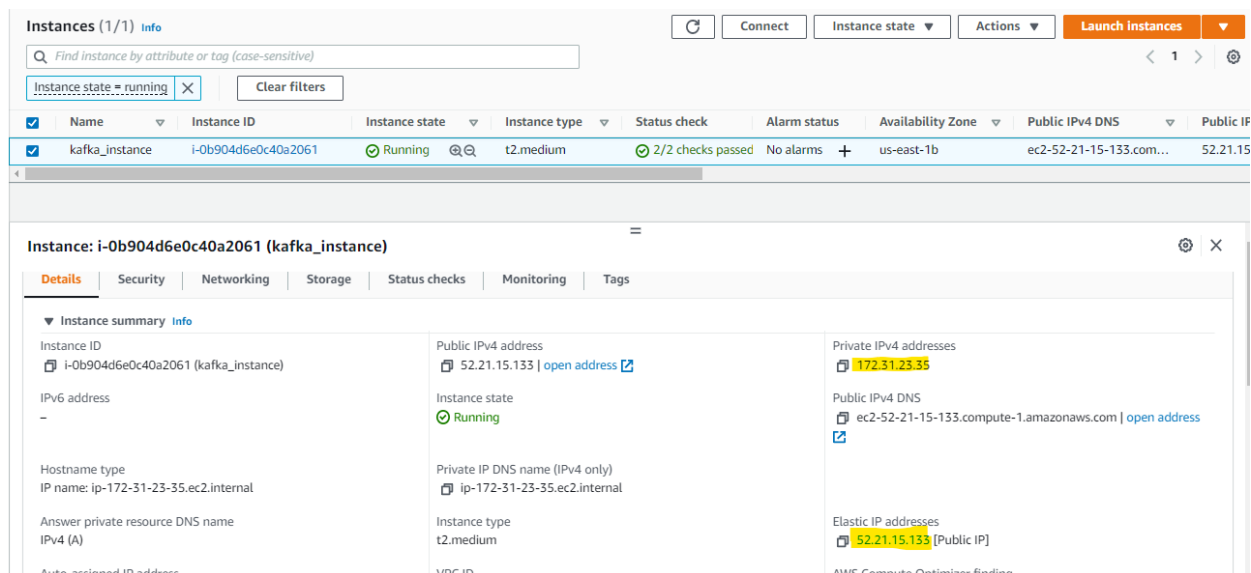
Cancel Associate

- You will be navigated to a new page where you will have confirmation that the elastic IP has been attached to your desired EC2 instance, as shown below and click on **Close**.

**Elastic IP address associated successfully.**  
Elastic IP address 52.21.15.133 has been associated with instance i-0b904d6e0c40a2061

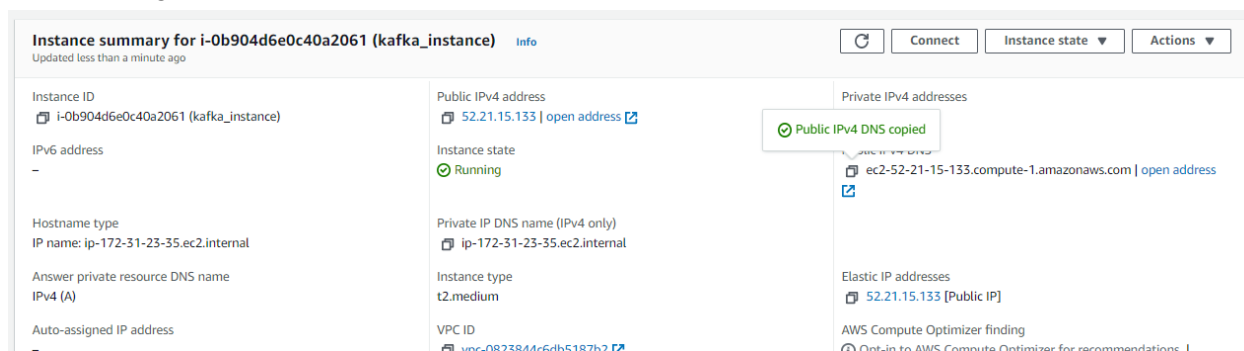
Name	Allocated IPv4 add...	Type	Allocation ID	Reverse DNS record
-	52.21.15.133	Public IP	eipalloc-046779924386e333c	-

- Now, go to the Instances page using the Scroll bar on the left-hand side of the screen and click on the tick box against the EC2 instance with which you associated the elastic IP. You will see that the elastic IP has been attached to the EC2 instance, as shown below.



The screenshot shows the AWS Management Console 'Instances' page. The instance 'kafka\_instance' (ID: i-0b904d6e0c40a2061) is in a 'Running' state. The 'Details' tab is selected, showing the instance summary. The 'Public IPv4 address' is 52.21.15.133, and the 'Private IPv4 address' is 172.31.23.35. The 'Elastic IP addresses' section shows the public IP 52.21.15.133 is attached to the instance.

- You can see that the elastic IP has been attached to the EC2 instance and your Private IPv4 address has been updated as well. Now that you have setup your EC2 instance, you will now need to SSH into the instance using the Public IPv4 DNS of the EC2 instance.



The screenshot shows the 'Instance summary' for 'kafka\_instance'. A notification bubble indicates 'Public IPv4 DNS copied'. The 'Public IPv4 address' is 52.21.15.133, and the 'Private IPv4 address' is 172.31.23.35. The 'Elastic IP addresses' section shows the public IP 52.21.15.133 is attached to the instance.

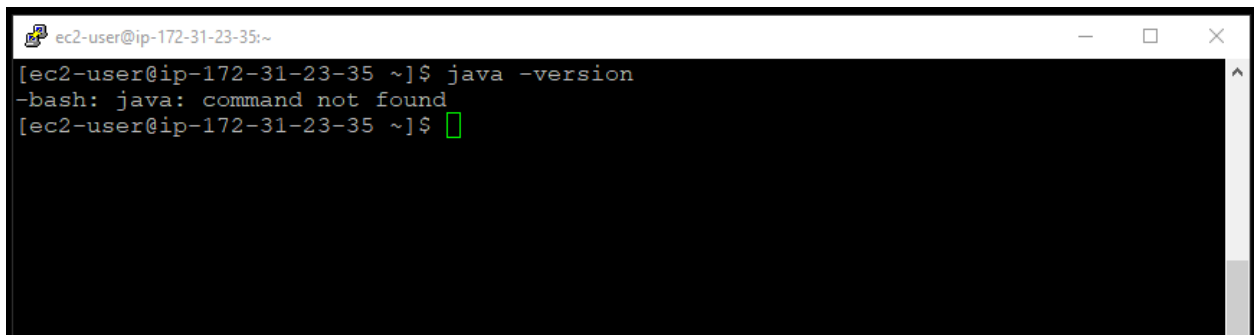
Now we'll instance all the required packages to work with Kafka. As mentioned in the beginning, the following services are required for working with the Kafka service in the module:

- Java 1.8
- Zookeeper
- Kafka
- Jupyter Notebook

Before you begin working, use the **sudo yum update** command to update the packages.

## UPDATING JAVA ENVIRONMENT

- Check the Java version by using the **java -version** command



```
ec2-user@ip-172-31-23-35:~  
[ec2-user@ip-172-31-23-35 ~]$ java -version  
-bash: java: command not found  
[ec2-user@ip-172-31-23-35 ~]$
```

- If you find the above error, it means that the java package has not been installed in the instance yet.

Use the following command to download the RPM package of Oracle JDK (8u121)

```
wget --no-check-certificate --no-cookies --header "Cookie:  
oraclelicense=accept-securebackup-cookie"  
http://download.oracle.com/otn-pub/java/jdk/8u141-b15/336fa29ff2bb4ef291e34  
7e091f7f4a7/jdk-8u141-linux-x64.rpm
```

```
ec2-user@ip-172-31-23-35:~
Location: https://edelivery.oracle.com/otn-pub/java/jdk/8u141-b15/336fa29ff2bb4ef291e347e091f7f4a7/jdk-8u141-linux-x64.rpm [following]
--2022-09-23 11:13:31-- https://edelivery.oracle.com/otn-pub/java/jdk/8u141-b15/336fa29ff2bb4ef291e347e091f7f4a7/jdk-8u141-linux-x64.rpm
Resolving edelivery.oracle.com (edelivery.oracle.com)... 104.68.240.79, 2600:1408:5400:4b1::366, 2600:1408:5400:482::366
Connecting to edelivery.oracle.com (edelivery.oracle.com)|104.68.240.79|:443... connected.
HTTP request sent, awaiting response... 302 Moved Temporarily
Location: https://download.oracle.com/otn-pub/java/jdk/8u141-b15/336fa29ff2bb4ef291e347e091f7f4a7/jdk-8u141-linux-x64.rpm?AuthParam=1663931731_2b80049fa94835d641ab479f17fe2985 [following]
--2022-09-23 11:13:31-- https://download.oracle.com/otn-pub/java/jdk/8u141-b15/336fa29ff2bb4ef291e347e091f7f4a7/jdk-8u141-linux-x64.rpm?AuthParam=1663931731_2b80049fa94835d641ab479f17fe2985
Connecting to download.oracle.com (download.oracle.com)|23.56.12.90|:443... connected.
HTTP request sent, awaiting response... 200 OK
Length: 169980729 (162M) [application/x-redhat-package-manager]
Saving to: 'jdk-8u141-linux-x64.rpm'

100%[=====>] 169,980,729 166MB/s in 1.0s

2022-09-23 11:13:32 (166 MB/s) - 'jdk-8u141-linux-x64.rpm' saved [169980729/169980729]

[ec2-user@ip-172-31-23-35 ~]$
```

**NOTE:** It is important to manually instruct wget to send the cookie and here this is done using **–header “Cookie: oraclelicense=accept-securebackup-cookie”** where oraclelicense is the name of the cookie and accept-securebackup-cookie is its value.

- Now install JDK 8 using the command below:

```
sudo yum install -y jdk-8u141-linux-x64.rpm
```

When you’ve successfully installed the Java Development Kit, you’ll see the command prompt as the image below

```

ec2-user@ip-172-31-23-35:~
=====
Install 1 Package

Total size: 269 M
Installed size: 269 M
Downloading packages:
Running transaction check
Running transaction test
Transaction test succeeded
Running transaction
  Installing : 2000:jdk1.8.0_141-1.8.0_141-fcs.x86_64 1/1
Unpacking JAR files...
  tools.jar...
  plugin.jar...
  javaws.jar...
  deploy.jar...
  rt.jar...
  jsse.jar...
  charsets.jar...
  localedata.jar...
Verifying : 2000:jdk1.8.0_141-1.8.0_141-fcs.x86_64 1/1

Installed:
  jdk1.8.0_141.x86_64 2000:1.8.0_141-fcs

Complete!
[ec2-user@ip-172-31-23-35 ~]$

```

- Verify oracle JDK version using the following command

```
java -version
```

If the JDK installation has been successful, you will see the image below

```

[ec2-user@ip-172-31-23-35 ~]$ java -version
java version "1.8.0_141"
Java(TM) SE Runtime Environment (build 1.8.0_141-b15)
Java HotSpot(TM) 64-Bit Server VM (build 25.141-b15, mixed mode)
[ec2-user@ip-172-31-23-35 ~]$

```

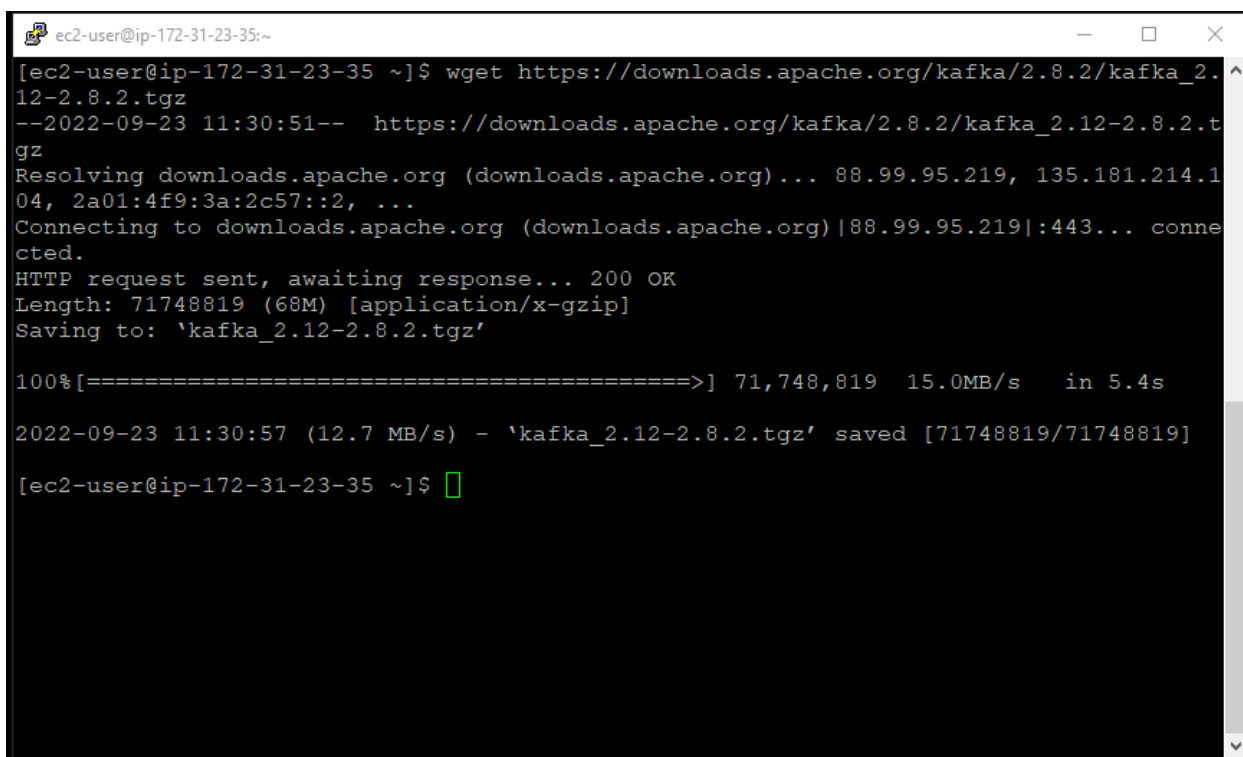
- The java version is now updated. After installing Java, follow the steps mentioned below in sequence to start Kafka on your instance.

## INSTALLING KAFKA

We'll now install kafka\_2.12-2.8.2 to the EC2 instance. Please note that in the module though we've installed a different version of kafka\_2.12-2.3.0, the commands for working with Kafka will still remain the same.

- Using the following wget command, download the Kafka package

```
wget https://downloads.apache.org/kafka/2.8.2/kafka_2.12-2.8.2.tgz
```



```
ec2-user@ip-172-31-23-35:~  
[ec2-user@ip-172-31-23-35 ~]$ wget https://downloads.apache.org/kafka/2.8.2/kafka_2.12-2.8.2.tgz  
--2022-09-23 11:30:51-- https://downloads.apache.org/kafka/2.8.2/kafka_2.12-2.8.2.tgz  
Resolving downloads.apache.org (downloads.apache.org)... 88.99.95.219, 135.181.214.104, 2a01:4f9:3a:2c57::2, ...  
Connecting to downloads.apache.org (downloads.apache.org)|88.99.95.219|:443... connected.  
HTTP request sent, awaiting response... 200 OK  
Length: 71748819 (68M) [application/x-gzip]  
Saving to: 'kafka_2.12-2.8.2.tgz'  
  
100%[=====>] 71,748,819 15.0MB/s in 5.4s  
2022-09-23 11:30:57 (12.7 MB/s) - 'kafka_2.12-2.8.2.tgz' saved [71748819/71748819]  
[ec2-user@ip-172-31-23-35 ~]$
```

- The kafka file will be saved to your instance as .tgz file. Now use the following command to extract the .tgz package

```
tar -xvzf kafka_2.12-2.8.2.tgz
```

- Since the original file is of no use, we can remove it using the following command:

```
rm kafka_2.12-2.8.2.tgz
```

Use the **ls** command to see the files in your EC2 instance.

```
ec2-user@ip-172-31-23-35:~
[ec2-user@ip-172-31-23-35 ~]$ ls
jdk-8u141-linux-x64.rpm  kafka_2.12-2.8.2
[ec2-user@ip-172-31-23-35 ~]$
```

## KAFKA SERVER CONFIGURATION

The following steps list the process to configure the Kafka Server so that you can connect your Kafka server through your SSH client or an IDE such as Eclipse.

To configure the Kafka server, you need to make changes to the `server.properties` file for which you need to go through the following steps:

- Login to the EC2 instance and go inside the Kafka directory present there. The command for the same is **cd /kafka\_2.12-2.8.2**. Once you are inside this directory, you need to go inside the config directory. The command for the same is **cd config/**. Once you are inside the directory, enter the command **ls**. This will list you the different files present inside the directory.

```
[ec2-user@ip-172-31-23-35 ~]$ cd kafka_2.12-2.8.2
[ec2-user@ip-172-31-23-35 kafka_2.12-2.8.2]$ ls
bin  config  libs  LICENSE  licenses  NOTICE  site-docs
[ec2-user@ip-172-31-23-35 kafka_2.12-2.8.2]$ cd config/
[ec2-user@ip-172-31-23-35 config]$ ls
connect-console-sink.properties  consumer.properties
connect-console-source.properties  kraft
connect-distributed.properties  log4j.properties
connect-file-sink.properties  producer.properties
connect-file-source.properties  server.properties
connect-log4j.properties  tools-log4j.properties
connect-mirror-maker.properties  trogdor.conf
connect-standalone.properties  zookeeper.properties
[ec2-user@ip-172-31-23-35 config]$
```

- Here you need to make changes to the **server.properties** file. To edit this file enter the command **vi server.properties**. You would get a screen, as shown below.



```
ec2-user@ip-172-31-23-35:~/kafka_2.12-2.8.2/config
# Licensed to the Apache Software Foundation (ASF) under one or more
# contributor license agreements. See the NOTICE file distributed with
# this work for additional information regarding copyright ownership.
# The ASF licenses this file to You under the Apache License, Version 2.0
# (the "License"); you may not use this file except in compliance with
# the License. You may obtain a copy of the License at
#
# http://www.apache.org/licenses/LICENSE-2.0
#
# Unless required by applicable law or agreed to in writing, software
# distributed under the License is distributed on an "AS IS" BASIS,
# WITHOUT WARRANTIES OR CONDITIONS OF ANY KIND, either express or implied.
# See the License for the specific language governing permissions and
# limitations under the License.
#
# see kafka.server.KafkaConfig for additional details and defaults
##### Server Basics #####
# The id of the broker. This must be set to a unique integer for each broker.
broker.id=0
##### Socket Server Settings #####
# The address the socket server listens on. It will get the value returned from
# java.net.InetAddress.getCanonicalHostName() if not configured.
"server.properties" 136L, 6849B 1,1 Top
```

- Here you need to make changes in the line which reads as follows:

```
#advertised.listeners=PLAINTEXT://your.host.name:9092
```

- You need to uncomment this line and in place of your.host.name you need to enter the IPv4 Public IP of your EC2 Instance. This is the same as the Elastic IP associated with your EC2 instance. So press i, and you will enter insert mode.

Instance: i-0b904d6e0c40a2061 (kafka\_instance)

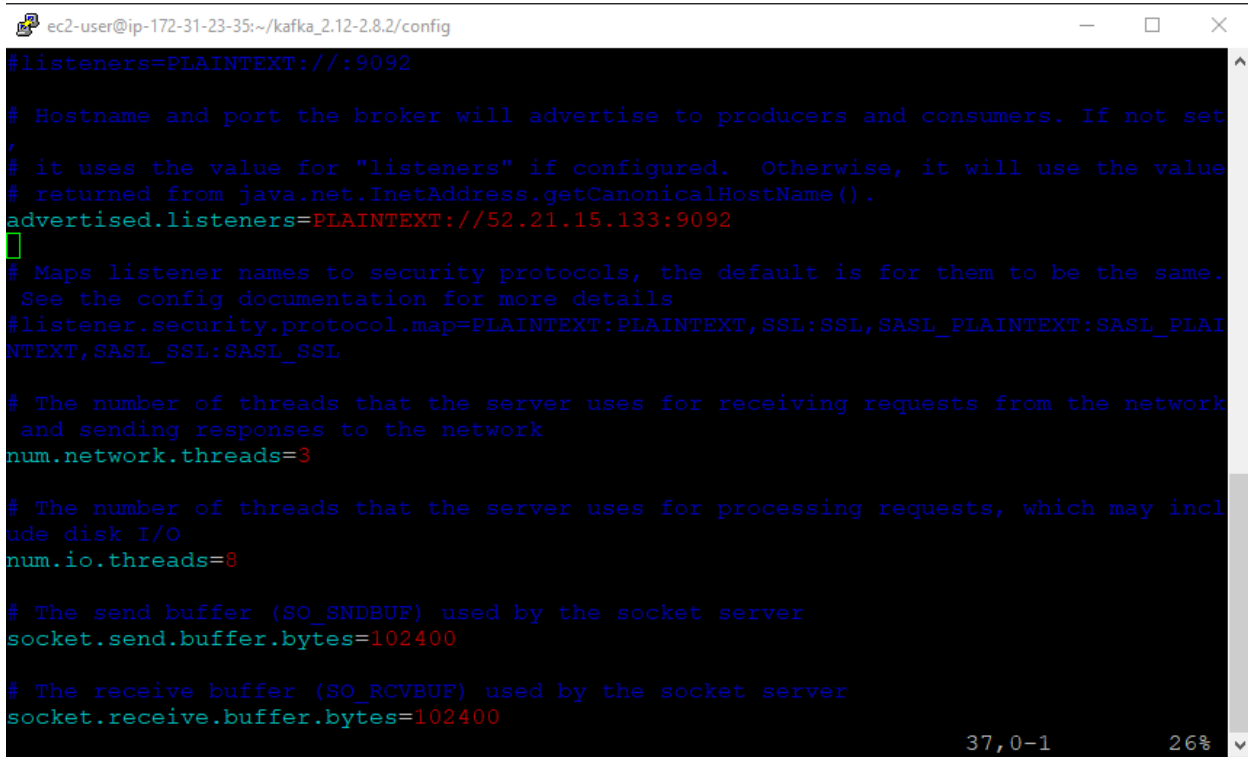
Details	Security	Networking	Storage	Status checks	Monitoring	Tags												
<p>▼ Instance summary info</p> <table border="1"> <tr> <td>Instance ID i-0b904d6e0c40a2061 (kafka_instance)</td> <td>Public IPv4 address 52.21.15.133   open address</td> <td>Private IPv4 addresses 172.31.23.35</td> </tr> <tr> <td>IPv6 address -</td> <td>Instance state Running</td> <td>Public IPv4 DNS ec2-52-21-15-133.compute-1.amazonaws.com   open address</td> </tr> <tr> <td>Hostname type IP name: ip-172-31-23-35.ec2.internal</td> <td>Private IP DNS name (IPv4 only) ip-172-31-23-35.ec2.internal</td> <td>Elastic IP addresses 52.21.15.133 [Public IP]</td> </tr> <tr> <td>Answer private resource DNS name IPv4 (A)</td> <td>Instance type t2.medium</td> <td></td> </tr> </table>							Instance ID i-0b904d6e0c40a2061 (kafka_instance)	Public IPv4 address 52.21.15.133   open address	Private IPv4 addresses 172.31.23.35	IPv6 address -	Instance state Running	Public IPv4 DNS ec2-52-21-15-133.compute-1.amazonaws.com   open address	Hostname type IP name: ip-172-31-23-35.ec2.internal	Private IP DNS name (IPv4 only) ip-172-31-23-35.ec2.internal	Elastic IP addresses 52.21.15.133 [Public IP]	Answer private resource DNS name IPv4 (A)	Instance type t2.medium	
Instance ID i-0b904d6e0c40a2061 (kafka_instance)	Public IPv4 address 52.21.15.133   open address	Private IPv4 addresses 172.31.23.35																
IPv6 address -	Instance state Running	Public IPv4 DNS ec2-52-21-15-133.compute-1.amazonaws.com   open address																
Hostname type IP name: ip-172-31-23-35.ec2.internal	Private IP DNS name (IPv4 only) ip-172-31-23-35.ec2.internal	Elastic IP addresses 52.21.15.133 [Public IP]																
Answer private resource DNS name IPv4 (A)	Instance type t2.medium																	

You can see the IP in my case is **52.21.15.133**

- Then go to the above line and uncomment it by removing the #.  
Next in place of your.host.name enter the IPv4 Public IP of your EC2 instance. For instance, the line would read as follows:

```
advertised.listeners=PLAINTEXT://52.21.15.133:9092
```

Your screen should look something as shown below:



```
ec2-user@ip-172-31-23-35:~/kafka_2.12-2.8.2/config
#listeners=PLAINTEXT://:9092

# Hostname and port the broker will advertise to producers and consumers. If not set
# it uses the value for "listeners" if configured. Otherwise, it will use the value
# returned from java.net.InetAddress.getCanonicalHostName().
advertised.listeners=PLAINTEXT://52.21.15.133:9092

# Maps listener names to security protocols, the default is for them to be the same.
# See the config documentation for more details
#listener.security.protocol.map=PLAINTEXT:PLAINTEXT,SSL:SSL,SASL_PLAINTEXT:SASL_PLAINTEXT,SASL_SSL:SASL_SSL

# The number of threads that the server uses for receiving requests from the network
# and sending responses to the network
num.network.threads=3

# The number of threads that the server uses for processing requests, which may include disk I/O
num.io.threads=8

# The send buffer (SO_SNDBUF) used by the socket server
socket.send.buffer.bytes=102400

# The receive buffer (SO_RCVBUF) used by the socket server
socket.receive.buffer.bytes=102400
```

**Make sure that you enter the IP address of your EC2 instance.**

Once you have done these steps, press **esc** and then type **:wq** to come out of the file.

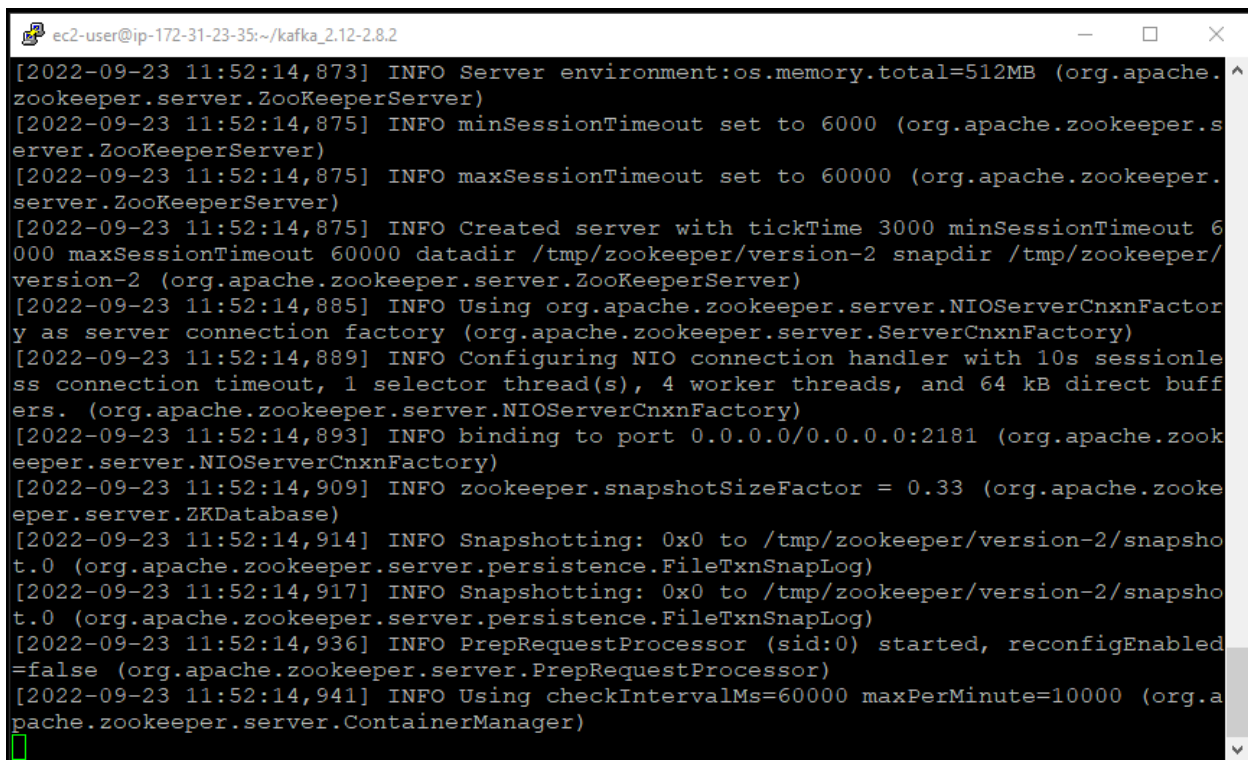
With these steps, you have configured the Kafka server to connect it with an IDE such as Eclipse.

## STEPS TO VERIFY THE INSTALLATION

### Zookeeper

To verify the zookeeper installation, follow the steps listed below.

- You need to get inside the Kafka directory. Go to the Kafka directory using the **cd kafka\_2.12-2.8.2/** command and then start the Zookeeper server using the **bin/zookeeper-server-start.sh config/zookeeper.properties** command. You should get the following output.

A terminal window titled 'ec2-user@ip-172-31-23-35:~/kafka\_2.12-2.8.2' displays the output of the 'bin/zookeeper-server-start.sh config/zookeeper.properties' command. The logs show various configuration parameters being set, such as 'os.memory.total=512MB', 'minSessionTimeout set to 6000', 'maxSessionTimeout set to 60000', and 'tickTime 3000'. The final line of the visible log is 'INFO binding to port 0.0.0.0/0.0.0.0:2181 (org.apache.zookeeper.server.NIOServerCnxnFactory)', which is highlighted in green in the original image.

```
[2022-09-23 11:52:14,873] INFO Server environment:os.memory.total=512MB (org.apache.zookeeper.server.ZooKeeperServer)
[2022-09-23 11:52:14,875] INFO minSessionTimeout set to 6000 (org.apache.zookeeper.server.ZooKeeperServer)
[2022-09-23 11:52:14,875] INFO maxSessionTimeout set to 60000 (org.apache.zookeeper.server.ZooKeeperServer)
[2022-09-23 11:52:14,875] INFO Created server with tickTime 3000 minSessionTimeout 6000 maxSessionTimeout 60000 datadir /tmp/zookeeper/version-2 snapdir /tmp/zookeeper/version-2 (org.apache.zookeeper.server.ZooKeeperServer)
[2022-09-23 11:52:14,885] INFO Using org.apache.zookeeper.server.NIOServerCnxnFactory as server connection factory (org.apache.zookeeper.server.ServerCnxnFactory)
[2022-09-23 11:52:14,889] INFO Configuring NIO connection handler with 10s sessionless connection timeout, 1 selector thread(s), 4 worker threads, and 64 kB direct buffers. (org.apache.zookeeper.server.NIOServerCnxnFactory)
[2022-09-23 11:52:14,893] INFO binding to port 0.0.0.0/0.0.0.0:2181 (org.apache.zookeeper.server.NIOServerCnxnFactory)
[2022-09-23 11:52:14,909] INFO zookeeper.snapshotSizeFactor = 0.33 (org.apache.zookeeper.server.ZKDatabase)
[2022-09-23 11:52:14,914] INFO Snapshotting: 0x0 to /tmp/zookeeper/version-2/snapshot.0 (org.apache.zookeeper.server.persistence.FileTxnSnapLog)
[2022-09-23 11:52:14,917] INFO Snapshotting: 0x0 to /tmp/zookeeper/version-2/snapshot.0 (org.apache.zookeeper.server.persistence.FileTxnSnapLog)
[2022-09-23 11:52:14,936] INFO PrepRequestProcessor (sid:0) started, reconfigEnabled=false (org.apache.zookeeper.server.PrepRequestProcessor)
[2022-09-23 11:52:14,941] INFO Using checkIntervalMs=60000 maxPerMinute=10000 (org.apache.zookeeper.server.ContainerManager)
```

At the bottom of the screen, you would get something like:

**INFO binding to port 0.0.0.0/0.0.0.0:2181**  
**(org.apache.zookeeper.server.NIOServerCnxnFactory)**

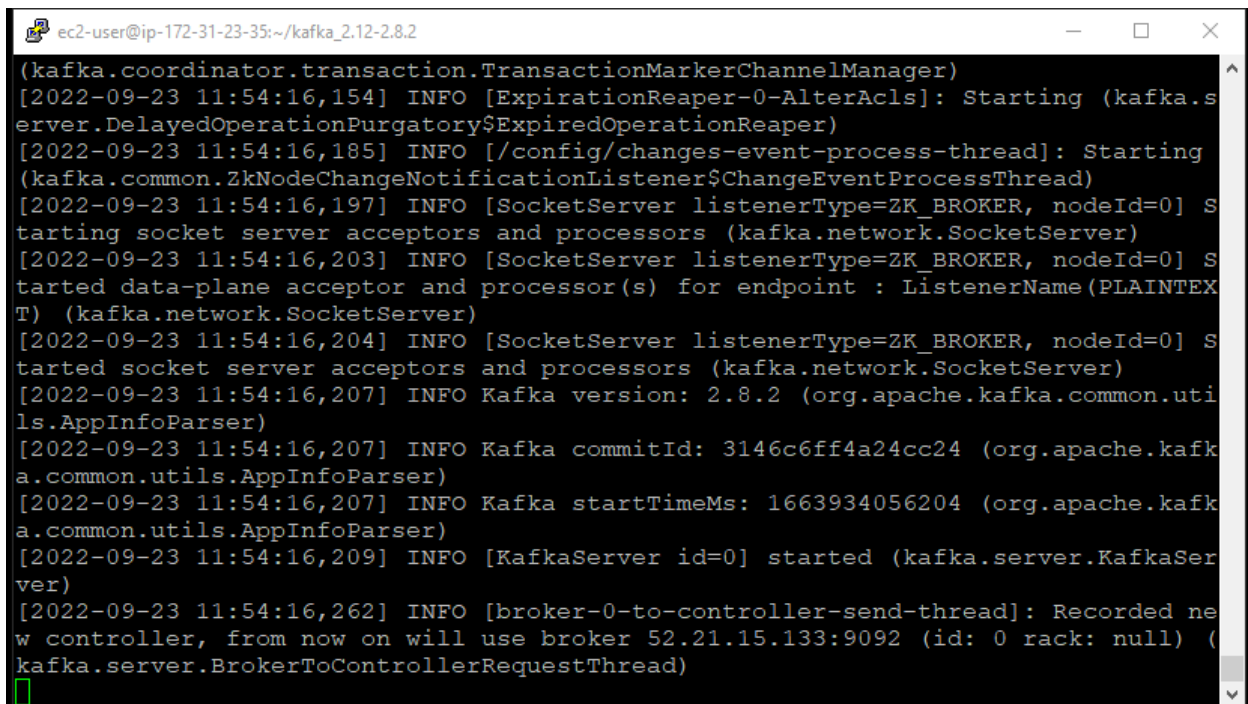
This means that your zookeeper server has successfully started.

## Kafka

Before going through this step, please ensure that the Zookeeper server is running.

To verify the Kafka installation, follow the steps listed below:

- Leave the previous terminal window as it is and login to your EC2 instance using another terminal.
- Go to the Kafka directory using the **cd downloads/kafka\_2.12-2.8.2** command.
- Start the Kafka server using the **bin/kafka-server-start.sh config/server.properties** command.
- You should get an output which displays a message something like “**INFO [KafkaServer id=0] started (kafka.server.KafkaServer)**”

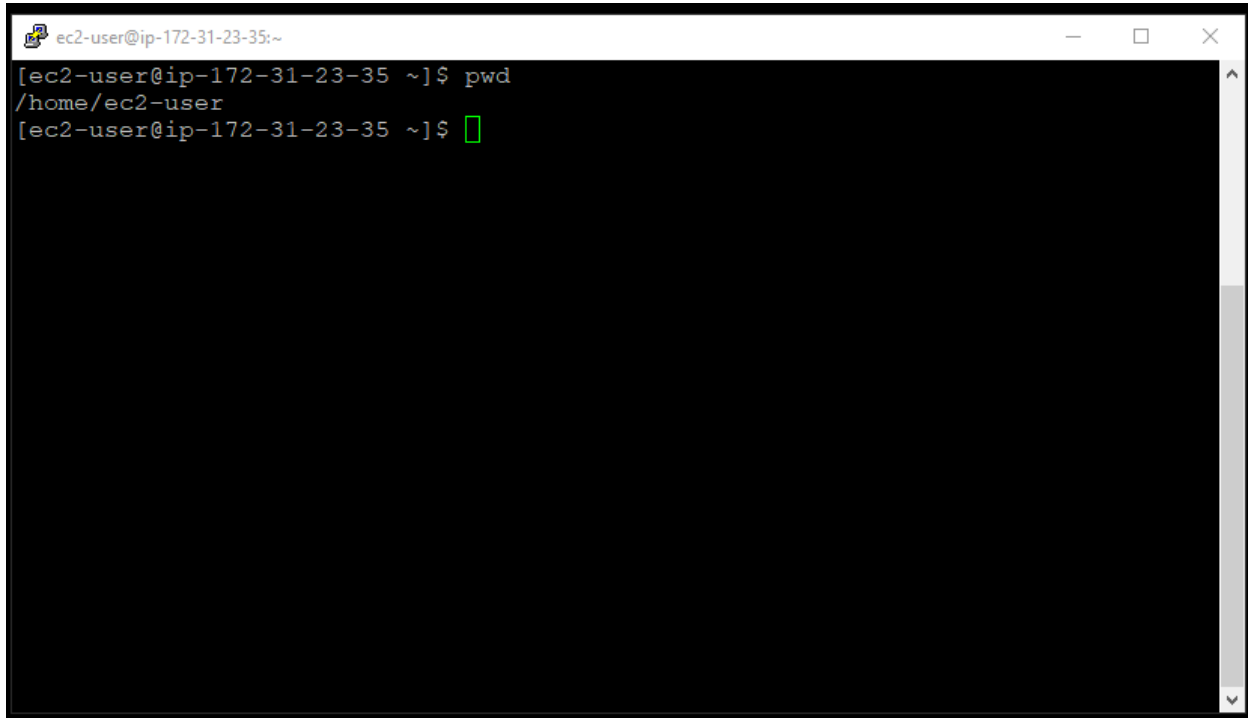


```
ec2-user@ip-172-31-23-35:~/kafka_2.12-2.8.2
(kafka.coordinator.transaction.TransactionMarkerChannelManager)
[2022-09-23 11:54:16,154] INFO [ExpirationReaper-0-AlterAcls]: Starting (kafka.s
server.DelayedOperationPurgatory$ExpiredOperationReaper)
[2022-09-23 11:54:16,185] INFO [/config/changes-event-process-thread]: Starting
(kafka.common.ZkNodeChangeNotificationListener$ChangeEventProcessThread)
[2022-09-23 11:54:16,197] INFO [SocketServer listenerType=ZK_BROKER, nodeId=0] S
tarting socket server acceptors and processors (kafka.network.SocketServer)
[2022-09-23 11:54:16,203] INFO [SocketServer listenerType=ZK_BROKER, nodeId=0] S
tarted data-plane acceptor and processor(s) for endpoint : ListenerName(PLAINTEX
T) (kafka.network.SocketServer)
[2022-09-23 11:54:16,204] INFO [SocketServer listenerType=ZK_BROKER, nodeId=0] S
tarted socket server acceptors and processors (kafka.network.SocketServer)
[2022-09-23 11:54:16,207] INFO Kafka version: 2.8.2 (org.apache.kafka.common.uti
ls.AppInfoParser)
[2022-09-23 11:54:16,207] INFO Kafka commitId: 3146c6ff4a24cc24 (org.apache.kafk
a.common.utils.AppInfoParser)
[2022-09-23 11:54:16,207] INFO Kafka startTimeMs: 1663934056204 (org.apache.kafk
a.common.utils.AppInfoParser)
[2022-09-23 11:54:16,209] INFO [KafkaServer id=0] started (kafka.server.KafkaSer
ver)
[2022-09-23 11:54:16,262] INFO [broker-0-to-controller-send-thread]: Recorded ne
w controller, from now on will use broker 52.21.15.133:9092 (id: 0 rack: null) (
kafka.server.BrokerToControllerRequestThread)
█
```

This means the instance creation is successful, and you can continue with the next steps.

## INSTALLING JUPYTER

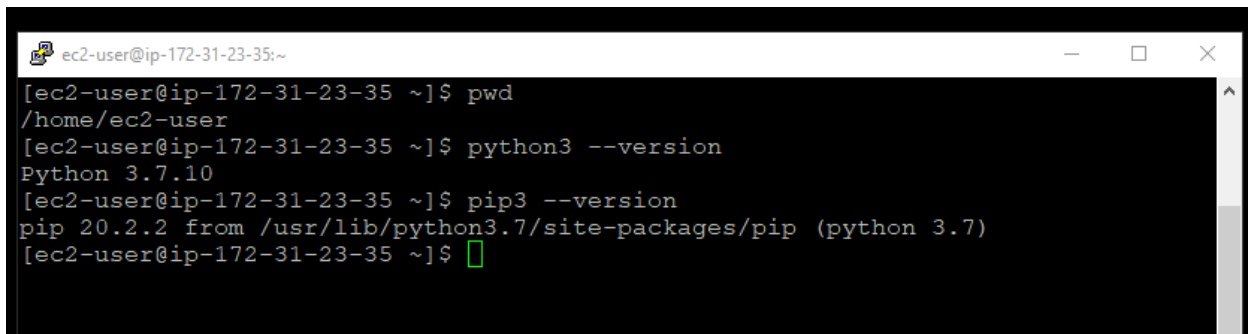
1. Open a new Terminal window and log in to the EC2 instance.  
Make sure that the current directory is **/home/ec2-user** by running the **pwd** command.  
Once you are inside the **/home/ec2-user** directory, follow the steps listed below.



```
ec2-user@ip-172-31-23-35:~  
[ec2-user@ip-172-31-23-35 ~]$ pwd  
/home/ec2-user  
[ec2-user@ip-172-31-23-35 ~]$
```

2. Once you have logged into the EC2 instance, the first step is to check for the prerequisites. You will require the pip3 package to install the Boto3 package. You can run the following commands:

```
python3 --version  
pip3 --version
```



```
ec2-user@ip-172-31-23-35:~  
[ec2-user@ip-172-31-23-35 ~]$ pwd  
/home/ec2-user  
[ec2-user@ip-172-31-23-35 ~]$ python3 --version  
Python 3.7.10  
[ec2-user@ip-172-31-23-35 ~]$ pip3 --version  
pip 20.2.2 from /usr/lib/python3.7/site-packages/pip (python 3.7)  
[ec2-user@ip-172-31-23-35 ~]$
```

If the packages are installed, you can skip the **3rd and 4th** steps and jump to the 5th step.

- As you can see in the image above, you must first install the Python3 and pip3 package. To install Python, you must run the following commands:

**sudo yum -y update**

This command will help you to update the Linux repository

Then to install python use

**sudo yum install python36**

This command will install the Python 3.6 on your instance

**Note:** In case you receive an error, you can run the following command:

**sudo yum -y install python37**

- Once you have installed Python, the next step is to install the pip package corresponding to Python3

**curl -O https://bootstrap.pypa.io/get-pip.py**

```
[ec2-user@ip-172-31-23-35 ~]$ curl -O https://bootstrap.pypa.io/get-pip.py
  % Total    % Received % Xferd  Average Speed   Time    Time     Time  Current
                                 Dload  Upload   Total   Spent    Left   Speed
100 2500k  100 2500k    0     0  98.8M      0  --:--:-- --:--:-- --:--:-- 101M
[ec2-user@ip-172-31-23-35 ~]$
```

This command will install the pip corresponding to Python3 on your instance.

**sudo python3 get-pip.py**

```
[ec2-user@ip-172-31-22-81 ~]$ sudo python3 get-pip.py
Collecting pip
  Downloading pip-20.3.3-py2.py3-none-any.whl (1.5 MB)
    | 1.5 MB 17.2 MB/s
Collecting wheel
  Downloading wheel-0.36.2-py2.py3-none-any.whl (35 kB)
Installing collected packages: pip, wheel
Attempting uninstall: pip
  Found existing installation: pip 9.0.3
  Uninstalling pip-9.0.3:
    Successfully uninstalled pip-9.0.3
WARNING: The scripts pip, pip3 and pip3.7 are installed in '/usr/local/bin' which is not on PATH.
Consider adding this directory to PATH or, if you prefer to suppress this warning, use --no-warn-script-location.
WARNING: The script wheel is installed in '/usr/local/bin' which is not on PATH.
Consider adding this directory to PATH or, if you prefer to suppress this warning, use --no-warn-script-location.
Successfully installed pip-20.3.3 wheel-0.36.2
[ec2-user@ip-172-31-22-81 ~]$
```

- Now, you have the required packages installed. However, there is one check required at this step. You must be able to call the functions associated with pip directly. You must not specify the complete path of the package to run the commands.

To check this, run the following command:

**pip3 --version**

```
[ec2-user@ip-172-31-23-35 ~]$ pip3 --version
pip 20.2.2 from /usr/lib/python3.7/site-packages/pip (python 3.7)
[ec2-user@ip-172-31-23-35 ~]$
```

You must follow the steps discussed ahead to make sure that you don't receive an error

- Run this command to open the bash file:

```
vi .bashrc
```

To edit the file, press the 'i' key.

- You are expected to add the following command at the end of the file.

```
export PATH="/usr/local/bin/:$PATH"
```

```
ec2-user@ip-172-31-23-35:~  
# .bashrc  
  
# Source global definitions  
if [ -f /etc/bashrc ]; then  
    . /etc/bashrc  
fi  
  
# Uncomment the following line if you don't like systemctl's auto-paging feature  
:  
# export SYSTEMD_PAGER=  
  
# User specific aliases and functions  
export PATH="/usr/local/bin/:$PATH"  
~  
~  
~  
~  
~  
~  
~  
~  
~
```

12,35 All

- Do not make any other changes in the file. Once you have added this, press Esc and type the command **:wq!** and press enter to exit the bash file.
- You must finally execute the file to load the changes using the following command:  
**source .bashrc**
- Now verify using the command

**pip3 -V**

```
[ec2-user@ip-172-31-23-35 ~]$ pip3 -V
pip 20.2.2 from /usr/lib/python3.7/site-packages/pip (python 3.7)
[ec2-user@ip-172-31-23-35 ~]$
```

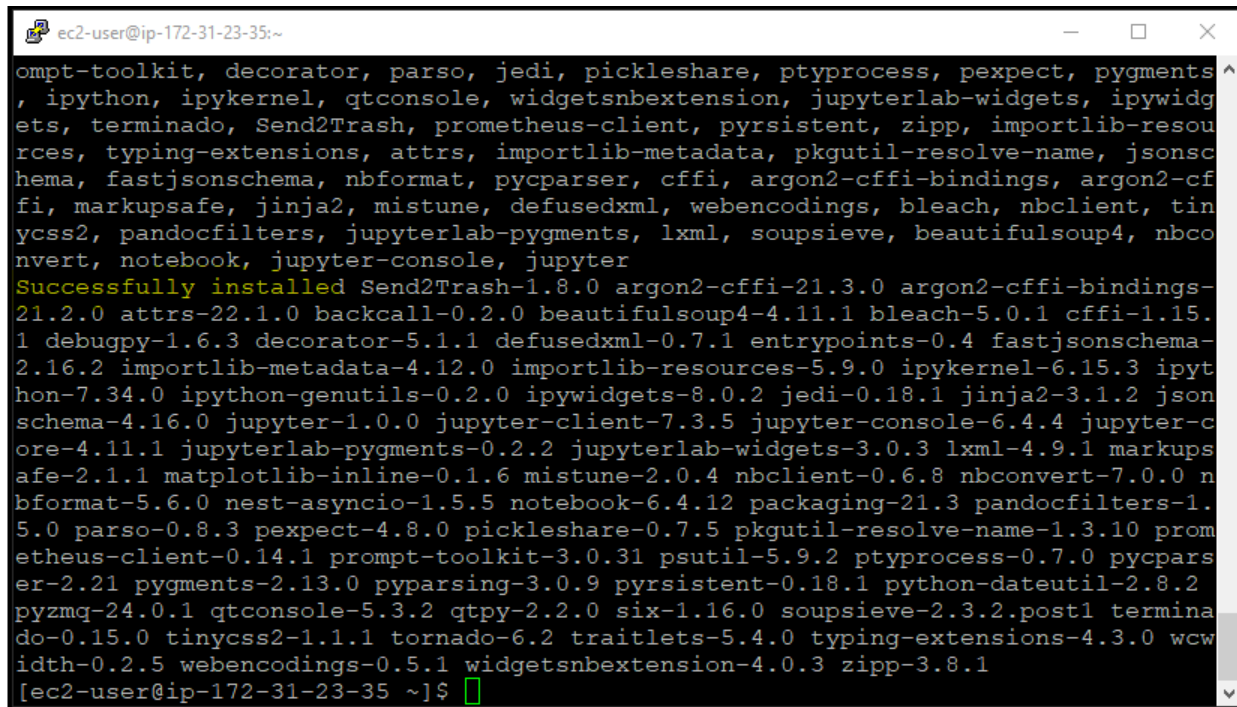
You are now ready with all the prerequisites to install the Jupyter library in the EC2 instance.

# INSTALLING JUPYTER ON EC2 INSTANCE

1. You can run the following command on the EC2 terminal to install the jupyter server:

**pip3 install jupyter --user**

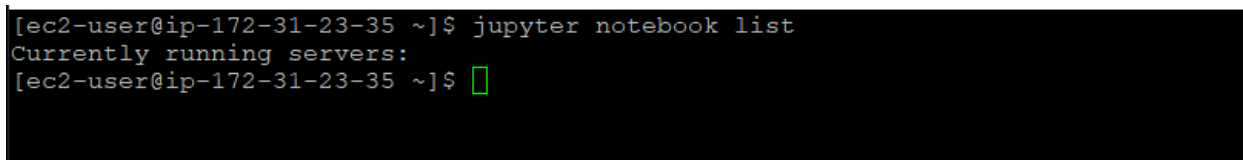
You'll get a "Successfully installed" message in your terminal.



```
ec2-user@ip-172-31-23-35:~$ pip3 install jupyter --user
ompt-toolkit, decorator, parso, jedi, pickleshare, ptyprocess, pexpect, pygments
, ipython, ipykernel, qtconsole, widgetsnbextension, jupyterlab-widgets, ipywidg
ets, terminado, Send2Trash, prometheus-client, pyrsistent, zipp, importlib-resou
rces, typing-extensions, attrs, importlib-metadata, pkgutil-resolve-name, jsonsc
hema, fastjsonschema, nbformat, pyparser, cffi, argon2-cffi-bindings, argon2-cf
fi, markupsafe, jinja2, mistune, defusedxml, webencodings, bleach, nbclient, tin
ycss2, pandocfilters, jupyterlab-pygments, lxml, soupsieve, BeautifulSoup4, nbco
nvert, notebook, jupyter-console, jupyter
Successfully installed Send2Trash-1.8.0 argon2-cffi-21.3.0 argon2-cffi-bindings-
21.2.0 attrs-22.1.0 backcall-0.2.0 BeautifulSoup4-4.11.1 bleach-5.0.1 cffi-1.15.
1 debugpy-1.6.3 decorator-5.1.1 defusedxml-0.7.1 entrypoints-0.4 fastjsonschema-
2.16.2 importlib-metadata-4.12.0 importlib-resources-5.9.0 ipykernel-6.15.3 ipyt
hon-7.34.0 ipython-genutils-0.2.0 ipywidgets-8.0.2 jedi-0.18.1 jinja2-3.1.2 json
schema-4.16.0 jupyter-1.0.0 jupyter-client-7.3.5 jupyter-console-6.4.4 jupyter-c
ore-4.11.1 jupyterlab-pygments-0.2.2 jupyterlab-widgets-3.0.3 lxml-4.9.1 markups
afe-2.1.1 matplotlib-inline-0.1.6 mistune-2.0.4 nbclient-0.6.8 nbconvert-7.0.0 n
bformat-5.6.0 nest-asyncio-1.5.5 notebook-6.4.12 packaging-21.3 pandocfilters-1.
5.0 parso-0.8.3 pexpect-4.8.0 pickleshare-0.7.5 pkgutil-resolve-name-1.3.10 prom
etheus-client-0.14.1 prompt-toolkit-3.0.31 psutil-5.9.2 ptyprocess-0.7.0 pypars
er-2.21 pygments-2.13.0 pyparsing-3.0.9 pyrsistent-0.18.1 python-dateutil-2.8.2
pyzmq-24.0.1 qtconsole-5.3.2 qtpy-2.2.0 six-1.16.0 soupsieve-2.3.2.post1 termina
do-0.15.0 tinycss2-1.1.1 tornado-6.2 traitlets-5.4.0 typing-extensions-4.3.0 wcw
idth-0.2.5 webencodings-0.5.1 widgetsnbextension-4.0.3 zipp-3.8.1
[ec2-user@ip-172-31-23-35 ~]$
```

2. Verify jupyter server status using below command:

**jupyter notebook list**



```
[ec2-user@ip-172-31-23-35 ~]$ jupyter notebook list
Currently running servers:
[ec2-user@ip-172-31-23-35 ~]$
```

Currently, there are no running jupyter servers.

Alternatively, you can view the list of installed libraries by using the following command:

**pip3 list**



```
ec2-user@ip-172-31-23-35:~
importlib-resources 5.9.0
ipykernel            6.15.3
ipython              7.34.0
ipython-genutils    0.2.0
ipywidgets           8.0.2
jedi                 0.18.1
Jinja2               3.1.2
jsonschema           4.16.0
jupyter              1.0.0
jupyter-client       7.3.5
jupyter-console      6.4.4
jupyter-core         4.11.1
jupyterlab-pygments  0.2.2
jupyterlab-widgets   3.0.3
lockfile             0.11.0
lxml                 4.9.1
MarkupSafe           2.1.1
matplotlib-inline    0.1.6
mistune              2.0.4
nbclient             0.6.8
nbconvert            7.0.0
nbformat             5.6.0
nest-asyncio         1.5.5
notebook             6.4.12
```

You can see that the jupyter library has been installed with its necessary components

3. You can start the Jupyter Notebook server using the command provided below:

### jupyter notebook

```
ec2-user@ip-172-31-23-35:~
[I 12:20:26.898 NotebookApp] Writing notebook server cookie secret to /home/ec2-
user/.local/share/jupyter/runtime/notebook_cookie_secret
[I 12:20:28.111 NotebookApp] Serving notebooks from local directory: /home/ec2-u
ser
[I 12:20:28.111 NotebookApp] Jupyter Notebook 6.4.12 is running at:
[I 12:20:28.111 NotebookApp] http://localhost:8888/?token=9736eab9490350ab3af01a
811126697ead2bb40027269f9d
[I 12:20:28.111 NotebookApp] or http://127.0.0.1:8888/?token=9736eab9490350ab3a
f01a811126697ead2bb40027269f9d
[I 12:20:28.112 NotebookApp] Use Control-C to stop this server and shut down all
kernels (twice to skip confirmation).
[W 12:20:28.116 NotebookApp] No web browser found: could not locate runnable bro
wser.
[C 12:20:28.116 NotebookApp]

To access the notebook, open this file in a browser:
file:///home/ec2-user/.local/share/jupyter/runtime/nbserver-4760-open.ht
ml
Or copy and paste one of these URLs:
http://localhost:8888/?token=9736eab9490350ab3af01a811126697ead2bb400272
69f9d
or http://127.0.0.1:8888/?token=9736eab9490350ab3af01a811126697ead2bb400272
69f9d
```

**NOTE:** Now, the Jupyter Notebook is running on the EC2 instance. However, you need to access it through a browser using SSH tunneling. This step differs for Windows and Linux/Mac users.

## **CONNECTING TO A JUPYTER NOTEBOOK - LINUX/MAC**

1. After you have launched the Jupyter notebook server, you must open a new terminal window and run the following commands:

```
ssh -i "keypair1.pem" -N -f -L 8888:localhost:8888 ec2-user@IPv4_address_of_EC2_instance
```

2. Next, you can select the url from the previous terminal window and paste to the local browser. You should not use **Ctrl + C** to copy as it raises the prompt to end the session. The text is automatically copied by selecting it in the window.

```
[ec2-user@ip-172-31-22-81 ~]$ jupyter notebook
[I 06:33:26.140 NotebookApp] Serving notebooks from local directory: /home/ec2-user
[I 06:33:26.140 NotebookApp] Jupyter Notebook 6.2.0 is running at:
[I 06:33:26.140 NotebookApp] http://localhost:8888/?token=297879d6ab64b1498601afe773244a055e20db11edf8bd59
[I 06:33:26.140 NotebookApp] or http://127.0.0.1:8888/?token=297879d6ab64b1498601afe773244a055e20db11edf8bd59
[I 06:33:26.140 NotebookApp] Use Control-C to stop this server and shut down all kernels (t
[W 06:33:26.144 NotebookApp] No web browser found: could not locate runnable browser.
[C 06:33:26.145 NotebookApp]

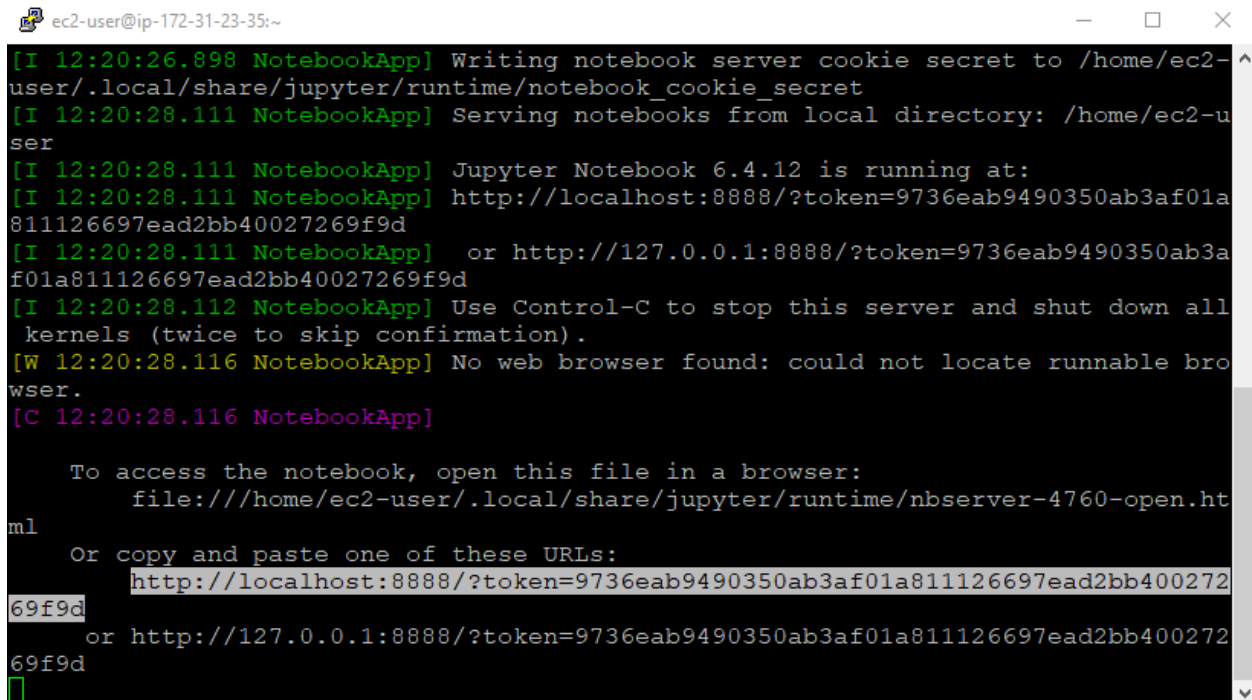
To access the notebook, open this file in a browser:
    file:///home/ec2-user/.local/share/jupyter/runtime/nbserver-32684-open.html
Or copy and paste one of these URLs:
    http://localhost:8888/?token=297879d6ab64b1498601afe773244a055e20db11edf8bd59
    or http://127.0.0.1:8888/?token=297879d6ab64b1498601afe773244a055e20db11edf8bd59
```

The Jupyter notebook will be launched in the browser

## CONNECTING TO A JUPYTER NOTEBOOK - WINDOWS


This is a step by step approach to connect to a Jupyter Notebook once it has been installed in the EC2 instance. For this task, you will require the PuTTY software when working with the Windows machine.

1. You can select the url from the PuTTY window and paste in the local browser. You should not use Cont + C to copy as it raises the prompt to end the session. The text is automatically copied by selecting it in the window.



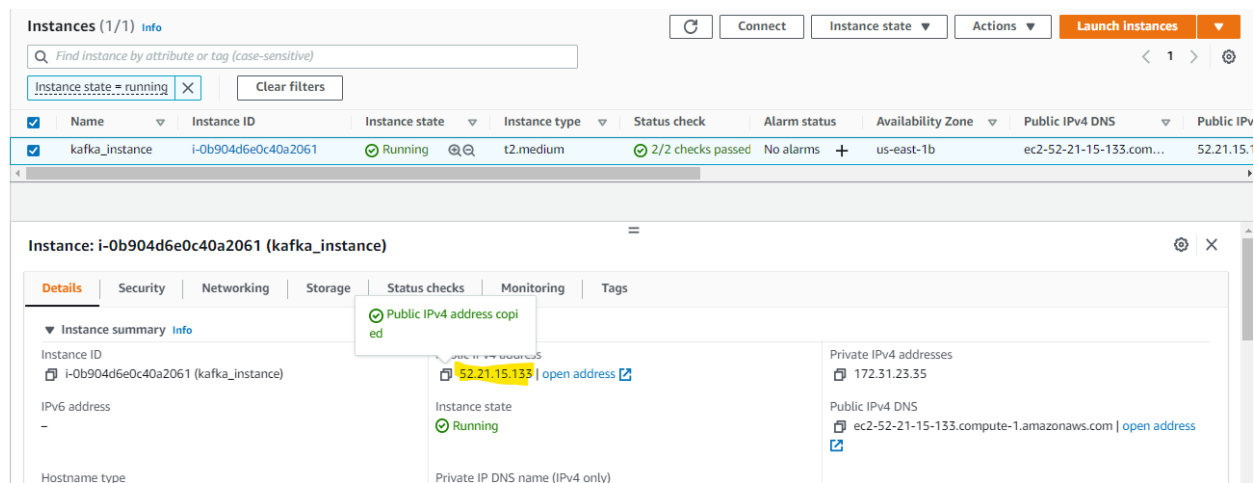
```
ec2-user@ip-172-31-23-35:~  
[I 12:20:26.898 NotebookApp] Writing notebook server cookie secret to /home/ec2-user/.local/share/jupyter/runtime/notebook_cookie_secret  
[I 12:20:28.111 NotebookApp] Serving notebooks from local directory: /home/ec2-user  
[I 12:20:28.111 NotebookApp] Jupyter Notebook 6.4.12 is running at:  
[I 12:20:28.111 NotebookApp] http://localhost:8888/?token=9736eab9490350ab3af01a811126697ead2bb40027269f9d  
[I 12:20:28.111 NotebookApp] or http://127.0.0.1:8888/?token=9736eab9490350ab3af01a811126697ead2bb40027269f9d  
[I 12:20:28.112 NotebookApp] Use Control-C to stop this server and shut down all kernels (twice to skip confirmation).  
[W 12:20:28.116 NotebookApp] No web browser found: could not locate runnable browser.  
[C 12:20:28.116 NotebookApp]  
  
To access the notebook, open this file in a browser:  
file:///home/ec2-user/.local/share/jupyter/runtime/nbserver-4760-open.html  
Or copy and paste one of these URLs:  
http://localhost:8888/?token=9736eab9490350ab3af01a811126697ead2bb40027269f9d  
or http://127.0.0.1:8888/?token=9736eab9490350ab3af01a811126697ead2bb40027269f9d
```

You will find that the Jupyter Notebook will not be launched.

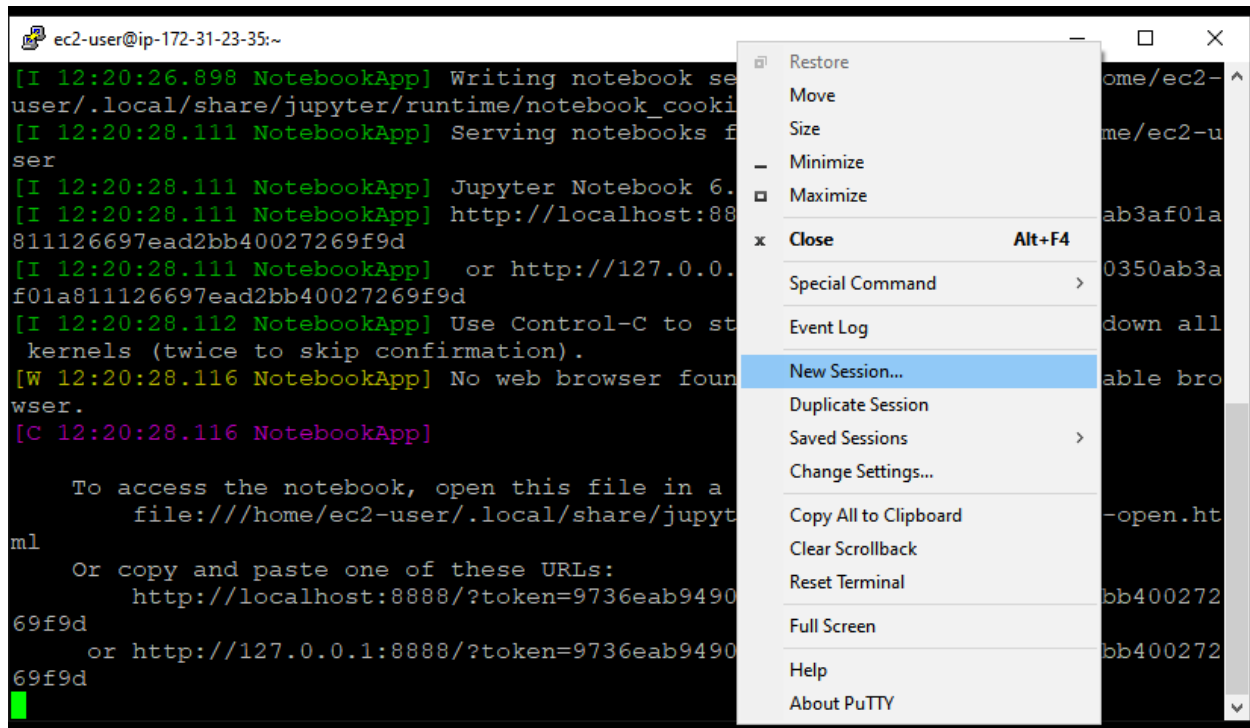


The conclusion is that we can not access this url without the SSH tunneling. For this, we will again use the services offered by PuTTY.

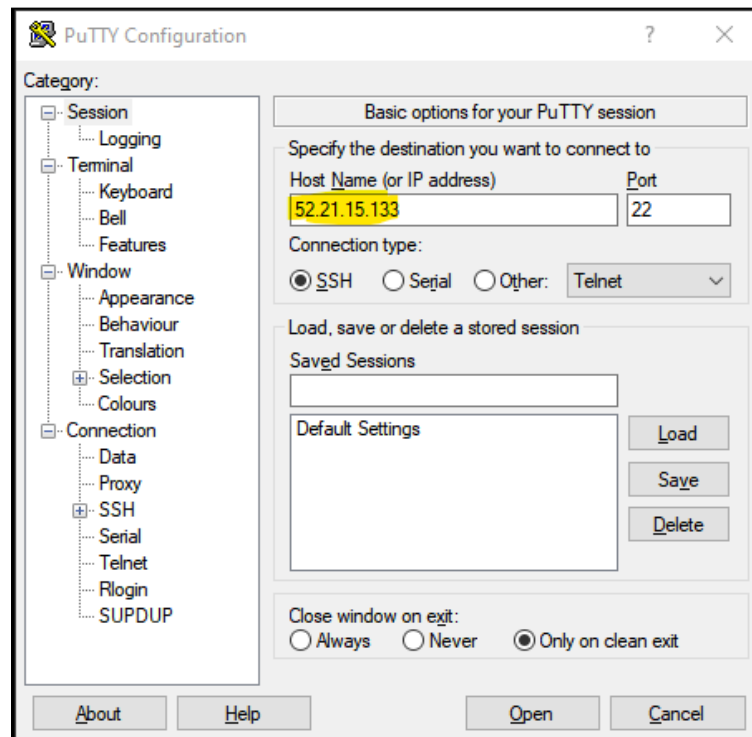
## 2. You must first copy the public IP address of the EC2 instance from the AWS Console



- Next, open a new PuTTY session by right clicking on the existing PuTTY window



- You must paste the copied Public IP address in the PuTTY window under the Host Name



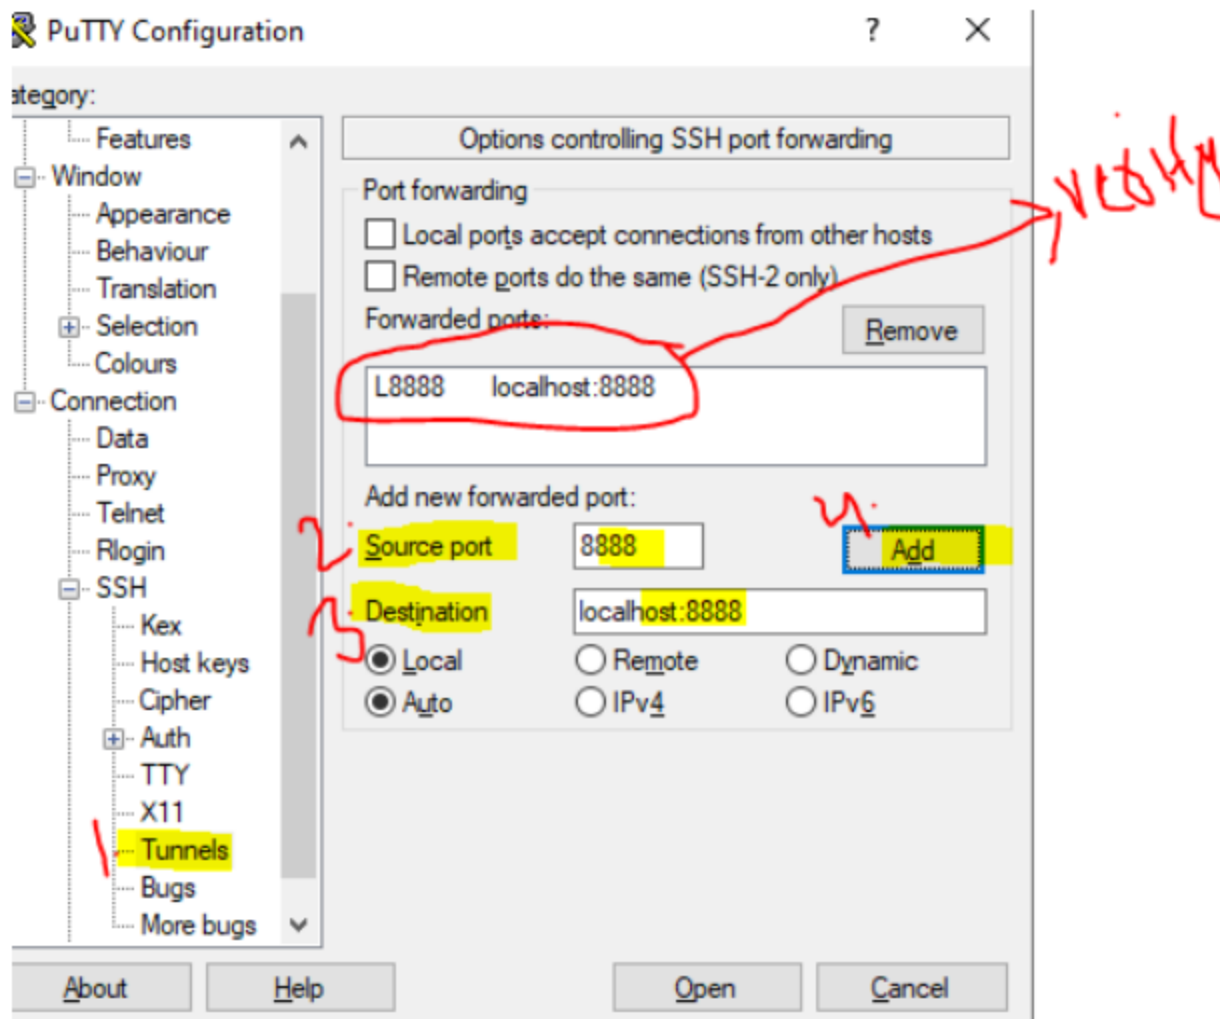
5. Next, you must establish a connection between the 8888 ports of both the machines. This port will be used in launching the Jupyter Notebook hosted on the EC2 instance on your local machine. For this, go to the 'Tunnels' section under 'SSH' to add the connection. Add the following elements as shown in the image:

### SSH > Tunnels

**Source port:** 8888

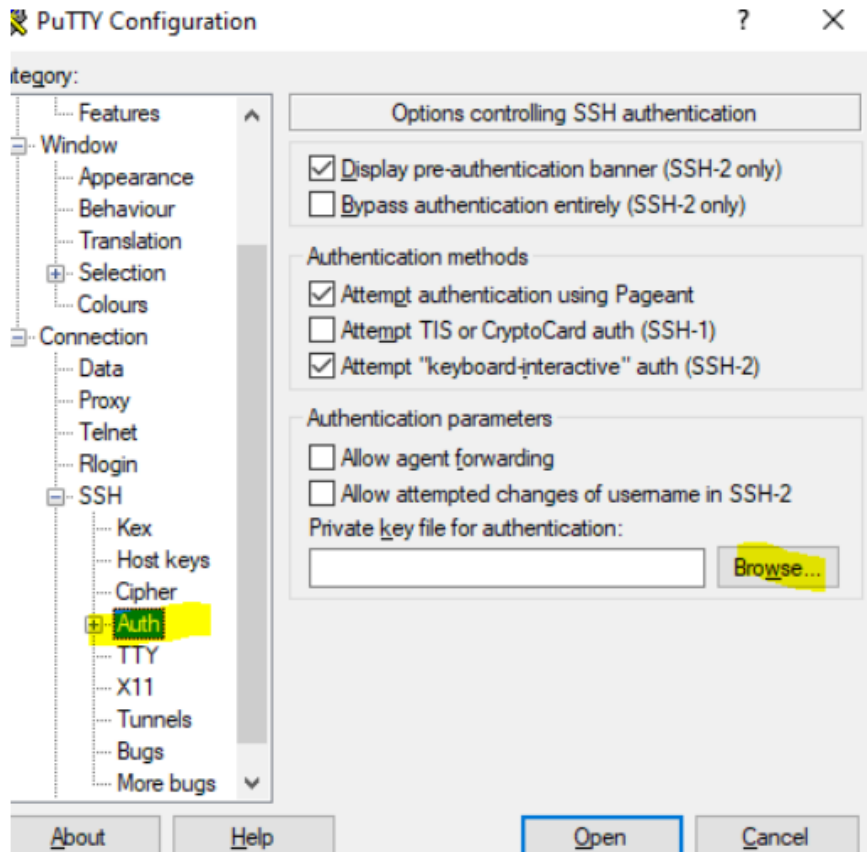
**Destination:** localhost:8888

Once you have added the details, click on the 'Add' button.

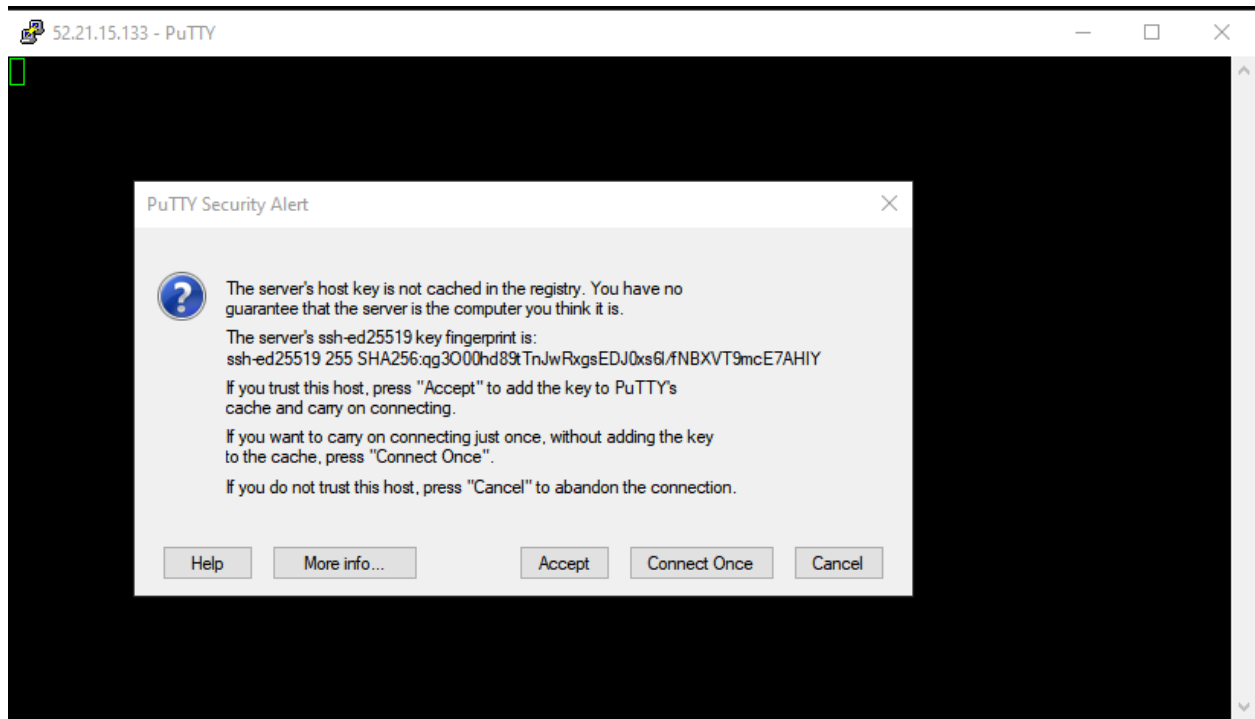


- The next step is to provide the Key pair file that you must have saved when the instance was created. You must provide the key under the Auth tab of the SSH Section.

**SSH > Auth > Browse > ppk file path**

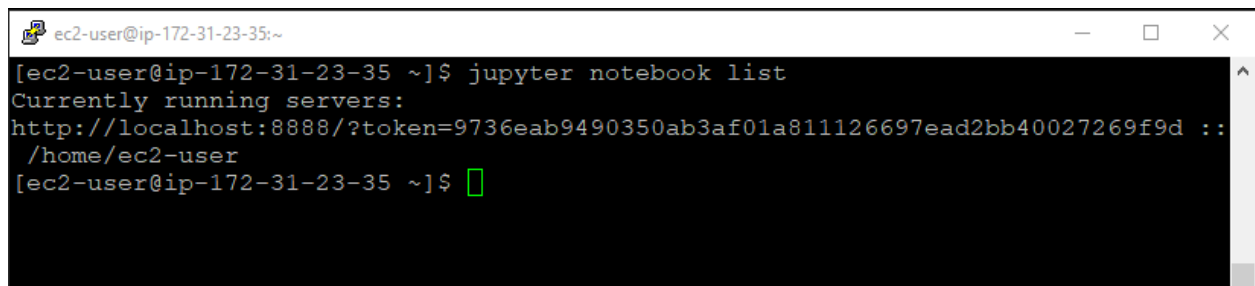


7. After you have successfully selected the file, click on 'Open' to launch the EC2 instance. Click on 'Accept' in the dialogue box when prompted.



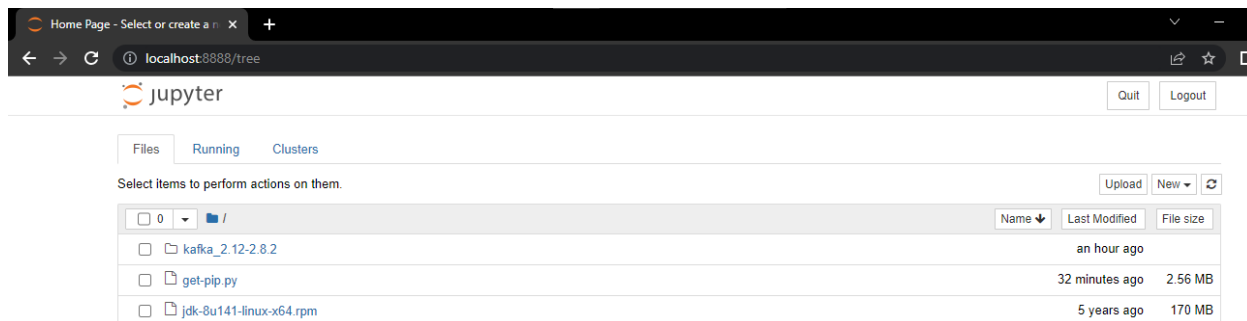
7. In the EC2 window, login with the username as 'ec2-user'. After logging in, run the following command to check the running servers:

### jupyter notebook list



You can now select the link mentioned in this window to copy, and then paste it in the local browser to access the Jupyter notebook. You can see that the notebook contains the files present in your EC2 instance.





Home Page - Select or create a notebook

localhost:8888/tree

jupyter

Files Running Clusters

Select items to perform actions on them.

Upload New

	Name	Last Modified	File size
<input type="checkbox"/>	/		
<input type="checkbox"/>	kafka_2.12-2.8.2	an hour ago	
<input type="checkbox"/>	get-pip.py	32 minutes ago	2.56 MB
<input type="checkbox"/>	jdk-8u141-linux-x64.rpm	5 years ago	170 MB

You can now upload or create any python notebooks and run codes on the EC2 instance.

- You can upload the Jupyter notebooks using the **Upload** button from your local machine to EC2 instance. Once you select the notebook, you need to upload and you get the screen as shown below.



jupyter

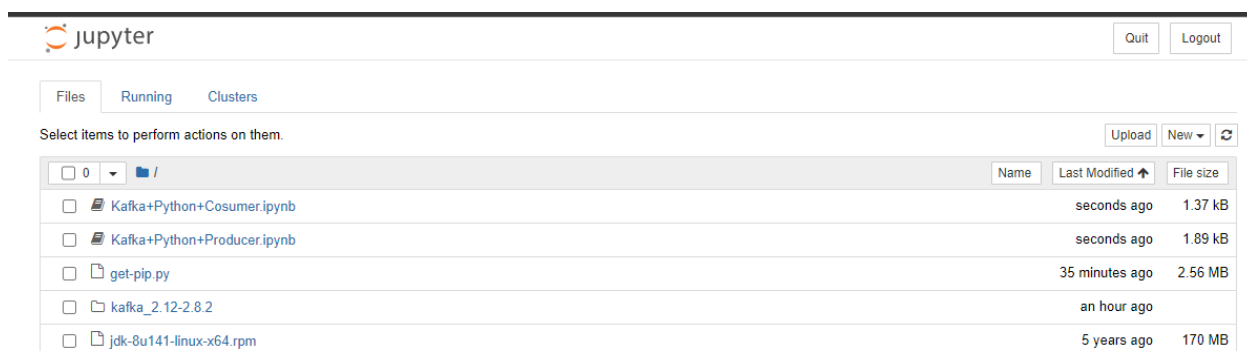
Files Running Clusters

Select items to perform actions on them.

Upload New

	Name	Last Modified	File size
<input type="checkbox"/>	/		
<input checked="" type="checkbox"/>	Kafka+Python+Producer.ipynb		
<input checked="" type="checkbox"/>	Kafka+Python+Consumer.ipynb		
<input type="checkbox"/>	kafka_2.12-2.8.2	an hour ago	
<input type="checkbox"/>	get-pip.py	35 minutes ago	2.56 MB
<input type="checkbox"/>	jdk-8u141-linux-x64.rpm	5 years ago	170 MB

- Click on the **Upload** button next to it and the notebook will get uploaded.



jupyter

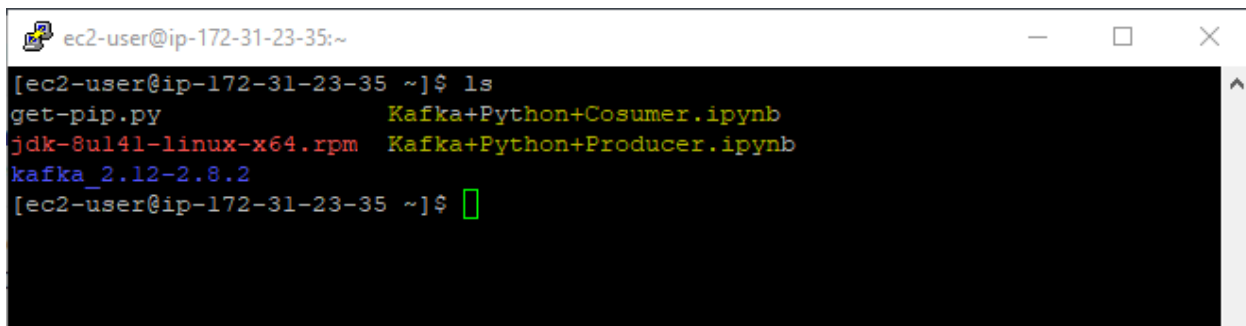
Files Running Clusters

Select items to perform actions on them.

Upload New

	Name	Last Modified	File size
<input type="checkbox"/>	/		
<input checked="" type="checkbox"/>	Kafka+Python+Consumer.ipynb	seconds ago	1.37 kB
<input checked="" type="checkbox"/>	Kafka+Python+Producer.ipynb	seconds ago	1.89 kB
<input type="checkbox"/>	get-pip.py	35 minutes ago	2.56 MB
<input type="checkbox"/>	kafka_2.12-2.8.2	an hour ago	
<input type="checkbox"/>	jdk-8u141-linux-x64.rpm	5 years ago	170 MB

- If you login to EC2 instance and enter the ls command you will see the notebook is present in the /home/ec2-user directory.



```
ec2-user@ip-172-31-23-35:~  
[ec2-user@ip-172-31-23-35 ~]$ ls  
get-pip.py          Kafka+Python+Consumer.ipynb  
jdk-8u141-linux-x64.rpm  Kafka+Python+Producer.ipynb  
kafka_2.12-2.8.2  
[ec2-user@ip-172-31-23-35 ~]$
```